

Version 2.0.0

Table of Contents

1. About This Document	2
1.1. Intended Audience	2
1.2. New and Changed Information	2
1.3. Notation Conventions	3
1.4. Comments Encouraged	5
2. Quick Start	7
2.1. Download Binaries	7
2.2. Unpack Installer	8
2.3. Collect Information	9
2.3.1. Location of Trafodion Server-Side Binary	9
2.3.2. Java Location	9
2.3.3. Data Nodes	. 10
2.3.4. Trafodion Runtime User Home Directory	. 10
2.3.5. Distribution Manager URL	. 10
2.4. Run Installer	. 11
3. Introduction	. 20
3.1. Security Considerations	. 20
3.2. Provisioning Options	. 21
3.3. Provisioning Activities	. 22
3.4. Provisioning Master Node	. 22
3.5. Trafodion Installer	. 23
3.5.1. Usage	. 24
3.5.2. Install vs. Upgrade	. 25
3.5.3. Guided Setup	. 25
3.5.4. Automated Setup	. 25
3.6. Trafodion Provisioning Directories	. 31
4. Requirements	
4.1. General Cluster and OS Requirements and Recommendations	
4.1.1. Hardware Requirements and Recommendations	. 32
4.1.2. OS Requirements and Recommendations	. 33
4.1.3. IP Ports	. 34
4.2. Prerequisite Software	
4.2.1. Hadoop Software	. 35
4.2.2. Software Packages	. 36
4.3. Trafodion User IDs and Their Privileges	. 37
4.3.1. Trafodion Runtime User	
4.3.2. Trafodion Provisioning User	. 37
4.4. Required Configuration Changes	
4.4.1. Operating System Changes	. 40
4.4.2. ZooKeeper Changes	. 41
4.4.3. HDFS Changes	. 41

111		
4.4.4	1. HBase Changes	42
4.5. Red	commended Configuration Changes	43
4.5.1	I. Recommended Security Changes	43
4.5.2	2. Recommended HDFS Configuration Changes	43
4.5.3	3. Recommended HBase Configuration Changes	44
5. Prepare		45
5.1. Inst	all Optional Workstation Software	45
5.2. Cor	nfigure Installation User ID	45
5.3. Disa	able requiretty	46
5.4. Ver	ify OS Requirements and Recommendations	46
	nfigure Kerberos	
5.6. Cor	nfigure LDAP Identity Store	47
5.7. Gat	her Configuration Information	47
5.8. Inst	all Required Software Packages	51
5.8.1	l. Download and Install Packages	51
5.9. Dov	vnload Trafodion Binaries	52
5.10. Pr	eparation for Recipe-Based Provisioning	53
5.10	.1. Modify OS Settings	53
5.10	.2. Modify ZooKeeper Configuration	53
5.10	.3. Modify HDFS Configuration	54
5.10	.4. Modify HBase Configuration	54
6. Install .		55
6.1. Unp	pack Installer	55
6.2. Gui	ded Install	56
6.3. Aut	omated Install	62
7. Upgrade		66
7.1. Dov	vnload Updated Trafodion Binaries	66
7.2. Unp	pack Installer	66
7.3. Sto	p Trafodion	66
7.4. Gui	ded Upgrade	69
7.5. Aut	omated Upgrade	73
	omatou opgituus	
8. Activate		76
8.1. Mar		76
8.1. Mar 8.2. Vali	nage Trafodion	76 76
8.1. Mar 8.2. Vali 8.2.1	nage Trafodion	76 76 76
8.1. Mar 8.2. Vali 8.2.1 8.3. Trou	nage Trafodion date Trafodion Installation Smoke Test	76 76 76 78
8.1. Mar 8.2. Vali 8.2.1 8.3. Trou 9. Remove	nage Trafodion Idate Trafodion Installation I. Smoke Test ubleshooting Tips	76 76 76 78 79
8.1. Mar 8.2. Vali 8.2.1 8.3. Trou 9. Remove 9.1. Sto	nage Trafodion date Trafodion Installation Smoke Test ubleshooting Tips	76 76 76 78 79 79
8.1. Mar 8.2. Vali 8.2.1 8.3. Trou 9. Remove 9.1. Stop 9.2. Rur	nage Trafodion date Trafodion Installation I. Smoke Test ubleshooting Tips p Trafodion	76 76 78 79 79
8.1. Mar 8.2. Vali 8.2.1 8.3. Trou 9. Remove 9.1. Sto 9.2. Rur 10. Enable	nage Trafodion date Trafodion Installation I. Smoke Test ubleshooting Tips p Trafodion n trafodion_uninstaller	76 76 78 79 79 80 81
8.1. Mar 8.2. Vali 8.2.1 8.3. Trou 9. Remove 9.1. Stop 9.2. Rur 10. Enable 10.1. Co	nage Trafodion date Trafodion Installation I. Smoke Test ubleshooting Tips p Trafodion Intrafodion_uninstaller Security	76 76 78 79 79 80 81
8.1. Mar 8.2. Vali 8.2.1 8.3. Trou 9. Remove 9.1. Stop 9.2. Rur 10. Enable 10.1. Co	nage Trafodion Idate Trafodion Installation I. Smoke Test ubleshooting Tips p Trafodion n trafodion_uninstaller Security onfiguring Trafodion for Kerberos	76 76 78 79 79 80 81 81 81

10.2. Configuring LDAP	. 8	3
10.2.1. Configuring LDAP Servers	. 8	3
10.2.2. Generate Trafodion Certificate	. 8	7
10.2.3. Creating the LDAP configuration file	. 8	8
10.2.4. Verifying configuration and users through Idapcheck	. 9	2
10.2.5. Verifying contents of configuration file through Idapconfigcheck	. 9	3
10.3. Manage Users	. 9	4
11. Install Recipe	. 9	6
12. Upgrade Recipe	. 9	7

License Statement

Licensed to the Apache Software Foundation (ASF) under one or more contributor license agreements. See the NOTICE file distributed with this work for additional information regarding copyright ownership. The ASF licenses this file to you under the Apache License, Version 2.0 (the "License"); you may not use this file except in compliance with the License. You may obtain a copy of the License at http://www.apache.org/licenses/LICENSE-2.0

Unless required by applicable law or agreed to in writing, software distributed under the License is distributed on an "AS IS" BASIS, WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied. See the License for the specific language governing permissions and limitations under the License.

Disclaimer: Apache Trafodion is an effort undergoing incubation at the Apache Software Foundation (ASF), sponsored by the Apache Incubator PMC. Incubation is required of all newly accepted projects until a further review indicates that the infrastructure, communications, and decision making process have stabilized in a manner consistent with other successful ASF projects. While incubation status is not necessarily a reflection of the completeness or stability of the code, it does indicate that the project has yet to be fully endorsed by the ASF. <<<

Revision History

Version	Date	
2.0.1	July 7, 2016	
2.0.0	June 6, 2016	
1.3.0	January, 2016	

Chapter 1. About This Document

This guide describes how to provision the end-user Trafodion binaries on top of an existing Hadoop environment. This install allows you to store and query data using Trafodion, either via Trafodion clients (see Trafodion Client Installation Guide) or via application code you write.

If you want to install a Trafodion developer-build environment, then please refer to the Trafodion Contributor Guide for instructions.

1.1. Intended Audience

This guide assumes that you are well-versed in Linux and Hadoop administration. If you don't have such experience, then you should consider going through the steps required to install a Hadoop environment before attempting to install Trafodion.

The instructions contained herein apply to the following environments.

- Single-Node Environments: Typically used when you want to evaluate Trafodion.
- Cluster (Multi-Node) Environments: Typically used when you deploy Trafodion for application usage.



Trafodion can be provisioned on a single-node or multi-node environment. Unless specifically noted, the term **cluster** is used to mean both single- and multi-node environments.

The provisioning instructions applies to a diverse set of platforms:

- Virtual Machines: Often used for evaluations and Trafodion development.
- Cloud: Used for Product Environments as well as for Developer Environments.
- Bare Metal: Used for Product Environments as well as for Developer Environments.



The term **node** is used to represent a computing platform on which operating system, Hadoop, and Trafodion software is running. Unless specifically qualified (bare-metal node, virtual-machine node, or cloud-node), **node** represents a computing platform in your cluster regardless of platform type.

1.2. New and Changed Information

This guide has been updated to include provisioning for LDAP and Kerberos.

1.3. Notation Conventions

This list summarizes the notation conventions for syntax presentation in this manual.

• UPPERCASE LETTERS

Uppercase letters indicate keywords and reserved words. Type these items exactly as shown. Items not enclosed in brackets are required.

SELECT

lowercase letters

Lowercase letters, regardless of font, indicate variable items that you supply. Items not enclosed in brackets are required.

file-name

• [] Brackets

Brackets enclose optional syntax items.

```
DATETIME [start-field TO] end-field
```

A group of items enclosed in brackets is a list from which you can choose one item or none.

The items in the list can be arranged either vertically, with aligned brackets on each side of the list, or horizontally, enclosed in a pair of brackets and separated by vertical lines.

For example:

```
DROP SCHEMA schema [CASCADE]
DROP SCHEMA schema [ CASCADE | RESTRICT ]
```

• { } Braces

Braces enclose required syntax items.

```
FROM { grantee [, grantee ] ... }
```

A group of items enclosed in braces is a list from which you are required to choose one item.

The items in the list can be arranged either vertically, with aligned braces on each side of the list, or horizontally, enclosed in a pair of braces and separated by vertical lines.

For example:

```
INTERVAL { start-field TO end-field }
{ single-field }
INTERVAL { start-field TO end-field | single-field }
```

• | Vertical Line

A vertical line separates alternatives in a horizontal list that is enclosed in brackets or braces.

```
{expression | NULL}
```

• ... Ellipsis

An ellipsis immediately following a pair of brackets or braces indicates that you can repeat the enclosed sequence of syntax items any number of times.

```
ATTRIBUTE[S] attribute [, attribute] ... {, sql-expression } ...
```

An ellipsis immediately following a single syntax item indicates that you can repeat that syntax item any number of times.

For example:

```
expression-n ...
```

Punctuation

Parentheses, commas, semicolons, and other symbols not previously described must be typed as shown.

```
DAY (datetime-expression)
@script-file
```

Quotation marks around a symbol such as a bracket or brace indicate the symbol is a required character that you must type as shown.

For example:

```
"{" module-name [, module-name] ... "}"
```

Item Spacing

Spaces shown between items are required unless one of the items is a punctuation symbol such as a parenthesis or a comma.

```
DAY (datetime-expression) DAY(datetime-expression)
```

If there is no space between two items, spaces are not permitted. In this example, no spaces are permitted between the period and any other items:

```
myfile.sh
```

· Line Spacing

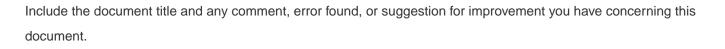
If the syntax of a command is too long to fit on a single line, each continuation line is indented three spaces and is separated from the preceding line by a blank line.

This spacing distinguishes items in a continuation line from items in a vertical list of selections.

```
match-value [NOT] LIKE _pattern
   [ESCAPE esc-char-expression]
```

1.4. Comments Encouraged

We encourage your comments concerning this document. We are committed to providing documentation that meets your needs. Send any errors found, suggestions for improvement, or compliments to user@trafodion.incubator.apache.org.



Chapter 2. Quick Start

This chapter provides a quick start for how to use the Trafodion Installer to install Trafodion.

You need the following before using the information herein:

- A supported and running Hadoop environment with HDFS, HBase, and Hive. Refer to the Release Notes for information about supported versions.
- A user ID with passwordless SSH among all the nodes in the cluster. This user ID must have sudo access.



The Trafodion Installer modifies and restarts your Hadoop environment.

2.1. Download Binaries

You download the Trafodion binaries from the Trafodion Download page. Download the following packages:

- Trafodion Installer (if planning to use the Trafodion Installer)
- Trafodion Server



You can download and install the Trafodion Clients once you've installed and activated Trafodion.

Refer to the Trafodion Client Install Guide for instructions.

Example

```
$ mkdir $HOME/trafodion-download
$ cd $HOME/trafodion-download
$ # Download the Trafodion Installer binaries
$ wget http://apache.cs.utah.edu/incubator/trafodion/trafodion-1.3.0.incubating/apache-
trafodion-installer-1.3.0-incubating-bin.tar.gz
Resolving http://apache.cs.utah.edu... 192.168.1.56
Connecting to http://apache.cs.utah.edu|192.168.1.56|:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 68813 (67K) [application/x-gzip]
Saving to: "apache-trafodion-installer-1.3.0-incubating-bin.tar.gz"
124K/s
                                                  in 0.5s
2016-02-14 04:19:42 (124 KB/s) - "apache-trafodion-installer-1.3.0-incubating-
bin.tar.gz" saved [68813/68813]
```

```
$ # Download the Trafodion Server binaries
$ wget http://apache.cs.utah.edu/incubator/trafodion/trafodion-1.3.0.incubating/apache-
trafodion-1.3.0-incubating-bin.tar.gz
Resolving http://apache.cs.utah.edu... 192.168.1.56
Connecting to http://apache.cs.utah.edu|192.168.1.56|:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 214508243 (205M) [application/x-gzip]
Saving to: "apache-trafodion-1.3.0-incubating-bin.tar.gz"
2016-02-14 04:22:14 (3.72 MB/s) - "apache-trafodion-1.3.0-incubating-bin.tar.gz" saved
[214508243/214508243]
$ ls -1
total 209552
-rw-rw-r-- 1 centos centos 214508243 Jan 12 20:10 apache-trafodion-1.3.0-incubating-
bin.tar.gz
-rw-rw-r-- 1 centos centos
                          68813 Jan 12 20:10 apache-trafodion-installer-1.3.0-
incubating-bin.tar.gz
```

2.2. Unpack Installer

The first step in the installation process is to unpack the Trafodion Installer tar file.

Example

```
$ mkdir $HOME/trafodion-installer
$ cd $HOME/trafodion-downloads
$ tar -zxf apache-trafodion-installer-1.3.0-incubating-bin.tar.gz -C $HOME/trafodion-
installer
$ ls $HOME/trafodion-installer/installer
bashrc_default
                                                           traf_config_check
trafodion_apache_hadoop_install traf_package_setup
build-version-1.3.0.txt traf_add_user
                                                           traf_config_setup
trafodion_config_default
                                 traf_setup
dcs_installer
                        traf_apache_hadoop_config_setup
                                                          traf_create_systemdefaults
trafodion_install
                                 traf_sqconfig
rest_installer
                         traf_authentication_conf_default traf_getHadoopNodes
trafodion_license
                                 traf_start
setup_known_hosts.exp
                         traf_cloudera_mods98
                                                           traf_hortonworks_mods98
trafodion uninstaller
```

2.3. Collect Information

Collect/decide the following information:

2.3.1. Location of Trafodion Server-Side Binary

You need the fully-qualified name of the Trafodion server-side binary.

Example

/home/trafodion-downloads/apache-trafodion-installer-1.3.0-incubating-bin.tar.gz

2.3.2. Java Location

You need to record the location of the Java. For example, use ps -ef | grep java | grep hadoop | grep hbase to determine what version HBase is running.

Example

```
ps -ef | grep java | grep hadoop | grep hbase
         17302 17288 1 20:35 ?
                                         00:00:10 /usr/jdk64/jdk1.7.0_67/bin/java
-Dproc_master -XX:OnOutOfMemoryError=kill -9 %p -Dhdp.version=2.3.6.0-3796
-XX:+UseConcMarkSweepGC -XX:ErrorFile=/var/log/hbase/hs_err_pid%p.log
-Djava.io.tmpdir=/tmp -verbose:gc -XX:+PrintGCDetails -XX:+PrintGCDateStamps
-Xloggc:/var/log/hbase/gc.log-201606302035 -Xmx1024m -XX:PermSize=128m
-XX:MaxPermSize=128m -Dhbase.log.dir=/var/log/hbase -Dhbase.log.file=hbase-hbase-master
-ip-172-31-56-238.log -Dhbase.home.dir=/usr/hdp/current/hbase-master/bin/..
-Dhbase.id.str=hbase -Dhbase.root.logger=INFO,RFA -Djava.library.path=:/usr/hdp/2.3.6.0
-3796/hadoop/lib/native/Linux-amd64-64:/usr/hdp/2.3.6.0-3796/hadoop/lib/native
-Dhbase.security.logger=INFO,RFAS org.apache.hadoop.hbase.master.HMaster start
```

The Java location is: /usr/jdk64/jdk1.7.0_67

2.3.3. Data Nodes

{projet-name} is installed on all data nodes in your Hadoop cluster. You need to record the fully-qualified domain name node for each node. For example, refer to /etc/hosts.

Example

```
$ cat /etc/hosts
127.0.0.1 localhost localhost.localdomain localhost4 localhost4.localdomain4
::1 localhost localhost.localdomain localhost6 localhost6.localdomain6

172.31.56.238 ip-172-31-56-238.ec2.internal node01
172.31.61.110 ip-172-31-61-110.ec2.internal node02
172.31.57.143 ip-172-31-57-143.ec2.internal node03
```

Record the node names in a space-separated list.

Example

```
ip-172-31-56-238.ec2.internal ip-172-31-61-110.ec2.internal ip-172-31-57-143.ec2.internal
```

2.3.4. Trafodion Runtime User Home Directory

The Installer creates the trafodion user ID. You need to decide the home directory for this user.

The default is: /home

2.3.5. Distribution Manager URL

The Installer interacts with the Distribution Manager (for example, Apache Ambari or Cloudera Manager) to modify the Hadoop configuration.

Example

Apache Ambari URL

```
http://myhost.com:8080
```

2.4. Run Installer

You run the Installer once you've collected the base information as described in Collect Information above.

The following example shows a guided install of Trafodion on a three-node Hortonworks Hadoop cluster.



By default, the Trafodion Installer invokes sqlci so that you can enter the initialize trafodion; command. This is shown in the example below.

Example

1. Run the Trafodion Installer in guided mode.

```
$ cd $HOME/trafodion-installer/installer
$ ./trafodion_install 2>&1 | tee install.log
*******
TRAFODION INSTALLATION START
********
***INFO: testing sudo access
***INFO: Log file located at /var/log/trafodion/trafodion_install_2016-06-30-21-02-
38.log
***INFO: Config directory: /etc/trafodion
***INFO: Working directory: /usr/lib/trafodion
**********
Trafodion Configuration File Setup
*********
***INFO: Please press [Enter] to select defaults.
Is this a cloud environment (Y/N), default is [N]: N
Enter trafodion password, default is [traf123]:
Enter list of data nodes (blank separated), default []: ip-172-31-56-
238.ec2.internal ip-172-31-61-110.ec2.internal ip-172-31-57-143.ec2.internal
Do you have a set of management nodes (Y/N), default is N: N
Enter Trafodion userid's home directory prefix, default is [/home]: /opt
Specify location of Java 1.7.0_65 or higher (JDK), default is []:
/usr/jdk64/jdk1.7.0_67
Enter full path (including .tar or .tar.gz) of trafodion tar file []:
/home/trafodion-downloads/apache-trafodion_server-2.0.1-incubating.tar.gz
Enter Backup/Restore username (can be Trafodion), default is [trafodion]:
Specify the Hadoop distribut ion installed (1: Cloudera, 2: Hortonworks, 3: Other):
Enter Hadoop admin username, default is [admin]: Enter Hadoop admin pas sword,
default is [admin]:
Enter full Hadoop external network URL:port (include 'http://' or 'https://),
default is []: http://ip-172-31-56-238.ec2.internal:8080
Enter HDFS username or username running HDFS, default is [hdfs]:
```

```
Enter HBase username or username running HBase, default is [hbase]:
Enter HBase group, default is [hbase]:
Enter Zookeeper username or username running Zookeeper, default is [zookeeper]:
Enter directory to install trafodion to, default is [/opt/trafodion/apache-
trafodion_server-2.0.1-incubating]:
Start Trafodion after install (Y/N), default is Y:
Total number of client connections per cluster, default [24]: 96
Enter the node of primary DcsMaste r, default [ip-172-31-56-238.ec2.internal]:
Enable High Availability (Y/N), default is N:
Enable simple LDAP security (Y/N), d efault is N:
***INFO: Trafodion configuration setup complete
***INFO: Trafodion Configuration File Check
***INFO: Testing sudo access on node ip-172-31-56-238
***INFO: Testing sudo access on node ip-172-31-61-110
***INFO: Testing sudo access on node ip-172-31-57-143
***INFO: Testing ssh on ip-172-31-56-238
***INFO: Testing ssh on ip-172-31-61-110
***INFO: Testing ssh on ip-172-31-57-143
#!/bin/bash
# @@@ START COPYRIGHT @@@
# Licensed to the Apache Software Foundation (ASF) under one
\ensuremath{\sharp} or more contributor license agreements. See the NOTICE file
# distributed with this work for additional information
# regarding copyright ownership. The ASF licenses this file
# to you under the Apache License, Version 2.0 (the
# "License"); you may not use this file except in compliance
# with the License. You may obtain a copy of the License at
#
#
   http://www.apache.org/licenses/LICENSE-2.0
# Unless required by applicable law or agreed to in writing,
# software distributed under the License is distributed on an
# "AS IS" BASIS, WITHOUT WARRANTIES OR CONDITIONS OF ANY
# KIND, either express or implied. See the License for the
# specific language governing permissions and limitations
# under the License.
9. Accepting Warranty or Additional Liability. While redistributing
the Work or Derivative Works thereof, You may choose to offer, and
charge a fee for, acceptance of support, warranty, indemnity, or
other liability obligations and/or rights consistent with this
License. However, in accepting such obligations, You may act only
on Your own behalf and on Your sole responsibility, not on behalf
of any other Contributor, and only if You agree to indemnify, defend,
and hold each Contributor harmless for any liability incurred by,
or claims asserted against, such Contributor by reason of your
accepting any such warranty or additional liability.
END OF TERMS AND CONDITIONS
```

```
BY TYPING "ACCEPT" YOU AGREE TO THE TERMS OF THIS AGREEMENT: ***INFO: testing sudo
***INFO: Starting Trafodion Package Setup (2016-06-30-21-06-40)
***INFO: Installing required packages
***INFO: Log file located in /var/log/trafodion
***INFO: ... pdsh on node ip-172-31-56-238
***INFO: ... pdsh on node ip-172-31-61-110
***INFO: ... pdsh on node ip-172-31-57-143
***INFO: Checking if apr is installed ...
***INFO: Checking if apr-util is installed ...
***INFO: Checking if sqlite is installed ...
***INFO: Checking if expect is installed ...
***INFO: Checking if perl-DBD-SQLite* is installed ...
***INFO: Checking if protobuf is installed ...
***INFO: Checking if xerces-c is installed ...
***INFO: Checking if perl-Params-Validate is installed ...
***INFO: Checking if perl-Time-HiRes is installed ...
***INFO: Checking if gzip is installed ...
***INFO: Checking if lzo is installed \dots
***INFO: Checking if lzop is installed ...
***INFO: Checking if unzip is installed ...
***INFO: modifying limits in /usr/lib/trafodion/trafodion.conf on all nodes
***INFO: create Trafodion userid "trafodion"
***INFO: Trafodion userid's (trafodion) home directory: /opt/trafodion
***INFO: testing sudo access
Generating public/private rsa key pair.
Created directory '/opt/trafodion/.ssh'.
Your identification has been saved in /opt/trafodion/.ssh/id rsa.
Your public key has been saved in /opt/trafodion/.ssh/id_rsa.pub.
The key fingerprint is:
12:59:ab:d7:59:a2:0e:e8:38:1c:e9:e1:86:f6:18:23 trafodion@ip-172-31-56-238
The key's randomart image is:
+--[ RSA 2048]----+
       ο.
      0...
    . . 0 0 +
  + . + S o
 = = =
E+B .
0.=.
***INFO: creating .bashrc file
***INFO: Setting up userid trafodion on all other nodes in cluster
***INFO: Creating known_hosts file for all nodes
ip-172-31-56-238
ip-172-31-56-238 ip-172-31-61-110 ip-172-31-57-143
ip-172-31-61-110
ip-172-31-56-238 ip-172-31-61-110 ip-172-31-57-143
ip-172-31-57-143
ip-172-31-56-238 ip-172-31-61-110 ip-172-31-57-143
***INFO: trafodion user added successfully
```

```
***INFO: Trafodion environment setup completed
***INFO: creating sqconfig file
***INFO: Reserving DCS ports
***INFO: Creating trafodion sudo access file
*********
TRAFODION MODS
*********
***INFO: Hortonworks installed will run traf_hortonworks_mods
***INFO: copying hbase-trx-hdp2_3-*.jar to all nodes
***INFO: hbase-trx-hdp2_3-*.jar copied correctly! Huzzah.
USERID=admin
PASSWORD=admin
PORT=:8080
######## Performing 'set' hbase.master.distributed.log.splitting:true on
(Site:hbase-site, Tag:version1)
######## PUTting json into: doSet_version1467320863286001262.json
 % Total % Received % Xferd Average Speed Time Time Current
                               Dload Upload Total Spent Left Speed
101 5757 102 3588 103 2169 98k 60930 --:---- --:-- 100k
  "resources" : [
     "href": "http://ip-172-31-56-
238.ec2.internal:8080/api/v1/clusters/trafodion/configurations/service_config_versio
vice_name=HBASE&service_config_version=2",
     "href" : "http://ip-172-31-56-
238.ec2.internal:8080/api/v1/clusters/trafodion/requests/12/tasks/128",
     "Tasks" : {
       "cluster_name" : "trafodion",
       "id" : 128,
       "request id" : 12,
       "stage_id" : 2
     }
   },
     "href": "http://ip-172-31-56-
238.ec2.internal:8080/api/v1/clusters/trafodion/requests/12/tasks/129",
     "Tasks" : {
       "cluster name" : "trafodion",
       "id" : 129,
       "request_id" : 12,
       "stage_id" : 2
     }
   },
```

```
"href": "http://ip-172-31-56-
238.ec2.internal:8080/api/v1/clusters/trafodion/requests/12/tasks/130",
      "Tasks" : {
        "cluster_name" : "trafodion",
        "id" : 130,
        "request_id" : 12,
        "stage_id" : 2
    }
  ],
  "stages" : [
      "href" : "http://ip-172-31-56-
238.ec2.internal:8080/api/v1/clusters/trafodion/requests/12/stages/0",
      "Stage" : {
        "cluster_name" : "trafodion",
        "request_id" : 12,
        "stage id" : 0
      }
    },
      "href" : "http://ip-172-31-56-
238.ec2.internal:8080/api/v1/clusters/trafodion/requests/12/stages/1",
      "Stage" : {
        "cluster name" : "trafodion",
        "request_id" : 12,
        "stage_id" : 1
    },
      "href" : "http://ip-172-31-56-
238.ec2.internal:8080/api/v1/clusters/trafodion/requests/12/stages/2",
      "Stage" : {
        "cluster_name" : "trafodion",
       "request_id" : 12,
        "stage_id" : 2
      }
  ]
}***INFO: ...polling every 30 seconds until HBase start is completed.
***INFO: HBase restart completed
***INFO: Setting HDFS ACLs for snapshot scan support
cp: `trafodion_config' and `/home/trafinstall/trafodion-
2.0.1/installer/trafodion_config' are the same file
***INFO: Trafodion Mods ran successfully.
 TRAFODION CONFIGURATION
*******
/usr/lib/trafodion/installer/..
/opt/trafodion/apache-trafodion_server-2.0.1-incubating
***INFO: untarring file to /opt/trafodion/apache-trafodion_server-2.0.1-incubating
***INFO: modifying .bashrc to set Trafodion environment variables
```

```
***INFO: copying .bashrc file to all nodes
***INFO: copying sqconfig file (/opt/trafodion/sqconfig) to /opt/trafodion/apache-
trafodion_server-2.0.1-incubating/sql/script
s/sqconfiq
***INFO: Creating /opt/trafodion/apache-trafodion_server-2.0.1-incubating directory
on all nodes
***INFO: Start of DCS install
***INFO: DCS Install Directory: /opt/trafodion/apache-trafodion_server-2.0.1-
***INFO: modifying /opt/trafodion/apache-trafodion_server-2.0.1-incubating/dcs-
2.0.1/conf/dcs-env.sh
***INFO: modifying /opt/trafodion/apache-trafodion_server-2.0.1-incubating/dcs-
2.0.1/conf/dcs-site.xml
***INFO: creating /opt/trafodion/apache-trafodion_server-2.0.1-incubating/dcs-
2.0.1/conf/servers file
***INFO: End of DCS install.
***INFO: Start of REST Server install
***INFO: Rest Install Directory: /opt/trafodion/apache-trafodion_server-2.0.1-
incubating/rest-2.0.1
***INFO: modifying /opt/trafodion/apache-trafodion_server-2.0.1-incubating/rest-
2.0.1/conf/rest-site.xml
***INFO: End of REST Server install.
***INFO: starting sqgen
ip-172-31-56-238, ip-172-31-57-143, ip-172-31-61-110
Creating directories on cluster nodes
/usr/bin/pdsh -R exec -w ip-172-31-56-238,ip-172-31-57-143,ip-172-31-61-110 -x ip-
172-31-56-238 ssh -q -n %h mkdir -p /opt/tra
fodion/apache-trafodion server-2.0.1-incubating/etc
/usr/bin/pdsh -R exec -w ip-172-31-56-238,ip-172-31-57-143,ip-172-31-61-110 -x ip-
172-31-56-238 ssh -q -n %h mkdir -p /opt/tra
fodion/apache-trafodion_server-2.0.1-incubating/logs
/usr/bin/pdsh -R exec -w ip-172-31-56-238,ip-172-31-57-143,ip-172-31-61-110 -x ip-
172-31-56-238 ssh -q -n %h mkdir -p /opt/tra
fodion/apache-trafodion_server-2.0.1-incubating/tmp
/usr/bin/pdsh -R exec -w ip-172-31-56-238,ip-172-31-57-143,ip-172-31-61-110 -x ip-
172-31-56-238 ssh -q -n %h mkdir -p /opt/tra
fodion/apache-trafodion_server-2.0.1-incubating/sql/scripts
Generating SQ environment variable file: /opt/trafodion/apache-trafodion_server-
2.0.1-incubating/etc/ms.env
Note: Using cluster.conf format type 2.
Generating SeaMonster environment variable file: /opt/trafodion/apache-
trafodion_server-2.0.1-incubating/etc/seamonster.env
Generated SQ startup script file: ./gomon.cold
Generated SQ startup script file: ./gomon.warm
Generated SQ cluster config file: /opt/trafodion/apache-trafodion_server-2.0.1-
incubating/tmp/cluster.conf
Generated SQ Shell
                          file: sqshell
Generated RMS Startup file: rmsstart
```

```
Generated RMS Stop file: rmsstop
Generated RMS Check
                        file: rmscheck.sql
Generated SSMP Startup
                        file: ssmpstart
Generated SSMP Stop
                        file: ssmpstop
Generated SSCP Startup file: sscpstart
Generated SSCP Stop file: sscpstop
Copying the generated files to all the nodes in the cluster
SQ Startup script (/opt/trafodion/apache-trafodion_server-2.0.1-
incubating/sql/scripts/gomon.cold) ran successfully. Performin
g further checks...
Checking if processes are up.
Checking attempt: 1; user specified max: 2. Execution time in seconds: 0.
The SQ environment is up!
         Configured Actual Down
Process
          -----
      3
             3
DTM
             6
RMS 6
DcsMaster 1 0
                        1
DcsServer 3
                 0
                         3
mxosrvr 96
                 0
                        96
Thu Jun 30 21:15:29 UTC 2016
Checking if processes are up.
Checking attempt: 1; user specified max: 1. Execution time in seconds: 0.
The SQ environment is up!
Process
         Configured Actual
                               Down
           _____
                                 ____
             3
      3
DTM
     6
             6
DcsMaster 1 0
                        1
DcsServer 3
                 0
                         3
mxosrvr 96
                 0
Starting the DCS environment now
starting master, logging to /opt/trafodion/apache-trafodion_server-2.0.1-
incubating/dcs-2.0.1/bin/../logs/dcs-trafodion-1-mast
er-ip-172-31-56-238.out
ip-172-31-56-238: starting server, logging to /opt/trafodion/apache-
trafodion_server-2.0.1-incubating/dcs-2.0.1/bin/../logs/dc
s-trafodion-1-server-ip-172-31-56-238.out
ip-172-31-57-143: starting server, logging to /opt/trafodion/apache-
trafodion_server-2.0.1-incubating/dcs-2.0.1/bin/../logs/dc
s-trafodion-3-server-ip-172-31-57-143.out
```

```
ip-172-31-61-110: starting server, logging to /opt/trafodion/apache-
trafodion_server-2.0.1-incubating/dcs-2.0.1/bin/../logs/dc
s-trafodion-2-server-ip-172-31-61-110.out
Checking if processes are up.
Checking attempt: 1; user specified max: 2. Execution time in seconds: 1.
The SQ environment is up!
Process
         Configured Actual
                                Down
           -----
                                ____
     3
6
             3
DTM
             6
RMS
DcsMaster 1 1
DcsServer 3 3
mxosrvr 96 7
                         89
Starting the REST environment now
starting rest, logging to /opt/trafodion/apache-trafodion_server-2.0.1-
incubating/rest-2.0.1/bin/../logs/rest-trafodion-1-rest
-ip-172-31-56-238.out
Zookeeper listen port: 2181
DcsMaster listen port: 23400
Configured Primary DcsMaster: "ip-172-31-56-238.ec2.internal"
Active DcsMaster : "ip-172-31-56-238"
Process Configured Actual Down
DcsMaster 1 1
DcsServer 3
                 3
mxosrvr 96
                 94 2
You can monitor the SQ shell log file : /opt/trafodion/apache-trafodion_server-
2.0.1-incubating/logs/sqmon.log
Startup time 0 hour(s) 2 minute(s) 19 second(s)
Apache Trafodion Conversational Interface 2.0.1
Copyright (c) 2015-2016 Apache Software Foundation
--- SQL operation complete.
>>
End of MXCI Session
***INFO: Installation setup completed successfully.
*******
 TRAFODION INSTALLATION END
```

```
*********
```

2. Switch to the Trafodion Runtime User and check the status of Trafodion.

```
$ sudo su - trafodion
$ sqcheck
Checking if processes are up.
Checking attempt: 1; user specified max: 2. Execution time in seconds: 0.
The SQ environment is up!
Process Configured Actual Down
DTM 3 3
RMS 6 6
DcsMaster 1 1
DcsServer 3
                 3
DcsServer 3 3
mxosrvr 96 96
$
```

Trafodion is now running on your Hadoop cluster. Please refer to the Activate chapter for basic instructions on how to verify the Trafodion management and how to perform basic management operations.

Chapter 3. Introduction

Trafodion is a Hadoop add-on service that provides transactional SQL on top of HBase. Typically, you use Trafodion as the database for applications that require Online Transaction Processing (OLTP), Operational Data Store (ODS), and/or strong reporting capabilities. You access Trafodion using standard JDBC and ODBC APIs.

You may choose whether to add Trafodion to an existing Hadoop environment or to create a standalone Hadoop environment specifically for Hadoop.

This guide assumes that a Hadoop environment exists upon which your provisioning Trafodion. Refer to Hadoop Software for information about what Hadoop software is required Trafodion.

3.1. Security Considerations

The following users and principals need be considered for Trafodion:

- Provisioning User: A Linux-level user that performs the Trafodion provisioning tasks. This user ID requires sudo
 access and passwordless ssh among the nodes where Trafodion is installed. In addition, this user ID requires access
 to Hadoop distribution, HDFS, and HBase administrative users to change respective environment's configuration
 settings per Trafodion requirements. Refer to Trafodion Provisioning User for more information about the requirements
 and usage associated with this user ID.
- Runtime User: A Linux-level user under which the Trafodion software runs. This user ID must be registered as a user in the Hadoop Distributed File System (HDFS) to store and access objects in HDFS, HBase, and Hive. In addition, this user ID requires passwordless access among the nodes where Trafodion is installed. Refer to Trafodion Runtime User for more information about this user ID.
- Trafodion Database Users: Trafodion users are managed by Trafodion security features (grant, revoke, etc.), which
 can be integrated with LDAP if so desired. These users are referred to as database users and do not have direct
 access to the operating system. Refer to LDAP for details on enabling LDAP for authenticating database users. Refer
 to Register User, Grant, and other SQL statements in the Trafodion SQL Reference Manual for more information about
 managing Trafodion Database Users.

If your environment has been provisioned with Kerberos, then the following additional information is required.

 KDC admin principal: Trafodion requires administrator access to Kerberos to create principals and keytabs for the trafodion user, and to look-up principal names for HDFS and HBase keytabs. Refer to Kerberos for more information about the requirements and usage associated with this principal.

- HBase keytab location: Trafodion requires administrator access to HBase to grant required privileges to the trafodion user. Refer to Kerberos for more information about the requirements and usage associated with this keytab.
- HDFS keytab location: Trafodion requires administrator access to HDFS to create directories that store files needed to perform SQL requests such as data loads and backups. Refer to Kerberos for more information about the requirements and usage associated with this keytab.

If your environment is using LDAP for authentication, then the following additional information is required.

- LDAP username for database root access: When Trafodion is installed, it creates a predefined database user referred to as the DB__ROOT user. In order to connect to the database as database root, there must be a mapping between the database user DB__ROOT and an LDAP user. Refer to LDAP for more information about this option.
- LDAP search user name: Trafodion optionally requests an LDAP username and password in order to perform LDAP operations such as LDAP search. Refer to LDAP for more information about this option.

3.2. Provisioning Options

Trafodion ships with a set of scripts that takes care of many of the installation and upgrade tasks associated with the Trafodion software and its requirements. The main features include:

- **Trafodion installer**: Performs installation and upgrade for Trafodion
- Trafodion uninstaller: Uninstalls Trafodion
- Trafodion security installer: Enables security features including Kerberos and LDAP for an existing Trafodion installation

Currently, the Trafodion Installer is able to install Trafodion on select Cloudera and Hortonworks Hadoop distributions, and for select vanilla Hadoop installations. The Trafodion Installer limitations are noted as they apply in the different chapters below. For example, the Trafodion Installer is less capable on SUSE than it is on RedHat/CentOS; you have to install the prerequisite software packages outside the Trafodion Installer.

The Trafodion Installer automates many of the tasks required to install/upgrade Trafodion, spanning from downloading and installing required software packages and making required changes to your Hadoop environment via creating the Trafodion runtime user ID to installing and starting Trafodion. It is, therefore, highly recommend that you use the Trafodion Installer for initial installation and upgrades of Trafodion. These steps are referred to as "Script-Based Provisioning" in this guide. Refer to Trafodion Installer that provides usage information.

If, for any reason, you choose not to use the Trafodion Installer, then separate chapters provide step-by-step recipes for

the tasks required to install/upgrade Trafodion. These steps are referred to as **Recipe-Based Provisioning** in this guide. It is assumed that you are well-versed in Linux and Hadoop administrative tasks if using Recipe-Based Provisioning.

3.3. Provisioning Activities

Trafodion provisioning is divided into the following main activities:

- Requirements: Activities and documentation required to install the Trafodion software. These activities include tasks such as understanding hardware and operating system requirements, Hadoop requirements, what software packages that need to be downloaded, configuration settings that need to be changed, and user ID requirements.
- Prepare: Activities to prepare the operating system and the Hadoop ecosystem to run Trafodion. These activities
 include tasks such as installing required software packages, configure the Trafodion Installation User, gather
 information about the Hadoop environment, and the modify configuration for different Hadoop services.
- Install: Activities related to installing the Trafodion software. These activities include tasks such as unpacking the Trafodion tar files, creating the Trafodion Runtime User, creating Trafodion HDFS directories, installing the Trafodion software, and enabling security features.
- **Upgrade**: Activities related to the upgrading the Trafodion software. These activities include tasks such as shutting down Trafodion and installing a new version of the Trafodion software. The upgrade tasks vary depending on the differences between the current and new release of Trafodion. For example, an upgrade may or may not include an upgrade of the Trafodion metadata.
- Activate: Activities related to starting the Trafodion software. These actives include basic management tasks such as starting and checking the status of the Trafodion components and performing basic smoke tests.
- Remove: Activities related to removing Trafodion from your Hadoop cluster.
- Enable Security: Activities related to enabling security features on an already installed Trafodion installation. These activities include tasks such as adding Kerberos principals and keytabs, and setting up the LDAP configuration files.

3.4. Provisioning Master Node

All provisioning tasks are performed from a single node in the cluster, which must be part of the Hadoop environment you're adding Trafodion to. This node is referred to as the "**Provisioning Master Node**" in this guide.

The Trafodion Provisioning User must have access to all other nodes from the Provisioning Master Node in order to perform provisioning tasks on the cluster.

3.5. Trafodion Installer

The Trafodion Installer is a set of scripts automates most of the tasks requires to install/upgrade Trafodion. You download the Trafodion Installer tar file from the Trafodion download page. Next, you unpack the tar file.

Example

```
$ mkdir $HOME/trafodion-installer
$ cd $HOME/trafodion-downloads
$ tar -zxf apache-trafodion-installer-1.3.0-incubating-bin.tar.gz -C $HOME/trafodion-
installer
```

The Trafodion Installer supports two different modes:

- Guided Setup: Prompts for information as it works through the installation/upgrade process. This mode is recommended for new users.
- 2. Automated Setup: Required information is provided in a pre-formatted bash-script configuration file, which is provided via a command argument when running the Trafodion Installer thereby suppressing all prompts. There is one exception, if Kerberos is enabled on the cluster, then you will always be prompted for the KDC admin password. We do not store the KDC admin password as part of installation anywhere.

A template of the configuration file is available here within the installer directory: trafodion_config_default. Make a copy of the file in your directory and populate the needed information.

Automated Setup is recommended since it allows you to record the required provisioning information ahead of time. Refer to Automated Setup for information about how to populate this file.

3.5.1. Usage

The following shows help for the Trafodion Installer.

3.5.2. Install vs. Upgrade

The Trafodion Installer automatically detects whether you're performing an install or an upgrade by looking for the Trafodion Runtime User in the /etc/passwd file.

- If the user ID doesn't exist, then the Trafodion Installer runs in install mode.
- If the user ID exists, then the Trafodion Installer runs in upgrade mode.

3.5.3. Guided Setup

By default, the Trafodion Installer runs in Guided Setup mode, which means that it prompts you for information during the install/upgrade process.

Refer to the following sections for examples:

- Guided Install
- Guided Upgrade

3.5.4. Automated Setup

The --config_file option runs the Trafodion in Automated Setup mode.

Before running the Trafodion Installer with this option, you do the following:

1. Copy the trafodion_config_default file.

Example

```
cp trafodion_config_default my_config
```

- 2. Edit the new file using information you collect in the Gather Configuration Information section in the Prepare chapter.
- 3. Run the Trafodion Installer in Automated Setup Mode

Example

```
./trafodion_installer --config_file my_config
```



Your Trafodion Configuration File contains the password for the Trafodion Runtime User and for the Distribution Manager. Therefore, we recommend that you secure the file in a manner that matches the security policies of your organization.



If you are installing Trafodion on a version of Hadoop that has been instrumented with Kerberos, you will be asked for a password associated with a Kerberos administrator.

Example: Creating a Trafodion Configuration File

Using the instructions in Gather Configuration Information in the Prepare chapter, you record the following information.

ADMIN Administrator user name for Apache Ambari or Cloudera admin Manager. ADMIN_PRINCIPAL Kerberos principal for the KDC admin user including the realm. BACKUP_DCS_NODES List of nodes where to start the backup DCS Master	
realm. BACKUP_DCS_NODES List of nodes where to start the backup DCS Master	
'	
components.	
CLOUD_CONFIG Whether you're installing Trafodion on a cloud environment. N	
CLOUD_TYPE What type of cloud environment you're installing Trafodion on.	
CLUSTER_NAME The name of the Hadoop Cluster. Cluster 1	
DB_ROOT_NAME LDAP name used to connect as database root user trafodion	
DCS_BUILD Tar file containing the DCS component.	
DCS_PRIMARY_MASTER_NODE The node where the primary DCS should run.	
DCS_SERVER_PARM Number of concurrent client sessions per node. 8	
ENABLE_HA Whether to run DCS in high-availability (HA) mode. N	
EPEL_RPM Location of EPEL RPM. Specify if you don't have access to the Internet.	
FLOATING_IP IP address if running DCS in HA mode.	
HADOOP_TYPE The type of Hadoop distribution you're installing Trafodion cloudera on.	
HBASE_GROUP Linux group name for the HBASE administrative user. hbase	
HBASE_KEYTAB Kerberos service keytab for HBase admin principal. Default based distribution	d on
HBASE_USER Linux user name for the HBASE administrative user. hbase	
HDFS_KEYTAB Kerberos service keytab for HDFS admin principal. Default based distribution	d on
HDFS_USER Linux user name for the HDFS administrative user. hdfs	
HOME_DIR Root directory under which the trafodion home directory /home should be created.	
INIT_TRAFODION Whether to automatically initialize the Trafodion database. Y	
INTERFACE Interface type used for \$FLOATING_IP.	

ID	Information	Setting
JAVA_HOME	Location of Java 1.7.0_65 or higher (JDK).	/usr/java/jdk1.7.0_67- cloudera
KDC_SERVER	Location of Kerberos server for admin access	
LDAP_CERT	Full path to TLS certificate.	
LDAP_HOSTS	List of nodes where LDAP Identity Store servers are running.	
LDAP_ID	List of LDAP unique identifiers.	
LDAP_LEVEL	LDAP Encryption Level.	
LDAP_PASSWORD	Password for LDAP_USER.	
LDAP_PORT	Port used to communicate with LDAP Identity Store.	
LDAP_SECURITY	Whether to enable LDAP authentication.	N
LDAP_USER	LDAP Search user name.	
LOCAL_WORKDIR	The directory where the Trafodion Installer is located.	/home/centos/trafodion-installer/installer
MANAGEMENT_ENABLED	Whether your installation uses separate management nodes.	N
MANAGEMENT_NODES	The FQDN names of management nodes, if any.	
MAX_LIFETIME	Kerberos ticket lifetime for Trafodion principal	24hours
NODE_LIST	The FQDN names of the nodes where Trafodion will be installed.	trafodion-1 trafodion-2
PASSWORD	Administrator password for Apache Ambari or Cloudera Manager.	admin
RENEW_LIFETIME	Kerberos ticket renewal lifetime for Trafodion principal	7days
REST_BUILD	Tar file containing the REST component.	
SECURE_HADOOP	Indicates whether Hadoop has Kerberos enabled	Based on whether Kerberos is enabled for your Hadoop installation
SQ_ROOT	Target directory for the Trafodion software.	/home/trafodion/apach e-trafodion-1.3.0- incubating-bin
START	Whether to start Trafodion after install/upgrade.	Υ
SUSE_LINUX	Whether your installing Trafodion on SUSE Linux.	false
TRAF_PACKAGE	The location of the Trafodion installation package tar file or core installation tar file.	/home/centos/trafodion- download/apache- trafodion-1.3.0- incubating-bin.tar.gz
TRAF_KEYTAB	Kerberos keytab for trafodion principal.	Default keytab based on distribution
TRAF_KEYTAB_DIR	Location of Kerberos keytab for the trafodion principal.	Default location based on distribution
TRAF_USER	The Trafodion runtime user ID. Must be trafodion in this release.	trafodion
TRAF_USER_PASSWORD	The password used for the ${\tt trafodion:trafodion:user}$ ID.	traf123
URL	FQDN and port for the Distribution Manager's REST API.	trafodion- 1.apache.org:7180

Next, you edit my_config to contain the following:

```
#!/bin/bash
# @@@ START COPYRIGHT @@@
# Licensed to the Apache Software Foundation (ASF) under one
# or more contributor license agreements. See the NOTICE file
# distributed with this work for additional information
# regarding copyright ownership. The ASF licenses this file
# to you under the Apache License, Version 2.0 (the
# "License"); you may not use this file except in compliance
# with the License. You may obtain a copy of the License at
#
   http://www.apache.org/licenses/LICENSE-2.0
# Unless required by applicable law or agreed to in writing,
# software distributed under the License is distributed on an
# "AS IS" BASIS, WITHOUT WARRANTIES OR CONDITIONS OF ANY
# KIND, either express or implied. See the License for the
# specific language governing permissions and limitations
# under the License.
# @@@ END COPYRIGHT @@@
# Trafodion Configuration File
# This file contains default values for the installer.
# Users can also edit this file and provide values for all parameters
# and then specify this file on the run line of trafodion_install.
# Example:
# ./trafodion_install --config_file <Trafodion-config-file>
# WARNING: This mode is for advanced users!
#Must be set to 'true' if on a SUSE linux system. If on another type of system
#this must be set to false.
export SUSE_LINUX="false"
# The working directory where Trafodion installer untars files, etc.
# do not change this unless you really know what you are doing
export TRAF_WORKDIR="/usr/lib/trafodion"
# This is the directory where the installer scripts were untarred to
export LOCAL_WORKDIR="/home/centos/trafodion-installer/installer"
# The maximum number of dcs servers, i.e. client connections
export DCS_SERVERS_PARM="8"
```

```
# "true" if this is an upgrade
export UPGRADE_TRAF="false"
# Trafodion userid, This is the userid the Trafodion instance will run under
export TRAF_USER="trafodion"
# Trafodion userid's password
export TRAF_USER_PASSWORD="traf123"
# a blank separated list of nodes in your cluster
# node names should include full domain names
#This can not be left blank!
export NODE_LIST="trafodion-1 trafodion-2"
# count of nodes in node list
export node_count="2"
# another list of the same nodes in NODE LIST but specified in a pdsh usable format
# i.e. "-w centos-cdh[1-6]" or "-w node1 -w node2 -w node3"
export MY_NODES="-w trafodion-[1-2]"
# the directory prefix for the trafodion userid's $HOME directory
# i.e. /opt/home, not /opt/home/trafodion
export HOME_DIR="/home"
#JAVA HOME must be a JDK. Must include FULL Path. Must be 1.7.0_65 or higher.
export JAVA_HOME="/usr/java/jdk1.7.0_67-cloudera"
# If your machine doesn't have external internet access then you must
# specify the location of the EPEL rpm, otherwise leave blank and it
# will be installed from the internet
export EPEL RPM=""
# full path of the Trafodion package tar file
export TRAF_PACKAGE="/home/centos/trafodion-download/apache-trafodion-1.3.0-incubating-
bin.tar.qz"
# if TRAF_PACKAGE wasn't specified then these two values must be specified
# TRAF BUILD is the trafodion server tar file
# DCS_BUILD is the DCS tar file
# REST_BUILD is the REST tar file
export TRAF_BUILD=""
export DCS_BUILD=""
export REST_BUILD=""
# Either "cloudera" or "hortonworks" (all lowercase)
export HADOOP_TYPE="cloudera"
# The URL for Cloudera/Hortonworks REST API (i.e. node1.host.com:8080)
export URL="trafodion-1.apache.org:7180"
# Cloudera/Hortonworks UI admin's userid and password
export ADMIN="admin"
export PASSWORD="admin"
```

```
# hadoop cluster name
export CLUSTER_NAME=""
# the Hadoop HDFS userid
export HDFS_USER="hdfs"
# the Hadoop HBase userid and group
export HBASE_USER="hbase"
export HBASE_GROUP="hbase"
# The hadoop HBase service name
export HBASE="hbase"
# full path of where to install Trafodion to
# Example is used below. If $HOME_DIR or $TRAF_USER have been changed
# then this will need to be changed.
# On an upgrade, it is recommend to choose a different directory.
# First time install : /home/trafodion/traf
# On Upgrade: /home/trafodion/traf_<date>
# By doing this the previous version will remain and allow for an easier rollback.
export SQ_ROOT="/home/trafodion/apache-trafodion-1.3.0-incubating-bin"
# Start Trafodion after install completes
export START="Y"
# initialize trafodion after starting
export INIT_TRAFODION="Y"
# full path to the sqconfig file
# Default is to leave as is and this file will be created.
export SQCONFIG=""
#----- security configuration information -----
#Enter in Kerberos details if Kerberos is enabled on your cluster
#Indicate Kerberos is enabled
export SECURE_HADOOP="N"
#Location of Kerberos server for admin access
export KDC_SERVER=""
#Kerberos Admin principal used to create Trafodion principals and keytabs
#Please include realm, for example: trafadmin/admin@MYREALM.COM
export ADMIN_PRINCIPAL=""
#Keytab for HBase admin user, used to grant Trafodion user CRWE privilege
export HBASE KEYTAB=""
#Keytab for HDFS admin user, used to create data directories for Trafodion
export HDFS_KEYTAB=""
#Kerberos ticket defaults for the Trafodion user
export MAX_LIFETIME="24hours"
```

```
export RENEW_LIFETIME="7days"
#Trafodion keytab information
export TRAF_KEYTAB=""
export TRAF_KEYTAB_DIR=""
#Enter in LDAP configuration information
#Turn on authentication - MUST have existing LDAP configured.
export LDAP_SECURITY="Y"
#Name of LDAP Config file
export LDAP_AUTH_FILE="traf_authentication_config_`hostname -s`"
#LDAP name to map to database user DB__ROOT
DB_ROOT_NAME="trafodion"
#---- end security configuration
export CONFIG_COMPLETE="true"
```

Once completed, run the Trafodion Installer with the --config_file option.

Refer to the following sections for examples:

- Automated Install
- Automated Upgrade

3.6. Trafodion Provisioning Directories

Trafodion stores its provisioning information in the following directories on each node in the cluster:

- /etc/trafodion: Configuration information.
- /usr/lib/trafodion: Copies of the files required by the installer.

Chapter 4. Requirements

Trafodion requires an x86 version of Linux.

The current release of Trafodion has been tested with:

- 64-bit Red Hat Enterprise Linux (RHEL) or CentOS 6.5, 6.6, and 6.7
- Cloudera CDH 5.4
- Hortonworks HDP 2.3

Other OS releases may work, too. The Trafodion project is currently working on better support for more distribution and non-distribution versions of Hadoop.

4.1. General Cluster and OS Requirements and Recommendations

64-bit x86 instruction set running a Linux distribution is required. Further, Trafodion assumes an environment based on the requirements of the tested Hadoop distributions/services.

4.1.1. Hardware Requirements and Recommendations

Single-Node Cluster

It is possible to run Trafodion on a single-node sandbox environment. Typically, any sandbox running a Hadoop distribution can be used. A typical single-node configuration uses 4-8 cores with 16 GB of memory, and 20 GB free disk space.

Multi-Node Cluster

For multi-node end-user clusters, your typical HBase environment should suffice for Trafodion. Typically, memory configuration range between 64-128 GB per node with minimum requirement of 16 GB. The cluster size can span from 1 to *n* nodes; a minimum of two nodes is recommended. A minimum of two cores is required regardless of whether you're deploying Trafodion on a bare-metal or virtual environment.

Recommended configurations:

Attribute	Guidance
Processors per Node	Small: 2 coresMedium: 4 coresLarge: 8+ cores
Memory per Node	Small: 16 GBMedium: 64 GBLarge: 128 GB
Concurrency:Nodes	Two Small Nodes: Four concurrent queriesTwo Medium Nodes: 64 concurrent queriesTwo Large Nodes: 256 concurrent queries

4.1.2. OS Requirements and Recommendations

Please verify these requirements on each node you will install Trafodion on:

Function	Requirement	Verification Guidance
Linux	64-bit version of Red Hat 6.5 or later, or SUSE SLES 11.3 or later.	
sshd	The ssh daemon is running on each node in the cluster.	•ps aux grep sshd •sudo netstat -plant grep :22
ntpd	The ntp daemon is running and synchronizing time on each node in the cluster.	•ps aux grep ntp •ntpq -p
FQDN	 /etc/hosts is set up for fully-qualified node names (FQDN). /etc/resolv.conf is configured to use a name server. 	 hostnamefqdn shows the fully-qualified node name, if any. The fully-qualified node name is part of the /etc/hosts file. host -T <fqdn> (responds if using a DNS server, times out otherwise)</fqdn> Simply ssh among nodes using ssh <fqdn>.</fqdn>
Port Availability	The Linux Kernel Firewall (iptables) has either been disabled or ports required by Trafodion have been opened.	• lsmod grep ip_tables checks whether iptables is loaded. If not, no further checking is needed. • sudo iptables -nL grep <port> checks the configuration of a port. An empty response indicates no rule for the port, which often means the port is not open.</port>
passwordless ssh	The user name used to provision Trafodion must have passwordless ssh access to all nodes.	ssh to the nodes, ensure that no password prompt appears.
sudo privileges	The user name used to provision Trafodion must sudo access to a number of root functions .	sudo echo "test" on each node.
bash	Available for shell-script execution.	bashversion
java	Available to run the Trafodion software. Same version as HBase is using.	javaversion
perl	Available for script execution.	perlversion
python	Available for script execution.	pythonversion

Function	Requirement	Verification Guidance
yum	Available for installs, updates, and removal of software packages.	yumversion
rpm	Available for installs, updates, and removal of software packages.	rpmversion
scp	Available to copy files among nodes in the cluster.	scphelp
curl	Available to transfer data with URL syntax.	curlversion
wget	Available to download files from the Web.	wgetversion
pdsh	Available to run shell commands in parallel.	pdsh -V
pdcp	Available to copy files among nodes in parallel. part of the pdsh package.	pdcp -V

4.1.3. IP Ports

The following table lists the default ports used by the different Trafodion components plus the configuration file and configuration attribute associated with each port setting.

Default Port	Configuration File	Configuration Entry	Required	Ran ge	Protocol	Comment
4200	rest-site.xml	trafodion.rest.port	Yes	1	REST	Trafodion REST Server.
4201	rest-site.xml	trafodion.rest.https.port	Yes	1	HTTPS	Trafodion REST Server (HTTPS).
23400	dcs-site.xml	dcs.master.port	Yes	n	binary	Start of Trafodion DCS port range. (37800 for Trafodion 1.1)
24400	dcs-site.xml	dcs.master.info.port	Yes	1	HTTP	DCS master web GUI. (40010 for Trafodion 1.1)
24410	dcs-site.xml	dcs.server.info.port	Yes	n	HTTP	Start of range for DCS server web GUIs. (40020 for Trafodion 1.1)
50030	mapred-site.xml	mapred.job.tracker.http .address	No	1	HTTP	MapReduce Job Tracker web GUI.
50070	hdfs-site.xml	dfs.http.address	No	1	HTTP	HDFS Name Node web GUI.
50075	hdfs-site.xml	dfs.datanode.http.addr ess	No	1	HTTP	HDFS Data Node web GUI.
50090	hdfs-site.xml	dfs.secondary.http.addr ess	No	1	HTTP	HDFS Secondary Name Node web GUI.
60010	hbase-site.xml	hbase.master.info.port	No	1	HTTP	HBase Master web GUI.
60030	hbase-site.xml	hbase.regionserver.info .port	No	1	HTTP	HBase Region Server web GUI.

There are two port ranges used by Trafodion.

• 23400 is a range, to allow multiple mxosrvr processes on each node. Allow a range of a few ports, enough to cover all the servers per node that are listed in the "servers" file in the DCS configuration directory.

24410 is a range as well, enough to cover the DCS servers per node, usually 1 or 2.

On top of the ports identified above, you also need the ports required by your Hadoop distribution. For example:

- Cloudera Ports
- Hortonworks Ports

If you have Kerberos or LDAP enabled, then ports required by these products need to be opened as well.

Although not all the ports will be used on every node of the cluster, you need to open most of them for all the nodes in the cluster that have Trafodion, HBase, or HDFS servers on them.

4.2. Prerequisite Software

4.2.1. Hadoop Software

Trafodion runs as an add-on service on Hadoop distributions. The following Hadoop services and their dependencies must be installed and running on the cluster where you intend to install Trafodion:

- Hadoop Distributed File System (HDFS)
- YARN with MapReduce version 2
- ZooKeeper
- HBase
- Hive
- Apache Ambari (Hortonworks) or Cloudera Manager (Cloudera) with associated embedded databases.

The following distributions have been tested with Trafodion.¹

Distribution	Version	HBase Version	Installation Documentation
Cloudera Distribution Including Apache Hadoop (CDH)	5.4	1.0	CHD 5.4 Installation
Hortonworks Data Platform (HDP)	2.3	1.1	HDP 2.3 Installation

1. Future releases of Trafodion will move away from distribution-specific integration. Instead, Trafodion will be tested with specific version of the Hadoop, HDFS, HBase, and other services/products only.

2. When possible, install using **parcels** to simply the installation process.



Trafodion does not yet support installation on a non-distribution version of Hadoop; that is, Hadoop downloaded from the Apache web site. This restriction will be lifted in a later release of Trafodion.

4.2.2. Software Packages

In addition to the software packages required to run different Hadoop services listed above (for example, Java), Trafodion requires supplementary software to be installed on the cluster before it is installed. These are Linux tools that are not typically packaged as part of the core Linux distribution.



For RedHat/CentOS, the Trafodion Installer automatically attempts get a subset of these packages over the Internet. If the cluster's access to the Internet is disabled, then you need to manually download the packages and make them available for installation. You need to build and install log4c++ manually.

Package	Usage	Installation
EPEL	Add-on packages to completed the Linux distribution.	Download Fedora RPM
pdsh	Parallelize shell commands during install and Trafodion runtime utilities.	yum install pdsh
log4cxx	Message logging.	Manual process ¹
sqlite	Internal configuration information managed by the Trafodion Foundation component.	yum install sqlite
expect	Not used?	yum install expect
perl-DBD-SQLite	Allows Perl scripts to connect to SQLite.	yum install perl-DBD-SQLite
perl-Params-Validate	Validates method/function parameters in Perl scripts.	yum install perl-Params-Validate
perl-Time-HiRes	High resolution alarm, sleep, gettimeofday, interval timers in Perl scripts.	yum install perl-Time-HiRes
protobuf	Data serialization.	yum install protobuf
xerces-c	C++ XML parsing.	yum install xerces-c
gzip	Data compress/decompress.	yum install gzip
rpm-build ²	Build binary and source software packages.	yum install rpm-build
apr-devel ²	Support files used to build applications using the APR library.	yum install apr-devel
apr-util-devel ²	Support files used to build applications using the APR utility library.	yum install apr-util-devel
doxygen ²	Generate documentation from annotated C++ sources.	yum install doxygen
gcc ²	GNU Compiler Collection	yum install gcc
gcc_c++2	GNU C++ compiler.	yum install gcc_c++

1. log4c++ was recently withdrawn from public repositories. Therefore, you will need to build the log4c++ RPM on your

system and then install the RPM using the procedure described in log4c++ Installation.

2. Software package required to build log4c++. Not required otherwise. These packages are not installed by the Trafodion Installer in this release.

The Trafodion Installer requires Internet access to install the required software packages.

4.3. Trafodion User IDs and Their Privileges

4.3.1. Trafodion Runtime User

The trafodion: trafodion user ID is created as part of the installation process. The default password is: traf123.

Trafodion requires that either HDFS ACL support or Kerberos is enabled. The Trafodion Installer will enable HDFS ACL and Kerberos support. Refer to Kerberos for more information about the requirements and usage of Kerberos in Trafodion. Refer to Apache HBase™ Reference Guide for security in HBase.

Also, Trafodion requires sudo access to ip and arping so that floating or elastic IP addresses can be moved from one node to another in case of node failures.



Do not create the trafodion: trafodion user ID in advance. The Trafodion Installer uses the presence of this user ID to determine whether you're doing an installation or upgrade.

4.3.2. Trafodion Provisioning User

Typically, the Trafodion Installer is used for Trafodion installations. It requires access to the user IDs documented below.

Linux Installation User

The user ID that performs the Trafodion installation steps. Typically, this User ID runs the Trafodion Installer.

Requirements:

- User name or group cannot be trafodion.
- · Passwordless ssh access to all nodes in the cluster.
- Internet access to download software packages.
- requiretty must be disabled in /etc/sudoers.

- sudo1 access to:
 - · Download and install software packages.
 - Modify /etc/sudoers.d (allow the trafodion user to modify floating IP: ip and arping).
 - Create the trafodion user ID and group.
 - Install Trafodion software into the HBase environment.
 - Run Java version command on each node in the cluster.
 - Run Hadoop version command on each node in the cluster.
 - Run HBase version command on each node in the cluster.
 - · Create directories and files in:
 - /etc
 - /usr/lib
 - /var/log
 - Invoke su to execute commands as other users; for example, trafodion.
 - Edit sysctl.conf and activate changes using sysctl -p:
 - Modify kernel limits.
 - · Reserve IP ports.

1 sudo is required in the current release of Trafodion. This restriction may be relaxed in later releases. Alternative mechanisms for privileged access (such as running as root or sudo alternative commands) are not supported.

Distribution Manager User

A user ID that can change the configuration using Apache Ambari or Cloudera Manager. The Trafodion Installer makes REST request to perform configuration and control functions to the distribution manager using this user ID.

Requirements:

- Administrator user name and password.
- URL to Distribution Manager's REST API.

HDFS Administrator User

The HDFS super user. Required to create directories and change security settings, as needed. The Trafodion Installer uses su to run commands under this user ID.

Requirements:

- HDFS Administrator user name.
- Write access to home directory on the node where the Distribution Manager is running.
- For Kerberos enabled installations, location of the keytab for the HDFS service principal.

HBase Administrator User

The HBase super user. Required to change directory ownership in HDFS. For Kerberos enabled installations, the HBase super user is needed to grant the trafodion user create, read, write, and execute privileges.

Requirements:

- · HBase Administrator user name and group.
- Read access to hbase-site.xml.
- For Kerberos enabled installations, location of the keytab for the HBase service principal.

Kerberos Administrator User

The Kerberos adminstrator. Required to create Trafodion principals and keytabs on a cluster where Kerberos is enabled.

Requirements:

- Kerberos Administrator admin name including the realm.
- · Kerberos Administrator password

4.4. Required Configuration Changes

Trafodion requires changes to a number of different areas of your system configuration: operating system, HDFS, and HBase.



These changes are performed by the Trafodion Installer, if used.

4.4.1. Operating System Changes

/etc/security/limits.d/trafodion.conf on each node in the cluster must contain the following settings:

```
# Trafodion settings
trafodion soft core unlimited
trafodion hard core unlimited
trafodion soft memlock unlimited
trafodion hard memlock unlimited
trafodion soft nofile 32768
trafodion hard nofile 65536
trafodion soft nproc 100000
trafodion hard nproc 100000
```

4.4.2. ZooKeeper Changes



These changes require a restart of ZooKeeper on all nodes in the cluster.

Trafodion requires the following changes to zoo.cfg:

Setting	New Value	Purpose
maxClientCnxns	0	Tell ZooKeeper to impose no limit to the number of connections to enable better Trafodion concurrency.



If Kerberos is enabled, it is not possible to secure the Trafodion data in ZooKeeper at this time.

4.4.3. HDFS Changes



These changes require a restart of HDFS on all nodes in the cluster.

Trafodion requires the following changes to the HDFS environment:

Action	Purpose
Create /hbase-staging directory.Change owner to HBase Administrator.	
 Create /bulkload directory. Change owner to trafodion. 	Used to stage data when processing the Trafodion LOAD INTO table statement and as a temporary directory to create links to actual HFile for snapshot scanning.

- Create /lobs directory.
- Change owner to trafodion.
- Create /apps/hbase/data/archive1.
- Change owner to: hbase: hbase (Cloudera) or hbase: hdfs (Hortonworks)
- Give the trafodion user RWX access to /apps/hbase/data/archive
- Set default user of /apps/hbase/data/archive to trafodion
- Recursively change setafcl of /apps/hbase/data/archive to **RWX**
- 1. These steps are performed after HDFS ACLs have been enabled.

The following changes are required in hdfs-site.xml:

Setting	New Value	Purpose
dfs.namenode.acls.enabled	true	Enable HDFS POSIX Access
		Control Lists (ACLs).

4.4.4. HBase Changes



These changes require a restart of ZooKeeper and HBase on all nodes in the cluster.

Trafodion requires that the following changes to the HBase environment:

Action	Affected Directories	Purpose
Install/replace Trafodion's version of hbase-trx	 /usr/lib/hbase/lib/ /usr/share/cmf/lib/plugins/ (Cloudera) /usr/hdp/current/hbase-regionserver/lib/ (Hortonworks) 	Trafodion transaction management relies on an enhanced version of hbase-trx.
Install/Replace Trafodion utility jar file.	 /usr/lib/hbase/lib/ /usr/share/cmf/lib/plugins/ (Cloudera) /usr/hdp/current/hbase-regionserver/lib (Hortonworks) 	TODO: Add purpose here.
For Kerberos enabled clusters, grant trafodion user privileges	not applicable	privileges: create, read, write, and execute access

The following changes are required in hbase-site.xml. Please refer to the Apache HBase™ Reference Guide for additional descriptions of these settings.

Setting	New Value	Purpose
hbase.master. distributed.log.splitting	false	Do not use the HBase Split Log Manager. Instead, the HMaster controls all log-splitting activities.
hbase.coprocessor. region.classes	org.apache.hadoop. hbase.coprocessor. transactional.TrxRegionObserver, org.apache.hadoop. hbase.coprocessor. transactional.TrxRegionEndpoint, org.apache.hadoop. hbase.coprocessor. AggregateImplementation	Install Trafodion coprocessor classes.
hbase.hregion.impl	org.apache.hadoop. hbase.regionserver. transactional.TransactionalRegion	Trafodion needs to be able to read the Write Ahead Log from a coprocessor using the getScanner method. This method is protected in standard HBase. This change overloads the getScanner method to be public thereby allowing coprocessor code to use it.
hbase.regionserver. region.split.policy	org.apache.hadoop. hbase.regionserver. ConstantSizeRegionSplitPolicy	Tell HBase to use the ConstantSizeRegionSplitPolicy for region splitting. This setting causes region splitting to occur only when the maximum file size is reached.
hbase.snapshot. enabled	true	Enable the HBase Snapshot feature. Used for Trafodion backup and restore.

Setting	New Value	Purpose
hbase.bulkload. staging.dir	hbase-staging	Use /hbase-staging as the bulk load staging directory.
hbase.regionserver.region. transactional.tlog	true	The HBase Regions requests that the Transaction Manager re-drives in-doubt transactions.
hbase.snapshot. master.timeoutMillis	600000	HMaster timeout when waiting for RegionServers involved in the snapshot operation.
hbase.snapshot. region.timeout	600000	RegionServer timeout when waiting for snapshot to be created.
hbase.client. scanner.timeout.period	600000	Time limit to perform a scan request.
hbase.regionserver. lease.period	600000	Clients must report within this time limit or they are considered dead by HBase.
hbase.namenode. java.heapsize ¹	1073741824 (1GB)	Java Heap Size for the HDFS NameNode.
hbase.secondary.namenode. java.heapsize ¹	1073741824 (1GB)	Java Heap Size for the HDFS Secondary NameNode.

^{1.} Applies to Cloudera distributions only.

4.5. Recommended Configuration Changes

The following configuration changes are recommended but not required.



The Trafodion Installer does **not** make these changes.

4.5.1. Recommended Security Changes

The trafodion user ID should not be given other sudo privileges than what's specified in this manual. Also, we recommend that this user ID is locked (sudo passwd -1 trafodion) once the installation/upgrade activity has been completed. Users that need issue commands as the trafodion ID should do so using sudo; for example, sudo -u trafodion -i.

4.5.2. Recommended HDFS Configuration Changes

These settings are configured in the hadoop-env.sh file.

Property	Recommended Setting	Guidance
DataNode Java Heap Size	2 GB	Use this setting for a large configuration.
NameNode Java Heap Size	2 GB	Use this setting for a large configuration.

Property	Recommended Setting	Guidance
Secondary NameNode Java Heap Size	2 GB	Use this setting for a large configuration.

4.5.3. Recommended HBase Configuration Changes

Configuration Property	Recommended Setting	Guidance
hbase.rpc.timeout	10 minutes	This setting depends on the tables' size. Sixty (60) seconds is the default. Increase this value for big tables. Make it the same value as hbase.client.scanner.timeout.period. We have found that increasing the setting to six-hundred (600) seconds will prevent many of the timeout-related errors we encountered, such as OutOfOrderNextException errors.
hbase.client.scanner.timeout.period	10 minutes	Similar to the hbase.rpc.timeout setting. Sixty (60) seconds is the default. Depending on the size of a user table, we have experienced timeout failures on count(*) and update statistics commands from this setting. The underlying issue is the length of the execution of the coprocessor within HBase. NOTE: HBase uses the smaller of hbase.rpc.timeout and hbase.client.scanner.timeout.period to calculate the scanner timeout.
hbase.snapshot.master.timeoutMilli s and hbase.snapshot.region.timeout	10 minutes	HBase's default setting is 60000 milliseconds. If you experience timeout issues with HBase snapshots when you use the Trafodion Bulk Loader or other statements, you can set the value for these two HBase properties to 10 minutes (600,000 milliseconds).
hbase.hregion.max.filesize	107374182400 bytes	HBase's default setting is 10737418240 (10 GB). We have increased the setting to 107374182400 (100 GB), which reduces the number of HStoreFiles per table and appears to reduce disruptions to active transactions from region splitting.
hbase.hstore.blockingStoreFiles	10	http://gbif.blogspot.com/2012/07/optimizing-writes-in- hbase.html
hbase.regionserver.handler.count	<num></num>	This setting should match the number of concurrent sessions (mxosrvr). The default is 10.

Chapter 5. Prepare

You need to prepare your Hadoop environment before installing Trafodion.

- Install Optional Workstation Software
- Configure Installation User ID
- 3. Disable requiretty
- 4. Verify OS Requirements and Recommendations
- Configure Kerberos
- 6. Configure LDAP Identity Store
- 7. Gather Configuration Information
- 8. Install Required Software Packages
- Perform Recipe-Based Provisioning Tasks

5.1. Install Optional Workstation Software

If you are using a Windows workstation, then the following optional software helps installation process. We recommended that you pre-install the software before continuing with the Trafodion installation:

- putty and puttygen (download from PuTTY web site)
- VNC client (download from RealVNC web site)
- Firefox or Chrome browser
- SFTP client to transfer files from your workstation to the Linux server: WinSCP or FileZilla

5.2. Configure Installation User ID

Trafodion installation requires a user ID with these attributes:

- sudo access per the requirements documented in Linux Installation User.
- passwordless ssh to all nodes on the cluster where Trafodion will be installed.



You may need to request permission from your cluster-management team to obtain this type of access.

The following example shows how to set up your user ID to have "passwordless ssh" abilities.

Do the following on the Provision Master Node:

```
echo -e 'y\n' | ssh-keygen -t rsa -N "" -f $HOME/.ssh/id_rsa cat $HOME/.ssh/id_rsa.pub >> $HOME/.ssh/authorized_keys echo localhost $(cat /etc/ssh/ssh_host_rsa_key.pub) >> $HOME/.ssh/known_hosts echo "NoHostAuthenticationForLocalhost=yes" >> $HOME/.ssh/config chmod 600 $HOME/.ssh/config chmod 600 $HOME/.ssh/authorized_keys chmod 700 $HOME/.ssh/
```

After running these commands, do the following:

- If necessary, create the \$HOME/.ssh directory on the other nodes in your cluster and secure it private to yourself (chmod 700).
- If necessary, create the \$HOME/.ssh/authorized_keys file on the other nodes in your cluster. Secure it with chmod 600 \$HOME/.ssh/authorized_keys.
- Copy the content of the \$HOME/.ssh/id_rsa.pub file on the Provisioning Master Node and append the to the \$HOME/.ssh/authorized_keys file on the other nodes in your cluster.
- ssh to the other nodes in the cluster. Answer y to the prompt asking you whether to continue the connection. This adds the node to the \$HOME/.ssh/known_hosts file completing the passwordless ssh setup.

5.3. Disable requiretty

You need to disable requiretty in /etc/sudoers on all nodes in the cluster to ensure that sudo commands can be run from inside the installation scripts.

Comment out the Defaults requiretty setting in the /etc/sudoers file to ensure that the requiretty option is NOT being used.

5.4. Verify OS Requirements and Recommendations

Please ensure that the OS Requirements and Recommendations are met for each node in the cluster where you intend to install Trafodion.

5.5. Configure Kerberos

If your Hadoop installation has enabled Kerberos, then Trafodion needs to have Kerberos enabled. If not, then Trafodion will not run. If you plan to enable Kerberos in Trafodion, then you need to have access to a KDC (Kerberos Key Distribution Center) and administration credentials so you can create the necessary Trafodion principals and keytabs.

If you wish to manually set up and activate Kerberos principals and keytabs, then refer to the section on Kerberos.

5.6. Configure LDAP Identity Store

If you plan to enable security features in Trafodion, then you need to have an LDAP identity store available to perform authentication. The Trafodion Installer prompts you to set up an authentication configuration file that points to an LDAP server (or servers), which enables security (that is, authentication and authorization) in the Trafodion database.

If you wish to manually set up the authentication configuration file and enable security, then refer to the section on LDAP.

5.7. Gather Configuration Information

You need to gather/decide information about your environment to aid installation Trafodion, both for the Trafodion Installer and for recipe-based provisioning. (Listed in alphabetical order to make it easier to find information when referenced in the install and upgrade instructions.)

ID¹	Information	Default	Notes
ADMIN	Administrator user name for Apache Ambari or Cloudera Manager.	admin	A user that can change configuration and restart services via the distribution manager's REST API.
ADMIN_PRINCIPAL ²	Kerberos admin principal to manage principals and keytabs	None	Required if Kerberos is enabled.
BACKUP_DCS_NODES	List of nodes where to start the backup DCS Master components.	None	Blank separated FQDN list. Not needed if \$ENABLE_HA = N.
CLOUD_CONFIG	Whether you're installing Trafodion on a cloud environment.	N	N = bare-metal or VM installation.
CLOUD_TYPE	What type of cloud environment you're installing Trafodion on.	None	{ AWS OpenStack Other }. Not applicable for bare-metal or VM installation.
CLUSTER_NAME	The name of the Hadoop Cluster.	None	From Apache Ambari or Cloudera Manager.
DB_ROOT_NAME ²	LDAP name used to connect as database root user	trafodion	Required when LDAP is enabled.
DCS_BUILD	Tar file containing the DCS component.	None	Not needed if using a Trafodion package installation tar file.

ID ¹	Information	Default	Notes
DCS_PRIMARY_MASTER_ NODE	The node where the primary DCS should run.	None	The DCS Master handles JDBC and ODBC connection requests.
DCS_SERVER_PARM	Number of concurrent client sessions per node.	16	This number specifies the concurrent sessions per node to be supported. Each session could require up to 1GB of physical memory. The number can be changed post-installation. For more information, refer to the Trafodion Client Installation Guide.
ENABLE_HA	Whether to run DCS in high-availability (HA) mode.	N	You need the floating IP address, the interface, and the backup nodes for DCS Master if enabling this feature.
EPEL_RPM	Location of EPEL RPM.	None	Specify if you don't have access to the Internet. Downloaded automatically by the Trafodion Installer.
FLOATING_IP	IP address if running DCS in HA mode.	None	Not needed if \$ENABLE_HA = N. An FQDN name or IP address.
HADOOP_TYPE	The type of Hadoop distribution you're installing Trafodion on.	None	Lowercase. cloudera or hadoop.
HBASE_GROUP	Linux group name for the HBASE administrative user.	hbase	Required in order to provide access to select HDFS directories to this user ID.
HBASE_KEYTAB ²	HBase credentials used to grant Trafodion CRWE privileges	based on distribution	Required if Kerberos is enabled.
HBASE_USER	Linux user name for the HBASE administrative user.	hbase	Required in order to provide access to select HDFS directories to this user ID.
HDFS_KEYTAB ²	HDFS credentials used to set privileges on HDFS directories	based on distribution	Required if Kerberos is enabled.
HDFS_USER	Linux user name for the HDFS administrative user.	hdfs	The Trafodion Installer uses sudo su to make HDFS configuration changes under this user.
HOME_DIR	Root directory under which the trafodion home	/home	Example
	directory should be created.		If the home directory of the trafodion user is /opt/home/trafodion, then specify the root directory as /opt/home.
INIT_TRAFODION	Whether to automatically initialize the Trafodion database.	N	Does not apply to Recipe-Based Provisioning. Applies if \$START=Y only.
INTERFACE	Interface type used for \$FLOATING_IP.	None	Not needed if \$ENABLE_HA = N.
JAVA_HOME	Location of Java 1.7.0_65 or higher (JDK).	\$JAVA_HOME setting	Fully qualified path of the JDK. For example: /usr/java/jdk1.7.0_67-cloudera
KDC_SERVER ²	Location of host where Kerberos server exists	None	Required if Kerberos enabled.
LDAP_CERT ²	Full path to TLS certificate.	None	Required if \$LDAP_LEVEL = 1 or 2.

ID¹	Information	Default	Notes
LDAP_HOSTS ²	List of nodes where LDAP Identity Store servers are running.	None	Blank separated. FQDN format.
LDAP_ID ²	List of LDAP unique identifiers.	None	Blank separated.
LDAP_LEVEL ²	LDAP Encryption Level.	0	0: Encryption not used, 1: SSL, 2: TLS
LDAP_PASSWORD ²	Password for LDAP_USER.	None	If LDAP_USER is required only.
LDAP_PORT ²	Port used to communicate with LDAP Identity Store.	None	Examples: 389 for no encryption or TLS, 636 for SSL.
LDAP_SECURITY ²	Whether to enable simple LDAP authentication.	N	If Y, then you need to provide LDAP_HOSTS.
LDAP_USER ²	LDAP Search user name.	None	Required if you need additional LDAP functionally such as LDAPSearch. If so, must provide LDAP_PASSWORD, too.
LOCAL_WORKDIR	The directory where the Trafodion Installer is located.	None	Full path, no environmental variables.
MANAGEMENT_ENABLED	Whether your installation uses separate management nodes.	N	Y if using separate management nodes for Apache Ambari or Cloudera Manager.
MANAGEMENT_NODES	The FQDN names of management nodes, if any.	None	Provide a blank-separated list of node names.
MAX_LIFETIME ²	Kerberos ticket lifetime for Trafodion principal	24hours	Can be specified when Kerberos is enabled.
NODE_LIST	The FQDN names of the nodes where Trafodion will be installed.	None	Provide a blank-separated list of node names. The Trafodion Provisioning ID must have passwordless and sudo access to these nodes.
PASSWORD	Administrator password for Apache Ambari or Cloudera Manager.	admin	A user that can change configuration and restart services via the distribution manager's REST API.
RENEW_LIFETIME ²	Number times Kerberos ticket is for the Trafodion principal	7days	Can be specified when Kerberos is enabled.
REST_BUILD	Tar file containing the REST component.	None	Not needed if using a Trafodion package installation tar file.
SECURE_HADOOP ²	Indicates whether Hadoop has enabled Kerberos	Y only if Kerberos enabled	Based on whether Kerberos is enabled for your Hadoop installation
SQ_ROOT	Target directory for the Trafodion software.	\$HOME_DIR/traf odion	Trafodion is installed in this directory on all nodes in \$NODE_LIST.
START	Whether to start Trafodion after install/upgrade.	N	Does not apply to Recipe-Based Provisioning.
SUSE_LINUX	Whether your installing Trafodion on SUSE Linux.	false	Auto-detected by the Trafodion Installer.
TRAF_KEYTAB ²	Name to use when specifying Trafodion keytab	based on distribution	Required if Kerberos is enabled.
TRAF_KEYTAB_DIR ²	Location of {project_name} keytab	based on distribution	Required if Kerberos is enabled.

ID¹	Information	Default	Notes
TRAF_PACKAGE	The location of the Trafodion installation package tar file or core installation tar file.	None	The package file contains the Trafodion server, DCS, and REST software while the core installation file contains the Trafodion server software only. If you're using a core installation file, then you need to record the location of the DCS and REST installation tar files, too. Normally, you perform Trafodion provisioning using a Trafodion package installation tar file.
TRAF_USER	The Trafodion runtime user ID.	trafodion	Must be trafodion in this release.
TRAF_USER_PASSWORD	The password used for the trafodion: trafodion user ID.	traf123	Must be 6-8 characters long.
URL	FQDN and port for the Distribution Manager's REST API.	None	<pre>Include http:// or https:// as applicable. Specify in the form: <ip- address="">:<port> or <node name="">:<port> Example: https://susevm- 1.yourcompany.local:8080</port></node></port></ip-></pre>

- 1. The ID matches the environmental variables used in the Trafodion Installation configuration file. Refer to Trafodion Installer for more information.
- 2. Refer to Enable Security for more information about these security settings.

5.8. Install Required Software Packages

5.8.1. Download and Install Packages

This step is required if you're:

- Installing Trafodion on SUSE.
- Using Recipe-Based Provisioning.
- Can't download the required software packages using the Internet.

If none of these situations exist, then we highly recommend that you use the Trafodion Installer.

You perform this step as a user with root or sudo access.

Install the packages listed in Software Packages above on all nodes in the cluster.

5.9. Download Trafodion Binaries

You download the Trafodion binaries from the Trafodion Download page. Download the following packages:

- Trafodion Installer (if planning to use the Trafodion Installer)
- Trafodion Server



You can download and install the Trafodion Clients once you've installed and activated Trafodion.

Refer to the Trafodion Client Install Guide for instructions.

Example

```
$ mkdir $HOME/trafodion-download
$ cd $HOME/trafodion-download
$ # Download the Trafodion Installer binaries
$ wget http://apache.cs.utah.edu/incubator/trafodion/trafodion-1.3.0.incubating/apache-
trafodion-installer-1.3.0-incubating-bin.tar.gz
Resolving http://apache.cs.utah.edu... 192.168.1.56
Connecting to http://apache.cs.utah.edu|192.168.1.56|:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 68813 (67K) [application/x-gzip]
Saving to: "apache-trafodion-installer-1.3.0-incubating-bin.tar.gz"
124K/s
                                                  in 0.5s
2016-02-14 04:19:42 (124 KB/s) - "apache-trafodion-installer-1.3.0-incubating-
bin.tar.gz" saved [68813/68813]
```

```
$ # Download the Trafodion Server binaries
$ wget http://apache.cs.utah.edu/incubator/trafodion/trafodion-1.3.0.incubating/apache-
trafodion-1.3.0-incubating-bin.tar.gz
Resolving http://apache.cs.utah.edu... 192.168.1.56
Connecting to http://apache.cs.utah.edu|192.168.1.56|:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 214508243 (205M) [application/x-gzip]
Saving to: "apache-trafodion-1.3.0-incubating-bin.tar.gz"
2016-02-14 04:22:14 (3.72 MB/s) - "apache-trafodion-1.3.0-incubating-bin.tar.gz" saved
[214508243/214508243]
$ ls -1
total 209552
-rw-rw-r-- 1 centos centos 214508243 Jan 12 20:10 apache-trafodion-1.3.0-incubating-
bin.tar.gz
-rw-rw-r-- 1 centos centos 68813 Jan 12 20:10 apache-trafodion-installer-1.3.0-
incubating-bin.tar.gz
```

5.10. Preparation for Recipe-Based Provisioning



This step should be skipped if you plan to use the Trafodion Installer

5.10.1. Modify OS Settings

Ensure that the /etc/security/limits.d/trafodion.conf on each node contains the limits settings required by Trafodion. Refer to Operating System Changes for the required settings.

5.10.2. Modify ZooKeeper Configuration

Do the following:

1. Modify the ZooKeeper configuration as follows:

Attribute	Setting	
maxClientCnxns	0	

2. Restart ZooKeeper to activate the new configuration setting.

5.10.3. Modify HDFS Configuration

Do the following:

1. Modify the HDFS configuration as follows:

Attribute	Setting
dfs.namenode.acls.enabled	true

2. Restart HDFS to activate the new configuration setting.

5.10.4. Modify HBase Configuration

Do the following:

1. Modify the HBase configuration as follows:

Attribute	Setting
hbase.master.distributed.log.splitting	false
hbase.coprocessor.region.classes ^b	org.apache.hadoop.hbase.coprocessor.transactional.TrxRegionObser ver,org.apache.hadoop.hbase.coprocessor.transactional.TrxRegionEn dpoint, org.apache.hadoop.hbase.coprocessor.AggregateImplementation
hbase.hregion.impl	org.apache.hadoop.hbase.regionserver.transactional.TransactionalRegion
hbase.regionserver.region.split.policy	org.apache.hadoop.hbase.regionserver.ConstantSizeRegionSplitPolic y
hbase.snapshot.enabled	true
hbase.bulkload.staging.dir	hbase-staging
hbase.regionserver.region.transactional.tlog	true
hbase.snapshot.master.timeoutMillis	600000
hbase.snapshot.region.timeout	600000
hbase.client.scanner.timeout.period	600000
hbase.regionserver.lease.period	600000
hbase.namenode.java.heapsizea	1073741824
hbase.secondary.namenode.java.heapsizea	1073741824

- a) Applies to Cloudera distributions only.
- b) Do not overwrite any coprocessors that may already exist.
- 2. Restart HBase to activate the new configuration setting.

Chapter 6. Install

This chapter describes how to use the Trafodion Installer to install Trafodion. You use the Trafodion Provisioning ID to run the Trafodion Installer.

6.1. Unpack Installer

You should already have downloaded the Trafodion Binaries per the instructions in the Download Trafodion Binaries in the Prepare chapter. If not, please do so now.

The first step in the installation process is to unpack the Trafodion Installer tar file.

Example

```
$ mkdir $HOME/trafodion-installer
$ cd $HOME/trafodion-downloads
$ tar -zxf apache-trafodion-installer-1.3.0-incubating-bin.tar.gz -C $HOME/trafodion-
installer
```

6.2. Guided Install

The Trafodion Installer prompts you for the information you collected in the Gather Configuration Information step in the Prepare chapter.

The following example shows a guided install of Trafodion on a two-node Cloudera Hadoop cluster that does not have Kerberos nor LDAP installed.

Example

1. Run the Trafodion Installer in guided mode.

```
$ cd $HOME/trafodion-installer/installer
$ ./trafodion_install
*******
TRAFODION INSTALLATION START
*******
***INFO: testing sudo access
***INFO: Log file located at /var/log/trafodion/trafodion_install_2016-02-15-04-45-
30.log
***INFO: Config directory: /etc/trafodion
***INFO: Working directory: /usr/lib/trafodion
*******
Trafodion Configuration Setup
**********
***INFO: Please press [Enter] to select defaults.
Enter trafodion password, default is [traf123]: traf123
Enter list of nodes (blank separated), default []: trafodion-1 trafodion-2
Enter Trafodion userid's home directory prefix, default is [/home]: /home
Specify full path to EPEL RPM (including .rpm), default is None:
***INFO: Will attempt to download RPM if EPEL is not installed on all nodes.
Specify location of Java 1.7.0_65 or higher (JDK), default is []:
/usr/java/jdk1.7.0_67-cloudera
Enter full path (including .tar or .tar.gz) of trafodion tar file []:
/home/centos/trafodion-download/apache-trafodion-1.3.0-incubating-bin.tar.gz
Enter Hadoop admin username, default is [admin]:
Enter Hadoop admin password, default is [admin]:
Enter Hadoop external network URL:port (no 'http://' needed), default is []:
trafodion-1.apache.org:7180
Enter HDFS username, default is [hdfs]:
Enter HBase username, default is [hbase]:
Enter HBase group, default is [hbase]:
Enter directory to install trafodion to, default is [/home/trafodion/apache-
trafodion-1.3.0-incubating-bin]:
Total number of client connections per node, default [16]: 8
```

```
Enable simple LDAP security (Y/N), default is N: N
***INFO: Configuration file: /etc/trafodion/trafodion_config
***INFO: Trafodion configuration setup complete
Trafodion Configuration File Check
*********
The authenticity of host 'trafodion-1 (10.1.30.71)' can't be established.
RSA key fingerprint is 83:96:d4:5e:c1:b8:b1:62:8d:c6:78:a7:7f:1f:6a:d7.
Are you sure you want to continue connecting (yes/no)? yes
***INFO: Testing sudo access on node trafodion-1
***INFO: Testing sudo access on node trafodion-2
***INFO: Testing ssh on trafodion-1
***INFO: Testing ssh on trafodion-2
***INFO: Getting list of all cloudera nodes
***INFO: Getting list of all cloudera nodes
***INFO: cloudera list of nodes: trafodion-1 trafodion-2
***INFO: Testing ssh on trafodion-1
***INFO: Testing ssh on trafodion-2
***INFO: Testing sudo access on trafodion-1
***INFO: Testing sudo access on trafodion-2
***DEBUG: trafodionFullName=trafodion_server-1.3.0.tgz
***INFO: Trafodion version = 1.3.0
***DEBUG: HBase's java_exec=/usr/java/jdk1.7.0_67-cloudera/bin/java
**********
TRAFODION SETUP
********
***INFO: Starting Trafodion environment setup (2016-02-15-07-09-58)
=== 2016-02-15-07-09-58 ===
# @@@ START COPYRIGHT @@@
# Licensed to the Apache Software Foundation (ASF) under one
# or more contributor license agreements. See the NOTICE file
# distributed with this work for additional information
# regarding copyright ownership. The ASF licenses this file
# to you under the Apache License, Version 2.0 (the
# "License"); you may not use this file except in compliance
# with the License. You may obtain a copy of the License at
and hold each Contributor harmless for any liability incurred by,
or claims asserted against, such Contributor by reason of your
accepting any such warranty or additional liability.
END OF TERMS AND CONDITIONS
BY TYPING "ACCEPT" YOU AGREE TO THE TERMS OF THIS AGREEMENT: ACCEPT
***INFO: testing sudo access
```

```
***INFO: Checking all nodes in specified node list
trafodion-1
trafodion-2
***INFO: Total number of nodes = 2
***INFO: Starting Trafodion Package Setup (2016-02-15-07-11-09)
***INFO: Installing required packages
***INFO: Log file located in /var/log/trafodion
***INFO: ... pdsh on node trafodion-1
***INFO: ... pdsh on node trafodion-2
***INFO: Checking if log4cxx is installed ...
***INFO: Checking if sqlite is installed ...
***INFO: Checking if expect is installed ...
***INFO: Installing expect on all nodes
***INFO: modifying limits in /usr/lib/trafodion/trafodion.conf on all nodes
***INFO: create Trafodion userid "trafodion"
***INFO: Trafodion userid's (trafodion) home directory: /home/trafodion
***INFO: testing sudo access
Generating public/private rsa key pair.
Created directory '/home/trafodion/.ssh'.
Your identification has been saved in /home/trafodion/.ssh/id_rsa.
Your public key has been saved in /home/trafodion/.ssh/id_rsa.pub.
The key fingerprint is:
4b:b3:60:38:c9:9d:19:f8:cd:b1:c8:cd:2a:6e:4e:d0 trafodion@trafodion-1
The key's randomart image is:
+--[ RSA 2048]----+
   o * X o
  . E X S
   . 0 + +
    0.0
   0..
   00
***INFO: creating .bashrc file
***INFO: Setting up userid trafodion on all other nodes in cluster
***INFO: Creating known_hosts file for all nodes
trafodion-1
trafodion-2
***INFO: trafodion user added successfully
***INFO: Trafodion environment setup completed
***INFO: creating sqconfig file
***INFO: Reserving DCS ports
******
TRAFODION MODS
***INFO: Cloudera installed will run traf_cloudera_mods98
***INFO: Detected JAVA version 1.7
```

```
***INFO: copying hbase-trx-cdh5_3-1.3.0.jar to all nodes
***INFO: Cloudera Manager is on trafodion-1
***INFO: Detected JAVA version 1.7
***INFO: copying hbase-trx-cdh5_3-1.3.0.jar to all nodes
***INFO: Cloudera Manager is on trafodion-1
 % Total % Received % Xferd Average Speed Time
                                               Time Time Current
***INFO: Hadoop restart completed successfully
***INFO: waiting for HDFS to exit safemode
Safe mode is OFF
***INFO: Setting HDFS ACLs for snapshot scan support
***INFO: Trafodion Mods ran successfully.
*******
TRAFODION START
********
/usr/lib/trafodion/installer/..
***INFO: Log file location /var/log/trafodion/trafodion_install_2016-02-15-07-08-
07.log
***INFO: traf_start
**********
**********
/home/trafodion/apache-trafodion-1.3.0-incubating-bin
***INFO: untarring build file /usr/lib/trafodion/apache-trafodion-1.3.0-incubating-
bin/trafodion_server-1.3.0.tgz to /home/trafodion/apache-trafodion-1.3.0-incubating-
hin
****** Generate public/private certificates ******
Cluster Name : Cluster%201
Generating Self Signed Certificate....
******************
Certificate file :server.crt
Private key file :server.key
Certificate/Private key created in directory :/home/trafodion/sqcert
*****************
Updating Authentication Configuration
******************
Creating folders for storing certificates
***INFO: copying /home/trafodion/sqcert directory to all nodes
***INFO: copying install to all nodes
***INFO: starting Trafodion instance
Checking orphan processes.
Removing old mpijob* files from /home/trafodion/apache-trafodion-1.3.0-incubating-
```

```
bin/tmp
Removing old monitor.port* files from /home/trafodion/apache-trafodion-1.3.0-
incubating-bin/tmp
Executing sqipcrm (output to sqipcrm.out)
Starting the SQ Environment (Executing /home/trafodion/apache-trafodion-1.3.0-
incubating-bin/sql/scripts/gomon.cold)
Background SQ Startup job (pid: 7276)
Zookeeper is listening on port 2181
DcsMaster is listening on port 23400
Process
             Configured Actual
                                           Down
_____
              _____
                            _____
DcsMaster
             1
                             1
             2
                            2
DcsServer
mxosrvr
             8
                             8
You can monitor the SQ shell log file: /home/trafodion/apache-trafodion-1.3.0-
incubating-bin/logs/sqmon.log
Startup time 0 hour(s) 1 minute(s) 9 second(s)
Apache Trafodion Conversational Interface 1.3.0
Copyright (c) 2015 Apache Software Foundation
>> initialize trafodion;
--- SQL operation complete.
>>
End of MXCI Session
***INFO: Installation completed successfully.
TRAFODION INSTALLATION COMPLETE
**********
$
```

2. Switch to the Trafodion Runtime User and check the status of Trafodion.

```
$ sudo su - trafodion
$ sqcheck
Checking if processes are up.
Checking attempt: 1; user specified max: 2. Execution time in seconds: 0.
The SQ environment is up!

        Process
        Configured
        Actual
        Down

        -----
        -----
        -----

        DTM
        2
        2

        RMS
        4
        4

        MXOSRVR
        8
        8

$
```

Trafodion is now running on your Hadoop cluster. Please refer to the Activate chapter for basic instructions on how to verify the Trafodion management and how to perform basic management operations.

6.3. Automated Install

The --config_file option runs the Trafodion in Automated Setup mode. Refer to Trafodion Installer in the Introduction chapter for instructions of how you edit your configuration file.

Edit your config file using the information you collected in the Gather Configuration Information step in the Prepare chapter.

The following example shows an automated install of Trafodion on a two-node Hortonworks Hadoop cluster that does not have Kerberos nor LDAP enabled.

Example

1. Run the Trafodion Installer in Automated Setup mode.

```
$ cd $HOME/trafodion-installer/installer
$ ./trafodion_install --config_file my
*******
TRAFODION INSTALLATION START
***INFO: testing sudo access
***INFO: Log file located at /var/log/trafodion/trafodion_install_2016-02-16-21-12-
03.log
***INFO: Config directory: /etc/trafodion
***INFO: Working directory: /usr/lib/trafodion
Trafodion Configuration File Check
*********
***INFO: Testing sudo access on node trafodion-1
***INFO: Testing sudo access on node trafodion-2
***INFO: Testing ssh on trafodion-1
***INFO: Testing ssh on trafodion-2
***INFO: Getting list of all hortonworks nodes
***INFO: Getting list of all hortonworks nodes
***INFO: hortonworks list of nodes: trafodion-1 trafodion-2
***INFO: Testing ssh on trafodion-1
***INFO: Testing ssh on trafodion-2
***INFO: Testing sudo access on trafodion-1
***INFO: Testing sudo access on trafodion-2
***DEBUG: trafodionFullName=trafodion_server-1.3.0.tgz
***INFO: Trafodion version = 1.3.0
***DEBUG: HBase's java_exec=/usr/jdk64/jdk1.7.0_67/bin/java
********
TRAFODION SETUP
```

```
********
***INFO: Starting Trafodion environment setup (2016-02-16-21-12-31)
=== 2016-02-16-21-12-31 ===
# @@@ START COPYRIGHT @@@
# Licensed to the Apache Software Foundation (ASF) under one
# or more contributor license agreements. See the NOTICE file
# distributed with this work for additional information
# regarding copyright ownership. The ASF licenses this file
# to you under the Apache License, Version 2.0 (the
# "License"); you may not use this file except in compliance
# with the License. You may obtain a copy of the License at
9. Accepting Warranty or Additional Liability. While redistributing
the Work or Derivative Works thereof, You may choose to offer, and
charge a fee for, acceptance of support, warranty, indemnity, or
other liability obligations and/or rights consistent with this
License. However, in accepting such obligations, You may act only
on Your own behalf and on Your sole responsibility, not on behalf
of any other Contributor, and only if You agree to indemnify, defend,
and hold each Contributor harmless for any liability incurred by,
or claims asserted against, such Contributor by reason of your
accepting any such warranty or additional liability.
END OF TERMS AND CONDITIONS
BY TYPING "ACCEPT" YOU AGREE TO THE TERMS OF THIS AGREEMENT: ***INFO: testing sudo
access
***INFO: Checking all nodes in specified node list
trafodion-1
trafodion-2
***INFO: Total number of nodes = 2
***INFO: Starting Trafodion Package Setup (2016-02-16-21-12-35)
***INFO: Installing required packages
***INFO: Log file located in /var/log/trafodion
***INFO: ... EPEL rpm
***INFO: ... pdsh on node trafodion-1
***INFO: ... pdsh on node trafodion-2
***INFO: Checking if log4cxx is installed ...
***INFO: Checking if sqlite is installed ...
***INFO: Checking if expect is installed ...
***INFO: trafodion user added successfully
***INFO: Trafodion environment setup completed
***INFO: creating sqconfig file
***INFO: Reserving DCS ports
*******
TRAFODION MODS
```

```
*********
***INFO: Hortonworks installed will run traf_hortonworks_mods98
***INFO: Detected JAVA version 1.7
***INFO: copying hbase-trx-hdp2_2-1.3.0.jar to all nodes
PORT=:8080
######## Performing 'set' hbase.master.distributed.log.splitting:false on
(Site:hbase-site, Tag:version1)
######## PUTting json into: doSet_version1455657199513777160.json
Starting the REST environment now
starting rest, logging to /home/trafodion/apache-trafodion-1.3.0-incubating-
bin/rest-1.3.0/bin/../logs/rest-trafodion-1-rest-trafodion-1.out
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/trafodion/apache-trafodion-1.3.0-incubating-
bin/rest-1.3.0/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/hdp/2.2.9.0-3393/hadoop/lib/slf4j-log4j12-
1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
DcsMaster is not started. Please start DCS using 'dcsstart' command...
              Configured Actual
                                            Down
Process
-----
              1
                              0
                                              1
DcsMaster
              2
                             0
DcsServer
                              8
mxosrvr
              8
You can monitor the SQ shell log file: /home/trafodion/apache-trafodion-1.3.0-
incubating-bin/logs/sqmon.log
Startup time 0 hour(s) 1 minute(s) 9 second(s)
Apache Trafodion Conversational Interface 1.3.0
Copyright (c) 2015 Apache Software Foundation
>> initialize trafodion;
--- SQL operation complete.
>>
End of MXCI Session
***INFO: Installation completed successfully.
**********
TRAFODION INSTALLATION COMPLETE
$
```

2. Switch to the Trafodion Runtime User and check the status of Trafodion.

Example

```
$ sudo su - trafodion
$ sqcheck
Checking if processes are up.
Checking attempt: 1; user specified max: 2. Execution time in seconds: 0.
The SQ environment is up!
RMS 4
MXOSRVR 8
                       8
$
```

Trafodion is now running on your Hadoop cluster. Please refer to the Activate chapter for basic instructions on how to verify the Trafodion management and how to perform basic management operations.

Chapter 7. Upgrade

This chapter describes how to use the Trafodion Installer to upgrade Trafodion. You use the Trafodion Provisioning ID to run the Trafodion Installer.

7.1. Download Updated Trafodion Binaries

You perform this step as the Trafodion Provisioning User.

You download the updated Trafodion binaries from the Trafodion Download page. Download the following packages:

- Trafodion Installer (if planning to use the Trafodion Installer)
- Trafodion Server

Refer to Download Trafodion Binaries in the Prepare chapter for examples.

7.2. Unpack Installer

You perform this step as the Trafodion Provisioning User.

You unpack the updated Trafodion Installer into a new directory.

Example

```
$ mkdir $HOME/trafodion-installer-2.0
$ cd $HOME/trafodion-downloads
$ tar -zxf apache-trafodion-installer-2.0.0-incubating-bin.tar.gz -C $HOME/trafodion-installer
$
```

7.3. Stop Trafodion

You perform this step as the Trafodion Runtime User.

Example

```
$ sudo su trafodion
$ sqstop
Shutting down the REST environment now
stopping rest.
Shutting down the DCS environment now
stopping master.
trafodion-1: stopping server.
trafodion-2: stopping server.
stopped $zlobsrv0
stopped $zlobsrv1
Shutting down (normal) the SQ environment!
Wed Feb 17 05:12:40 UTC 2016
Processing cluster.conf on local host trafodion-1
[$Z000KAE] Shell/shell Version 1.0.1 Apache_Trafodion Release 1.3.0 (Build release
[1.3.0-0-g5af956f_Bld2], date 20160112_1927)
ps
```

```
[$Z000KAE] %ps
[$Z000KAE] NID, PID(os) PRI TYPE STATES NAME
                                                  PARENT
                                                              PROGRAM
[$Z000KAE] ----- --- --- --- -----
[$Z000KAE] 000,00064198 000 WDG ES--A-- $WDG000
                                                  NONE
                                                              sqwatchdog
[$Z000KAE] 000,00064199 000 PSD ES--A-- $PSD000
                                                  NONE
                                                              pstartd
[$Z000KAE] 000,00064212 001 GEN ES--A-- $TSID0
                                                  NONE
                                                              idtmsrv
[$Z000KAE] 000,00064242 001 DTM ES--A-- $TM0
                                                  NONE
                                                             tm
[$Z000KAE] 000,00065278 001 GEN ES--A-- $ZSC000
                                                  NONE
                                                              mxsscp
[$Z000KAE] 000,00065305 001 SSMP ES--A-- $ZSM000
                                                  NONE
                                                              mxssmp
[$Z000KAE] 000,00001219 001 GEN ES--A-- $Z0000ZU
                                                  NONE
                                                              mxosrvr
[$Z000KAE] 000,00001235 001 GEN ES--A-- $Z00010A
                                                  NONE
                                                              mxosrvr
[$Z000KAE] 000,00001279 001 GEN ES--A-- $Z00011J
                                                  NONE
                                                              mxosrvr
[$Z000KAE] 000,00001446 001 GEN ES--A-- $Z00016B
                                                  NONE
                                                              mxosrvr
[$Z000KAE] 000,00024864 001 GEN ES--A-- $Z000KAE
                                                  NONE
                                                              shell
[$Z000KAE] 001,00025180 000 PSD ES--A-- $PSD001
                                                  NONE
                                                              pstartd
[$Z000KAE] 001,00025179 000 WDG ES--A-- $WDG001
                                                  NONE
                                                              sqwatchdog
[$Z000KAE] 001,00025234 001 DTM ES--A-- $TM1
                                                  NONE
                                                              tm
[$Z000KAE] 001,00025793 001 GEN ES--A-- $ZSC001
                                                  NONE
                                                              mxsscp
[$Z000KAE] 001,00025797 001 SSMP ES--A-- $ZSM001
                                                  NONE
                                                              mxssmp
[$Z000KAE] 001,00026587 001 GEN ES--A-- $Z010LPM
                                                  NONE
                                                              mxosrvr
[$Z000KAE] 001,00026617 001 GEN ES--A-- $Z010LQH
                                                  NONE
                                                              mxosrvr
[$Z000KAE] 001,00026643 001 GEN ES--A-- $Z010LR8
                                                  NONE
                                                              mxosrvr
[$Z000KAE] 001,00026644 001 GEN ES--A-- $Z010LR9
                                                  NONE
                                                              mxosrvr
shutdown
[$Z000KAE] %shutdown
exit
Issued a 'shutdown normal' request
Shutdown in progress
# of SQ processes: 0
SQ Shutdown (normal) from /home/trafodion Successful
Wed Feb 17 05:12:47 UTC 2016
$
```

7.4. Guided Upgrade

You perform this step as the Trafodion Provisioning User.

As in the case with an installation, the Trafodion Installer prompts you for the information you collected in the Gather Configuration Information step in the Prepare chapter. Some of the prompts are populated with the current values.

The following example shows a guided upgrade of Trafodion on a two-node Cloudera Hadoop cluster without Kerberos nor LDAP enabled.

Example

 Run the updated Trafodion Installer in Guided Setup mode to perform the upgrade. Change information at prompts as applicable.

```
$ cd $HOME/trafodion-installer-2.0/installer
$ ./trafodion_install
TRAFODION INSTALLATION START
***INFO: testing sudo access
***INFO: Log file located at /var/log/trafodion/trafodion_install_2016-02-17-08-15-
33.log
***INFO: Config directory: /etc/trafodion
***INFO: Working directory: /usr/lib/trafodion
*******
Trafodion Configuration Setup
*******
***INFO: Please press [Enter] to select defaults.
Enter trafodion password, default is [traf123]:
Enter list of nodes (blank separated), default []: trafodion-1.apache.org trafodion-
2.apache.org
Specify location of Java 1.7.0_65 or higher (JDK), default is
[/usr/java/jdk1.7.0_67-cloudera]:
Enter full path (including .tar or .tar.gz) of trafodion tar file []:
/home/centos/trafodion-download/apache-trafodion-2.0.0-incubating-bin.tar.gz
Enter Hadoop admin username, default is [admin]:
Enter Hadoop admin password, default is [admin]:
Enter Hadoop external network URL:port (no 'http://' needed), default is []:
trafodion-1.apache.org:7180
Enter HDFS username, default is [hdfs]:
Enter HBase username, default is [hbase]:
Enter HBase group, default is [hbase]:
Enter directory to install trafodion to, default is [/home/trafodion/apache-
trafodion-1.3.0-incubating-bin]: /home/centos/apache-trafodion-2.0.0-incubating-bin
```

```
Start Trafodion after install (Y/N), default is Y:
Total number of client connections per node, default [16]: 8
Enable simple LDAP security (Y/N), default is N:
***INFO: Configuration file: /etc/trafodion/trafodion_config
***INFO: Trafodion configuration setup complete
**********
Trafodion Configuration File Check
*********
***INFO: Testing sudo access on node trafodion-1
***INFO: Testing sudo access on node trafodion-2
***INFO: Testing ssh on trafodion-1
***INFO: Testing ssh on trafodion-2
***INFO: Getting list of all cloudera nodes
***INFO: Getting list of all cloudera nodes
***INFO: cloudera list of nodes: trafodion-1 trafodion-2
***INFO: Testing ssh on trafodion-1
***INFO: Testing ssh on trafodion-2
***INFO: Testing sudo access on trafodion-1
***INFO: Testing sudo access on trafodion-2
***INFO: Checking cloudera Version
***INFO: nameOfVersion=cdh5.3.0
***INFO: HADOOP_PATH=/usr/lib/hbase/lib
***INFO: Trafodion scanner will not be run.
***DEBUG: trafodionFullName=trafodion_server-1.3.0.tgz
***INFO: Trafodion version = 1.3.0
***DEBUG: HBase's java exec=/usr/java/jdk1.7.0 67-cloudera/bin/java
********
TRAFODION SETUP
*******
***INFO: Installing required RPM packages
***INFO: Starting Trafodion Package Setup (2016-02-17-08-16-11)
***INFO: Installing required packages
***INFO: Log file located in /var/log/trafodion
***INFO: ... pdsh on node trafodion-1
***INFO: ... pdsh on node trafodion-2
***INFO: Checking if log4cxx is installed ...
***INFO: Checking if sqlite is installed ...
***INFO: Checking if expect is installed ...
***INFO: Checking if perl-DBD-SQLite* is installed ...
***INFO: Checking if protobuf is installed ...
***INFO: Checking if xerces-c is installed ...
***INFO: Checking if perl-Params-Validate is installed ...
***INFO: Checking if perl-Time-HiRes is installed ...
***INFO: Checking if gzip is installed ...
***INFO: creating sqconfig file
***INFO: Reserving DCS ports
*******
TRAFODION MODS
```

```
********
***INFO: Cloudera installed will run traf_cloudera_mods98
***INFO: Detected JAVA version 1.7
***INFO: copying hbase-trx-cdh5_3-1.3.0.jar to all nodes
***INFO: Cloudera Manager is on trafodion-1
Zookeeper is listening on port 2181
DcsMaster is listening on port 23400
             Configured
                           Actual
Process
                                       Down
             _____
                            ____
_____
                                          ----
DcsMaster 1
DcsServer 2
mxosrvr 8
                            1
                            8
You can monitor the SQ shell log file : /home/trafodion/apache-trafodion-2.0.0-
incubating-bin/logs/sqmon.log
Startup time 0 hour(s) 1 minute(s) 9 second(s)
Apache Trafodion Conversational Interface 1.3.0
Copyright (c) 2015 Apache Software Foundation
>>
End of MXCI Session
***INFO: Installation completed successfully.
*******
TRAFODION INSTALLATION COMPLETE
*********
$
```

2. Switch to the Trafodion Runtime User and check the status of Trafodion.

Trafodion is now running on your Hadoop cluster. Please refer to the Activate chapter for basic instructions on how to verify the Trafodion management and how to perform basic management operations.

7.5. Automated Upgrade

You perform this step as the Trafodion Provisioning User.

The --config_file option runs the Trafodion in Automated Setup mode. Refer to Trafodion Installer in the Introduction chapter for instructions of how you edit your configuration file.

At a minimum, you need to change the following settings:

- LOCAL_WORKDIR
- TRAF_PACKAGE
- SQ ROOT

Example

```
$ cd $HOME/trafodion-configuration
$ cp my_config my_config_2.0
$ # Pre edit content
export LOCAL_WORKDIR="/home/centos/trafodion-installer/installer"
export TRAF_PACKAGE="/home/centos/trafodion-download/apache-trafodion-1.3.0-incubating-bin.tar.gz"
export SQ_ROOT="/home/trafodion/apache-trafodion-1.3.0-incubating-bin"
$ # Use your favorite editor to modify my_config_2.0
$ emacs my_config_2.0
$ # Post edit changes
export LOCAL_WORKDIR="/home/centos/trafodion-installer-2.0/installer"
export TRAF_PACKAGE="/home/centos/trafodion-download/apache-trafodion-2.0.0-incubating-bin.tar.gz"
export SQ_ROOT="/home/trafodion/apache-trafodion-2.0.0-incubating-bin"
```

The following example shows an upgrade of Trafodion on a two-node Hortonworks Hadoop cluster using Automated Setup mode without Kerberos nor LDAP enabled.



The Trafodion Installer performs the same configuration changes as it does for an installation, including restarting Hadoop services.

Example

1. Run the updated Trafodion Installer using the modified my config 2.0 file.

```
$ cd $HOME/trafodion-installer-2.0/installer
$ ./trafodion_install --config_file $HOME/trafodion-configuration/my_config_2.0
*******
TRAFODION INSTALLATION START
*******
***INFO: Testing sudo access on node trafodion-1
***INFO: Testing sudo access on node trafodion-2
***INFO: Testing ssh on trafodion-1
***INFO: Testing ssh on trafodion-2
***INFO: Getting list of all hortonworks nodes
***INFO: Getting list of all hortonworks nodes
***INFO: hortonworks list of nodes: trafodion-1 trafodion-2
***INFO: Testing ssh on trafodion-1
***INFO: Testing ssh on trafodion-2
***INFO: Testing sudo access on trafodion-1
***INFO: Testing sudo access on trafodion-2
***INFO: Trafodion scanner will not be run.
***DEBUG: trafodionFullName=trafodion_server-2.0.0.tgz
***INFO: Trafodion version = 2.0.0
***DEBUG: HBase's java_exec=/usr/jdk64/jdk1.7.0_67/bin/java
TRAFODION SETUP
*******
***INFO: Installing required RPM packages
***INFO: Starting Trafodion Package Setup (2016-02-17-05-33-29)
***INFO: Installing required packages
***INFO: Log file located in /var/log/trafodion
***INFO: ... pdsh on node trafodion-1
***INFO: ... pdsh on node trafodion-2
***INFO: Checking if log4cxx is installed ...
DcsMaster is not started. Please start DCS using 'dcsstart' command...
         Configured Actual
Process
                                           Down
_____
             _____
DcsMaster
                             0
                                            1
                            0
DcsServer
             2
mxosrvr
             8
                            8
You can monitor the SQ shell log file : /home/trafodion/apache-trafodion-2.0.0-
incubating-bin/logs/sqmon.log
Startup time 0 hour(s) 1 minute(s) 9 second(s)
Apache Trafodion Conversational Interface 1.3.0
Copyright (c) 2015 Apache Software Foundation
>>Metadata Upgrade: started
```

2. Switch to the Trafodion Runtime User and check the status of Trafodion.

```
$ sudo su - trafodion
$ sqcheck
Checking if processes are up.
Checking attempt: 1; user specified max: 2. Execution time in seconds: 0.
The SQ environment is up!
         Configured Actual Down
Process
            -----
-----
DTM
            2
            4
RMS
                          4
MXOSRVR 8
$
```

Trafodion is now running on your Hadoop cluster. Please refer to the Activate chapter for basic instructions on how to verify the Trafodion management and how to perform basic management operations.

Chapter 8. Activate

8.1. Manage Trafodion

You use the Trafodion runtime user ID to perform Trafodion management operations.

The following table provides an overview of the different Trafodion management scripts.

Component	Start	Stop	Status
All of Trafodion	sqstart	sqstop	sqcheck
RMS Server	rmsstart	rmsstop	rmscheck
REST Server	reststart	reststop	-
LOB Server	lobstart	lobstop	-
DCS (Database Connectivity Services)	dcsstart	dcsstop	dcscheck

Example: Start Trafodion

```
cd $MY_SQROOT/sql/scripts
sqstart
sqcheck
```

8.2. Validate Trafodion Installation

You can use sqlci (part of the base product) or trafci (requires separate install; see the Trafodion Client Installation Guide) to validate your installation.

8.2.1. Smoke Test

A simple smoke test to validate that Trafodion is functioning.

```
get schemas;
create table table1 (a int);
invoke table1;
insert into table1 values (1), (2), (3), (4);
select * from table1;
drop table table1;
exit;
```

Example

```
$ sqlci
Apache Trafodion Conversational Interface 1.3.0
Copyright (c) 2015 Apache Software Foundation
>>get schemas;
Schemas in Catalog TRAFODION
SEABASE
_MD_
_LIBMGR_
_REPOS_
--- SQL operation complete.
>>create table table1 (a int);
--- SQL operation complete.
>>invoke table1;
-- Definition of Trafodion table TRAFODION.SEABASE.TABLE1
-- Definition current Mon Feb 15 07:42:02 2016
   SYSKEY
                                    LARGEINT NO DEFAULT NOT NULL NOT DROPPABLE
     NOT SERIALIZED
  , A
                                    INT DEFAULT NULL SERIALIZED
  )
--- SQL operation complete.
>>insert into table1 values (1), (2), (3), (4);
--- 4 row(s) inserted.
>>select * from table1;
_____
         1
         2
         3
         4
--- 4 row(s) selected.
>>drop table table1;
--- SQL operation complete.
>>exit;
$
```

Assuming no errors, your installation has been successful. Next, do the following:

- Install the Trafodion client software. Refer to the Trafodion Client Installation Guide.
- Use Trafodion. Refer to the Trafodion SQL Reference Manual.
- Review other Trafodion documentation. Refer to the Trafodion documentation.

8.3. Troubleshooting Tips

If you are not able to start up the environment or if there are problems running sqlci or trafci, then verify that the all the processes are up and running.

• sqcheck should indicate all processes are running.

If processes are not running as expected, then:

- sqstop to shut down Trafodion. If some Trafodion processes do not terminate cleanly, then run ckillall.
- sqstart to restart Trafodion.

If problems persist please review logs:

• \$MY_SQROOT/logs: Trafodion logs.

Chapter 9. Remove

You use the Trafodion Provisioning User for these instructions.



You do not need to use the trafodion_uninstaller script if upgrading Trafodion. Instead, use the trafodion_install script, which automatically upgrades the version of Trafodion. Please refer to the Install chapter for further instructions.

Run the commands from the first node of the cluster. Do not run them from a machine that is not part of the Trafodion cluster.

9.1. Stop Trafodion

Do the following:

```
su trafodion
cd $MY_SQROOT/sql/scripts or cds
sqstop
exit
```

Example

```
[admin@trafodion-1 ~]$ su trafodion
[trafodion@trafodion-1 scripts]$ cds
[trafodion@trafodion-1 scripts]$ sqstop
Shutting down the REST environment now
stopping rest.
Shutting down the DCS environment now
stopping master.
trafodion-1: stopping server.
trafodion-2: stopping server.
stopped $zlobsrv0
stopped $zlobsrv1
Shutting down (normal) the SQ environment!
Mon Feb 15 07:49:18 UTC 2016
Processing cluster.conf on local host trafodion-1
```

```
[$Z000HDS] 001,00024772 001 GEN ES--A-- $Z010K7S
                                                     NONE
                                                                  mxosrvr
[$Z000HDS] 001,00024782 001 GEN ES--U-- $ZLOBSRV1
                                                     NONE
                                                                  mxlobsrvr
shutdown
[$Z000HDS] %shutdown
exit
Issued a 'shutdown normal' request
Shutdown in progress
# of SQ processes: 0
SQ Shutdown (normal) from /home/trafodion/apache-trafodion-1.3.0-incubating-
bin/sql/scripts Successful
Mon Feb 15 07:49:26 UTC 2016
[trafodion@trafodion-1 scripts]$ exit
[admin@trafodion-1 ~]$
```

9.2. Run trafodion uninstaller

The trafodion_uninstaller completely removes Trafodion.

Example

```
[admin@trafodion-1 ~]$ cd $HOME/trafodion-installer/installer
[admin@trafodion-1 installer]$ ./trafodion_uninstaller
Do you want to uninstall Trafodion (Everything will be removed)? (Y/N) y
***INFO: testing sudo access
***INFO: NOTE, rpms that were installed will not be removed.
***INFO: stopping Trafodion instance
SQ environment is not up.
Going to execute ckillall
Can't find file /home/trafodion/.vnc/trafodion-1:1.pid
You'll have to kill the Xvnc process manually
***INFO: restoring linux system files that were changed
***INFO: removing hbase-trx* from Hadoop directories
pdsh@trafodion-1: trafodion-1: ssh exited with exit code 1
pdsh@trafodion-1: trafodion-2: ssh exited with exit code 1
pdsh@trafodion-1: trafodion-1: ssh exited with exit code 1
pdsh@trafodion-1: trafodion-2: ssh exited with exit code 1
***INFO remove the Trafodion userid and group
***INFO: removing all files from /home/trafodion/apache-trafodion-1.3.0-incubating-bin
***INFO: removing all files from /usr/lib/trafodion and /var/log/trafodion
***INFO: removing all files from /etc/trafodion
***INFO: Trafodion uninstall complete.
[admin@trafodion-1 installer]$
```

Chapter 10. Enable Security

Trafodion supports user authentication with LDAP, integrates with Hadoop's Kerberos environment and supports authorization through database grant and revoke requests (privileges).

If this is an initial installation, both LDAP and Kerberos can be configured by running the Trafodion installer. If Trafodion is already installed, then both LDAP and Kerberos can be configured by running the Trafodion security installer.

- If Hadoop has enabled Kerberos, then Trafodion must also enable Kerberos.
- If Kerberos is enabled, then LDAP must be enabled.
- If LDAP is enabled, then database authorization (privilege support) is automatically enabled.
- If Kerberos is not enabled, then enabling LDAP is optional.

10.1. Configuring Trafodion for Kerberos

Kerberos is a protocol for authenticating a request for a service or operation. It uses the notion of a ticket to verify accessibility. The ticket is proof of identity encrypted with a secret key for the particular requested service. Tickets exist for a short time and then expire. Therefore, you can use the service as long as your ticket is valid (i.e. not expired). Hadoop uses Kerberos to provide security for its services, as such Trafodion needs to function properly with Hadoop that have Kerberos enabled.

10.1.1. Kerberos configuration file

It is assumed that Kerberos has already been set up on all the nodes by the time Trafodion is installed. This section briefly discusses the Kerberos configuration file for reference.

The Kerberos configuration file defaults to /etc/krb5.conf and contains, among other attributes:

- * log location: location where Kerberos errors and other information are logged
- * KDC location: host location where the KDC (Key Distribution Center) is located
- * admin server location: host location where the Kerberos admin server is located
- * realm: the set of nodes that share a Kerberos database
- * ticket defaults: contains defaults for ticket lifetimes, encoding, and other attributes

You need to have access to a Kerberos administrator account to enable Kerberos for Trafodion. The following is an example request that lists principals defined in the Kerberos database that can be used to test connectivity:

```
kadmin -p 'kdcadmin/admin' -w 'kdcadmin123' -s 'kdc.server' -q 'listprincs'
* -p (principal): please replace 'kdcadmin/admin' with your admin principal
* -w (password): please replace 'kdadmin123' with the password for the admin principal
* -s (server location): please replace 'kdc.server' with your KDC admin server location
* -q (command): defines the command to run, in this case principals are returned
```

10.1.2. Ticket Management

When Kerberos is enabled in Trafodion, the security installation process:

- Adds a Trafodion principal in Kerberos, one per node with the name trafodion/hostname@realm.
- Creates a keytab for each principal and distributes the keytab to each node. The keytab name is the same for all nodes
 and defaults to a value based on the distribution, for example: etc/trafodion/keytabs/trafodion.service.keytab.
- Performs a "kinit" on all nodes in the cluster for the trafodion user.
- Adds commands to perform "kinit" and to start the ticket renewal procedure to the trafodion bashrc scripts on each node.

The ticket renewal service renews tickets up until the maximum number of renewals allowed. So if your ticket lifetime is one day and the number of renewals is seven days, the ticket renewal service automatically renews tickets six times. Once the ticket expires, it must be initialized again to continue running Trafodion. Connecting to each node as the trafodion user initializes the ticket if one does not exist.

TBD - A future update will include details on how tickets can be managed at the cluster level.

10.1.3. Kerberos installation

The Trafodion installation scripts automatically determine if Kerberos is enabled on the node. If it is enabled, then the environment variable SECURE_HADOOP is set to "Y".

The following are questions that will be asked related to Kerberos:

- Enter KDC server address, default is []: no default
- Enter admin principal (include realm), default is []: no default
- Enter fully qualified name for HBase keytab, default is []: Installer searches for a valid keytab based on the distribution
- Enter fully qualified name for HDFS keytab, default is []: Installer searches for a valid keytab based on the distribution
- Enter max lifetime for Trafodion principal (valid format required), default is [24hours]:

- Enter renew lifetime for Trafodion principal (valid format required), default is [7days]:
- Enter Trafodion keytab name, default is []: Installer determines default name based on the distribution
- Enter keytab location, default is []: Installer determins default name based on the distribution



The Trafodion installer always asked for the KDC admin password when enabling Kerberos independent on whether running in Automated of Guided mode. It does not save this password.

10.2. Configuring LDAP

Trafodion does not manage user names and passwords internally but supports authentication via directory servers using the OpenLDAP protocol, also known as LDAP servers. You can configure the LDAP servers during installation by answering the Trafodion Installer's prompts. To configure LDAP after installation run the Trafodion security installer directly. Installing LDAP also enables database authorization (privilege support).

Once authentication and authorization are enabled, Trafodion allows users to be registered in the database and allows privileges on objects to be granted to users and roles (which are granted to users). Trafodion also supports componentlevel (or system-level) privileges, such as MANAGE_USERS, which can be granted to users and roles. Refer to Manage Users below.



If you do not enable LDAP in Trafodion, then a client interface to Trafodion may request a user name and password, but Trafodion ignores the user name and password entered in the client interface, and the session runs as the database root user, DB__ROOT, without restrictions. If you want to restrict users, restrict access to certain users only, or restrict access to an object or operation, then you must enable security, which enforces authentication and authorization.

10.2.1. Configuring LDAP Servers

To specify the LDAP server(s) to be used for authentication, you need to configure the text file

.traf_authentication_config, located (by default) in \$MY_SQROOT/sql/scripts. This file is a flat file, organized as a series of attribute/value pairs. Details on all the attributes and values accepted in the authentication configuration file and how to configure alternate locations can be found in .traf_authentication_config below.

A sample template file is located in \$MY_SQROOT/sql/scripts/traf_authentication_config.

Attributes and values in the authentication configuration file are separated with a colon immediately following the attribute name. In general, white space is ignored but spaces may be relevant in some values. Attribute names are always case insensitive. Multiple instances of an attribute are specified by repeating the attribute name and providing the new value.

For attributes with only one instance, if the attribute is repeated, the last value provided is used.

Attribute1: valueA Attribute2: valueB Attribute1: valueC

If Attribute1 has only one instance, valueC is used, otherwise, valueA and valueC are both added to the list of values for Attribute1.

Attributes are grouped into sections; this is for future enhancements. Attributes are declared in the LOCAL section, unless otherwise specified.



Section names, attribute names, and the general layout of the authentication configuration file are subject to change in future versions of Trafodion and backward compatibility is not guaranteed.

Specification of your directory server(s) requires at a minimum:

Setting	Description	Example
LDAP Host Name(s)	One or more names of hosts that support the OpenLDAP protocol must be specified. Trafodion attempts to connect to all provided host names during the authentication process. The set of user names and passwords should be identical on all hosts to avoid unpredictable results. The attribute name is LDAPHostName.	LDAPHostName: ldap.company.com
LDAP Port Number	Port number of the LDAP server. Typically this is 389 for servers using no encryption or TLS, and 636 for servers using SSL. The attribute name is LDAPPort.	LDAPPort: 389
LDAP Unique Identifier	Attribute(s) used by the directory server that uniquely identifies the user name. You may provide one or more unique identifier specifiers.	UniqueIdentifier: uid=,ou=users,dc=com

Setting	Description	Example
Encryption Level	A numeric value indicating the encryption scheme used by your LDAP server. Values are:	LDAPSSL: 2
	 0: Encryption not used 1: SSL 2: TLS 	If your LDAP server uses TLS you must specify a file containing the certificate used to encrypt the password. By default the Trafodion software looks for this file in \$MY_SQROOT/cacerts, but you may specify a fully qualified filename, or set the environment variable CACERTS_DIR to another directory. To specify the file containing the certificate, you set the value of the attribute TLS_CACERTFilename, located in the Defaults section.
		Example `TLS_CACERTFilename: mycert.pem TLS_CACertFilename: /usr/etc/cert.pem `
Search username and password	Some LDAP servers require a known user name and password to search the directory of user names. If your environment has that requirement, provide these "search" values.	LDAPSearchDN: lookup@company.com LDAPSearchPwd: Lookup123

There are additional optional attributes that can be used to customize Trafodion authentication. As mentioned earlier, they are described in .traf_authentication_config below.

You can test the authentication configuration file for syntactic errors using the ldapconfigcheck tool. If you have loaded the Trafodion environment (sqenv.sh), then the tool automatically checks the file at \$MY_SQROOT/sql/scripts/.traf_authentication_config. If not, you can specify the file to be checked.

Example

```
ldapconfigcheck -file myconfigfile
File myconfigfile is valid.
```

If an error is found, then the line number with the error is displayed along with the error. Please refer to Idapconfigcheck below for more information.



The authentication configuration file needs to be propagated to all nodes, but there is a script that does that for you described later. For now, you can test your changes on the local node.

You can test the LDAP connection using the utility ldapcheck. To use this utility the Trafodion environment must be loaded (sqenv.sh), but the Trafodion instance does not need to be running. To test the connection only, you can specify

any user name, and a name lookup is performed using the attributes in .traf_authentication_config.

```
ldapcheck --username=fakename@company.com
User fakename@company.com not found
```

If ldapcheck reports either that the user was found or the user was not found, the connection was successful. However, if an error is reported, either the configuration file is not setup correctly, or there is a problem either with your LDAP server or the connection to the server. You can get additional error detail by including the --verbose option. Please refer to Idapcheck for more information.

If you supply a password, ldapcheck attempts to authenticate the specified username and password. The example below shows the password for illustrative purposes, but to avoid typing the password on the command line, leave the password blank (--password=) and the utility prompts for the password with no echo.

ldapcheck --username=realuser@company.com --password=StrongPassword Authentication successful

10.2.2. Generate Trafodion Certificate

Trafodion clients such as trafci encrypt the password before sending it to Trafodion. A self-signed certificate is used to encrypt the password. The certificate and key are generated when the sagen script is invoked. By default, the files server.key and server.crt are located in \$HOME/sqcert. If those files are not present and since Trafodion clients does not send unencrypted passwords, then you need to manually generate those files. To do so, run the script sqcertgen located in \$MY_SQROOT/sql/scripts. The script runs openss1 to generate the certificate and key.

To run openss! manually, follow the example:

```
openssl req -x509 -nodes -days 365 -subj
'/C=US/ST=California/L=PaloAlto/CN=host.domain.com/O=Some Company/OU=Service
- newkey rsa:2048 -keyout server.key -out server.crt
```

Option	Description
-x509	Generate a self-signed certificate.
-days <validity certificate="" of=""></validity>	Make the certificate valid for the days specified.
-newkey rsa: bytes>	Generate a new private key of type RSA of length 1024 or 2048 bytes.
-subj <certificateinfo></certificateinfo>	Specify the information that is incorporated in the certificate. Each instance in a cluster should have a unique common name(CN).
-keyout <filename></filename>	Write the newly generated RSA private key to the file specified.
-nodes	It is an optional parameter that specifies NOT to encrypt the private key. If you encrypt the private key, then you must enter the password every time the private key is used by an application.
-out <filename></filename>	Write the self-signed certificate to the specified file.

Both the public (server.crt) and private (server.key) files should be placed in the directory \$HOME/sqcert. If you do not want to use the HOME directory or if you want to use different names for the private and/or public key files, then Trafodion supports environment variables to specific the alternate locations or names.

• Trafodion first checks the environment variables SQCERT_PRIVKEY and SQCERT_PUBKEY. If they are set, Trafodion uses the fully qualified filename value of the environment variable.

You can specify either one filename environment variable or both.

- If at least one filename environment variable is not set, Trafodion checks the value of the environment variable SQCERT_DIR. If set, then the default filename server.key or server.crt is appended to the value of the environment variable SQCERT_DIR.
- If the filename environment variable is not set and the directory environment variable is not set, then Trafodion uses the default location (\$HOME/sqcert) and the default filename.

10.2.3. Creating the LDAP configuration file

The .traf_authentication_config file is user to enable the Trafodion security features.

File Location

By default, the Trafodion authentication configuration file is located in

\$MY_SQROOT/sql/scripts/.traf_authentication_config. If you want to store the configuration file in a different location and/or use a different filename, then Trafodion supports environment variables to specify the alternate location/name.

Trafodion firsts checks the environment variable TRAFAUTH_CONFIGFILE. If set, the value is used as the fully-qualified Trafodion authentication configuration file.

If the environment variable is not set, then Trafodion next checks the variable TRAFAUTH_CONFIGDIR. If set, the value is prepended to .traf_authentication_config and used as the Trafodion authentication file.

If neither is set, Trafodion defaults to \$MY_SQROOT/sql/scripts/.traf_authentication_config.

Template

```
# To use authentication in Trafodion, this file must be configured
# as described below and placed in $MY_SQROOT/sql/scripts and be named
# .traf_authentication_config.
# NOTE: the format of this configuration file is expected to change in the
# next release of Trafodion. Backward compatibility is not quaranteed.
SECTION: Defaults
 DefaultSectionName: local
 RefreshTime: 1800
 TLS_CACERTFilename:
SECTION: local
# If one or more of the LDAPHostName values is a load balancing host, list
# the name(s) here, one name: value pair for each host.
 LoadBalanceHostName:
# One or more identically configured hosts must be specified here,
# one name: value pair for each host.
 LDAPHostName:
# Default is port 389, change if using 636 or any other port
 LDAPPort:389
# Must specify one or more unique identifiers, one name: value pair for each
 UniqueIdentifier:
# If the configured LDAP server requires a username and password to
# to perform name lookup, provide those here.
 LDAPSearchDN:
 LDAPSearchPwd:
# If configured LDAP server requires TLS(1) or SSL (2), update this value
 LDAPSSL:0
# Default timeout values in seconds
 LDAPNetworkTimeout: 30
 LDAPTimeout: 30
 LDAPTimeLimit: 30
# Default values for retry logic algorithm
 RetryCount: 5
 RetryDelay: 2
 PreserveConnection: No
 ExcludeBadHosts: Yes
 MaxExcludeListSize: 3
```

Configuration Attributes

Attribute Name	Purpose	Example Value	Notes
LDAPHostName	Host name of the local LDAP server.	ldap.master.com	If more than one LDAPHostName entry is provided, then Trafodion attempts to connect with each LDAP server before returning an authentication error. Also see the description related to RetryCount and RetryDelay entries.
LDAPPort	Port number of the local LDAP server.	345	Must be numeric value. Related to LDAPSSL entry. Standard port numbers for OpenLDAP are as follows: • Non-secure: 389 • SSL: 636 • TLS: 389
LDAPSearchDN	If a search user is needed, the search user distinguished name is specified here.	cn=aaabbb, dc=demo, dc=net	If anonymous search is allowed on the local server, then this attribute does not need to be specified or can be specified with no value (blank). To date, anonymous search is the normal approach used.
LDAPSearchPWD	Password for the LDAPSearchDN value. See that entry for details.	welcome	None.
LDAPSSL	A numeric value specifying whether the local LDAP server interface is unencrypted or TLS or SSL. Legal values are 0 for unencrypted, 1 for SSL, and 2 for TLS. For SSL/TLS, see the section below on Encryption Support.		None.
UniqueIdentifier	The directory attribute that contains the user's unique identifier.		To account for the multiple forms of DN supported by a given LDAP server, specify the UniqueIdentifier parameter multiple times with different values. During a search, each UniqueIdentifier is tried in the order it is listed in the configuration file.

Attribute Name	Purpose	Example Value	Notes
LDAPNetworkTimeout	Specifies the timeout (in seconds) after which the next LDAPHOSTNAME entry is tried, in case of no response for a connection request. This parameter is similar to NETWORK_TIMEOUT in ldap_conf(5). Default value is 30 seconds.	20	The value must be a positive number or -1. Setting this to -1 results in an infinite timeout.
LDAPTimelimit	Specifies the time to wait when performing a search on the LDAP server for the user name. The number must be a positive integer. This parameter is similar to TIMELIMIT in ldap_conf(5). Default value is 30 seconds.	15	The server may still apply a lower server- side limit on the duration of a search operation.
LDAPTimeout	Specifies a timeout (in seconds) after which calls to synchronous LDAP APIs aborts if no response is received. This parameter is similar to TIMEOUT in ldap_conf(5). Default value is 30 seconds.	15	The value must be a positive number or -1. Setting this to -1 results in an infinite timeout.
RetryCount	Number of attempts to establish a successful LDAP connection. Default is 5 retries before returning an error.	10	When a failed operation is retried, it is attempted with each configured LDAP server, until the operation is successful or the number of configured retries is exceeded.
RetryDelay	Specifies the number of seconds to delay between retries. Default value is 2 seconds. See description of RetryCount.	1	None.
PreserveConnection	Specifies whether the connection to LDAP server is maintained (YES) or closed (NO) once the operation finishes. Default value is NO.	YES	None.

Attribute Name	Purpose	Example Value	Notes
RefreshTime	Specifies the number of seconds that must have elapsed before the configuration file is reread. Default is 1800 (30 minutes).	f 3600	If set to zero, the configuration file is never read. The connectivity servers must be restarted for changes to take effect if this value is zero. This attribute is not specific to either configuration and must be defined in the DEFAULTS section.
TLS_CACERTFilename	Specifies the location of the certificate file for the LDAP server(s). Filename can either be fully qualified or relative to \$CACERTS_DIR.	cert.pem	This attribute applies to both configurations. If a configuration does not require a certificate, then this attribute is ignored. This attribute must be defined in the DEFAULTS section.
DefaultSectionName	Specifies the configuration type that is assigned to a user by the REGISTER USER command if no authentication type is specified. In the initial Trafodion release, only one configuration is supported.	LOCAL	This attribute must be defined in the DEFAULTS section. If the DefaultSectionName attribute is specified, then a section by that name (or equivalent) must be defined in .traf_ldapconfig. Legal values are LOCAL and ENTERPRISE. This syntax is likely to change.

10.2.4. Verifying configuration and users through Idapcheck

Usage

```
| Idapcheck [<option>]...
| Idapcheck | Coption>]...
| Idapcheck | Coption>
```

Considerations

- Aliases for primary include enterprise and local. Aliases for secondary include cluster and remote. If no configuration is specified, primary is assumed.
- The equals sign is required when supplying a value to username or password.

- To be prompted for a password value with no echo, specify the password argument but omit the equals sign and value.
- · Passwords that contain special characters may need to be escaped if the password is specified on the command line or within a script file.
- If the password keyword is not specified, only the username is checked. The tool can therefore be used to test the LDAP configuration and connection to the configured LDAP server(s) without knowing a valid username or password.

10.2.5. Verifying contents of configuration file through Idapconfigcheck

This page describes the ldapconfigcheck tool, which validates the syntactic correctness of a Trafodion authentication configuration file. Trafodion does not need to be running to run the tool.

Considerations

If the configuration filename is not specified, then the tool looks for a file using environment variables. Those environment variables and the search order are:

1. TRAFAUTH CONFIGFILE

A fully qualified name is expected.

2. TRAFAUTH_CONFIGDIR

Filename .traf_authentication_config/ is appended to the specified directory

3. MY_SQROOT

/sql/scripts/.traf_authentication_config is appended to the value of MY_SQROOT.

Errors

One of the following is output when the tool is run. Only the first error encountered is reported.

Code	Text
0	File filename is valid.
1	File filename not found.
2	File: filename
	Invalid attribute name on line line-number.
3	File: filename
	Missing required value on line line-number.
4	File: filename
	Value out of range on line line-number .
5	File: filename
	Open of traf_authentication_config file failed.
6	File: filename
	Read of traf_authentication_config file failed.
7	No file provided. Either specify a file parameter or verify environment variables.
8	TLS was requested in at least one section, but TLS_CACERTFilename was not provided.
9	Missing host name in at least one section.
	Each LDAP connection configuration section must provide at least one host name.
10	Missing unique identifier in at least one section.
	Each LDAP connection configuration section must provide at least one unique identifier.
11	At least one LDAP connection configuration section must be specified.
12	<pre>Internal error parsing .traf_authentication_config.</pre>

10.3. Manage Users

Kerberos is enabled for installations that require a secure Hadoop environment. LDAP is enabled to enforce authentication for any user connecting to Trafodion. The Trafodion database enforces privileges on the database, database schemas, database objects (table, views, etc) and database operations. Privileges are enforced when authorization is enabled. When LDAP or Kerberos is enabled, authorization is automatically enabled.

To determine the status of authentication and authorization, bring up sqlci and perform "env;".

>>env; Current Environment _____ AUTHENTICATION enabled AUTHORIZATION enabled CURRENT DIRECTORY /.../incubator-trafodion/install/installer LIST_COUNT 4294967295 LOG FILE MESSAGEFILE /.../incubator-trafodion/core/sqf/export/ ...
MESSAGEFILE LANG US English MESSAGEFILE VRSN {2016-06-14 22:27 LINUX:host/user} SQL CATALOG TRAFODION SCH SQL SCHEMA SQL USER CONNECTED user not connected SQL USER DB NAME SQLUSER1 SQL USER ID 33367 TERMINAL CHARSET ISO88591 TRANSACTION ID TRANSACTION STATE not in progress on WARNINGS

Once authorization is enabled, there is one predefined database user called DB_ROOT associated with your specified LDAP username. Please connect to the database and this user and register users that will perform database admin management. The database admin can then connect and setup required users, roles, and privileges.

TBD - A future update should include a pointer to the security best practices guide.

To learn more about how to register users, grant object and component privileges, and manage users and roles, please see the Trafodion SQL Reference Manual.

Chapter 11. Install Recipe

To be written.

Chapter 12. Upgrade Recipe

To be written.