



Provisioning Guide

Version 2.3.0

Table of Contents

1. About This Document	2
1.1. Intended Audience	2
1.2. New and Changed Information	2
1.3. Notation Conventions	3
1.4. Comments Encouraged	5
2. Quick Start	7
2.1. Download Binaries	7
2.2. Unpack Installer and Server package	8
2.3. Collect Information	9
2.3.1. Java Location	9
2.3.2. Data Nodes	10
2.3.3. Distribution Manager URL	10
2.4. Run Installer	11
3. Introduction	12
3.1. Security Considerations	12
3.2. Provisioning Options	13
3.3. Provisioning Activities	13
3.4. Provisioning Master Node	14
3.5. Trafodion Installer	14
3.5.1. Usage	15
3.5.2. Install vs. Upgrade	16
3.5.3. Guided Setup	16
3.5.4. Automated Setup	16
3.6. Trafodion Provisioning Directories	19
4. Requirements	20
4.1. General Cluster and OS Requirements and Recommendations	20
4.1.1. Hardware Requirements and Recommendations	20
4.1.2. OS Requirements and Recommendations	21
4.1.3. IP Ports	22
4.2. Prerequisite Software	23
4.2.1. Hadoop Software	23
4.2.2. Software Packages	23
4.3. Trafodion User IDs and Their Privileges	24
4.3.1. Trafodion Runtime User	24
4.3.2. Trafodion Provisioning User	24
4.4. Recommended Configuration Changes	26
4.4.1. Recommended Security Changes	27
4.4.2. Recommended HDFS Configuration Changes	27
4.4.3. Recommended HBase Configuration Changes	27
5. Prepare	29
5.1. Install Optional Workstation Software	29

5.2. Configure Installation User ID	29
5.3. Disable requiretty	30
5.4. Verify OS Requirements and Recommendations	30
5.5. Configure Kerberos	31
5.6. Configure LDAP Identity Store	31
5.7. Gather Configuration Information	31
5.8. Install Required Software Packages	34
5.8.1. Download and Install Packages	34
5.9. Download Trafodion Binaries	35
6. Install with Ambari	36
6.1. Local Repository	36
6.2. Install Ambari Management Pack for Trafodion	37
6.3. Install Trafodion	37
6.3.1. Initial Cluster Creation	37
6.3.2. Existing Cluster	38
6.4. Initialize Meta-Data	39
6.5. Upgrade from Prior Version	40
Upgrade Entire HDP Stack	40
Upgrade Only Trafodion	40
7. Install	41
7.1. Unpack Installer	41
7.2. Guided Install	42
7.3. Automated Install	47
8. Upgrade	51
8.1. Download Updated Trafodion Binaries	51
8.2. Unpack Installer	51
8.3. Stop Trafodion	51
8.4. Guided Upgrade	53
8.5. Automated Upgrade	54
9. Activate	55
9.1. Manage Trafodion	55
9.2. Validate Trafodion Installation	55
9.2.1. Smoke Test	55
9.3. Troubleshooting Tips	57
10. Remove	58
10.1. Stop Trafodion	58
10.2. Run db_uninstall.py	59
11. Enable Security	61
11.1. Configuring Trafodion for Kerberos	61
11.1.1. Kerberos configuration file	61
11.1.2. Ticket Management	62
11.1.3. Kerberos installation	62
11.2. Configuring LDAP	63
11.2.1. Configuring LDAP Servers	63

11.2.2. Generate Trafodion Certificate	67
11.2.3. Creating the LDAP configuration file	68
11.2.4. Verifying configuration and users through ldapcheck	72
11.2.5. Verifying contents of configuration file through ldapconfigcheck	73
11.3. Manage Users	74

License Statement

Licensed to the Apache Software Foundation (ASF) under one or more contributor license agreements. See the NOTICE file distributed with this work for additional information regarding copyright ownership. The ASF licenses this file to you under the Apache License, Version 2.0 (the "License"); you may not use this file except in compliance with the License. You may obtain a copy of the License at

<http://www.apache.org/licenses/LICENSE-2.0>

Unless required by applicable law or agreed to in writing, software distributed under the License is distributed on an "AS IS" BASIS, WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied. See the License for the specific language governing permissions and limitations under the License.

Disclaimer: *Apache Trafodion is an effort undergoing incubation at the Apache Software Foundation (ASF), sponsored by the Apache Incubator PMC. Incubation is required of all newly accepted projects until a further review indicates that the infrastructure, communications, and decision making process have stabilized in a manner consistent with other successful ASF projects. While incubation status is not necessarily a reflection of the completeness or stability of the code, it does indicate that the project has yet to be fully endorsed by the ASF. <<<*

Revision History

Version	Date
2.2.0	TBD
2.1.0	May 1, 2017
2.0.1	July 7, 2016
2.0.0	June 6, 2016
1.3.0	January, 2016

Chapter 1. About This Document

This guide describes how to provision the end-user Trafodion binaries on top of an existing Hadoop environment. This install allows you to store and query data using Trafodion, either via Trafodion clients (see [Trafodion Client Installation Guide](#)) or via application code you write.

If you want to install a Trafodion developer-build environment, then please refer to the [Trafodion Contributor Guide](#) for instructions.

1.1. Intended Audience

This guide assumes that you are well-versed in Linux and Hadoop administration. If you don't have such experience, then you should consider going through the steps required to install a Hadoop environment before attempting to install Trafodion.

The instructions contained herein apply to the following environments.

- **Single-Node Environments:** Typically used when you want to evaluate Trafodion.
- **Cluster (Multi-Node) Environments:** Typically used when you deploy Trafodion for application usage.



Trafodion can be provisioned on a single-node or multi-node environment. Unless specifically noted, the term **cluster** is used to mean both single- and multi-node environments.

The provisioning instructions applies to a diverse set of platforms:

- **Virtual Machines:** Often used for evaluations and Trafodion development.
- **Cloud:** Used for Product Environments as well as for Developer Environments.
- **Bare Metal:** Used for Product Environments as well as for Developer Environments.



The term **node** is used to represent a computing platform on which operating system, Hadoop, and Trafodion software is running. Unless specifically qualified (bare-metal node, virtual-machine node, or cloud-node), **node** represents a computing platform in your cluster regardless of platform type.

1.2. New and Changed Information

This guide has been updated to include Ambari installation.

1.3. Notation Conventions

This list summarizes the notation conventions for syntax presentation in this manual.

- UPPERCASE LETTERS

Uppercase letters indicate keywords and reserved words. Type these items exactly as shown. Items not enclosed in brackets are required.

```
SELECT
```

- lowercase letters

Lowercase letters, regardless of font, indicate variable items that you supply. Items not enclosed in brackets are required.

```
file-name
```

- [] Brackets

Brackets enclose optional syntax items.

```
DATETIME [start-field TO] end-field
```

A group of items enclosed in brackets is a list from which you can choose one item or none.

The items in the list can be arranged either vertically, with aligned brackets on each side of the list, or horizontally, enclosed in a pair of brackets and separated by vertical lines.

For example:

```
DROP SCHEMA schema [CASCADE]  
DROP SCHEMA schema [ CASCADE | RESTRICT ]
```

- {} Braces

Braces enclose required syntax items.

```
FROM { grantee [, grantee ] ... }
```

A group of items enclosed in braces is a list from which you are required to choose one item.

The items in the list can be arranged either vertically, with aligned braces on each side of the list, or horizontally, enclosed in a pair of braces and separated by vertical lines.

For example:

```
INTERVAL { start-field TO end-field }  
{ single-field }  
INTERVAL { start-field TO end-field | single-field }
```

- | Vertical Line

A vertical line separates alternatives in a horizontal list that is enclosed in brackets or braces.

```
{expression | NULL}
```

- ... Ellipsis

An ellipsis immediately following a pair of brackets or braces indicates that you can repeat the enclosed sequence of syntax items any number of times.

```
ATTRIBUTE[S] attribute [, attribute] ...  
{, sql-expression } ...
```

An ellipsis immediately following a single syntax item indicates that you can repeat that syntax item any number of times.

For example:

```
expression-n ...
```

- Punctuation

Parentheses, commas, semicolons, and other symbols not previously described must be typed as shown.

```
DAY (datetime-expression)
@script-file
```

Quotation marks around a symbol such as a bracket or brace indicate the symbol is a required character that you must type as shown.

For example:

```
"{" module-name [, module-name] ... "}"
```

- Item Spacing

Spaces shown between items are required unless one of the items is a punctuation symbol such as a parenthesis or a comma.

```
DAY (datetime-expression) DAY(datetime-expression)
```

If there is no space between two items, spaces are not permitted. In this example, no spaces are permitted between the period and any other items:

```
myfile.sh
```

- Line Spacing

If the syntax of a command is too long to fit on a single line, each continuation line is indented three spaces and is separated from the preceding line by a blank line.

This spacing distinguishes items in a continuation line from items in a vertical list of selections.

```
match-value [NOT] LIKE _pattern
    [ESCAPE esc-char-expression]
```

1.4. Comments Encouraged

We encourage your comments concerning this document. We are committed to providing documentation that meets your needs. Send any errors found, suggestions for improvement, or compliments to user@trafodion.incubator.apache.org.

Include the document title and any comment, error found, or suggestion for improvement you have concerning this document.

Chapter 2. Quick Start

This chapter provides a quick start for how to use the command-line Trafodion Installer to install Trafodion. **If you prefer to install on HDP distribution using Ambari, refer to the [Ambari Install](#) section.**

You need the following before using the information herein:

- A supported and running Hadoop environment with HDFS, HBase, and Hive. Refer to the [Release Notes](#) for information about supported versions.
- A user ID with passwordless SSH among all the nodes in the cluster. This user ID must have sudo access.



The Trafodion Installer modifies and restarts your Hadoop environment.

2.1. Download Binaries

You download the Trafodion binaries from the Trafodion [Download](#) page. Download the following packages:

- Trafodion Installer (if planning to use the Trafodion Installer)
- Trafodion Server



You can download and install the Trafodion Clients once you've installed and activated Trafodion. Refer to the [Trafodion Client Install Guide](#) for instructions.

Example

Download the Trafodion Installer and Server binaries:

```

$ mkdir $HOME/trafodion-download
$ cd $HOME/trafodion-download
$ # Download the Trafodion Installer binaries
$ wget http://apache.cs.utah.edu/incubator/trafodion/trafodion-2.1.0.incubating/apache-
trafodion-pyinstaller-2.1.0-incubating.tar.gz
Resolving http://apache.cs.utah.edu... 192.168.1.56
Connecting to http://apache.cs.utah.edu|192.168.1.56|:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 68813 (67K) [application/x-gzip]
Saving to: "apache-trafodion-pyinstaller-2.1.0-incubating.tar.gz"

100%[=====
=====>] 68,813          124K/s   in 0.5s

2016-02-14 04:19:42 (124 KB/s) - "apache-trafodion-pyinstaller-2.1.0-incubating.tar.gz"
saved [68813/68813]

$ wget http://apache.cs.utah.edu/incubator/trafodion/trafodion-2.1.0.incubating/apache-
trafodion_server-2.1.0-RH-x86_64-incubating.tar.gz
Resolving http://apache.cs.utah.edu... 192.168.1.56
Connecting to http://apache.cs.utah.edu|192.168.1.56|:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 214508243 (205M) [application/x-gzip]
Saving to: "apache-trafodion_server-2.1.0-RH-x86_64-incubating.tar.gz"

100%[=====
=====>] 214,508,243 3.90M/s   in 55s

2016-02-14 04:22:14 (3.72 MB/s) - "apache-trafodion_server-2.1.0-RH-x86_64-
incubating.tar.gz" saved [214508243/214508243]

$ ls -l
-rw-rw-r--. 1 centos centos      74237 Feb 13 14:53 apache-trafodion_pyinstaller-2.1.0-
incubating.tar.gz
-rw-rw-r--. 1 centos centos 183114066 Feb 10 22:34 apache-trafodion_server-2.1.0-RH-
x86_64-incubating.tar.gz
$

```

2.2. Unpack Installer and Server package

The first step in the installation process is to unpack the Trafodion Installer tar file. Trafodion server package tar file can be auto detected by installer if put it in installer's folder.

Example

```
$ mkdir $HOME/trafodion-installer
$ cd $HOME/trafodion-downloads
$ tar -zxvf apache-trafodion-pyinstaller-2.1.0-incubating.tar.gz -C $HOME/trafodion-
installer
$ cp -f apache-trafodion_server-2.1.0-RH-x86_64-incubating.tar.gz $HOME/trafodion-
installer
$ ls $HOME/trafodion-installer/python-installer
apache-trafodion_server-2.1.0-RH-x86_64-incubating.tar.gz  db_install.py      DISCLAIMER
LICENSE  prettytable.py  scripts
configs                                          db_uninstall.py
discovery.py  NOTICE  README.md
$
```

2.3. Collect Information

Collect/decide the following information:

2.3.1. Java Location

Java location can be automatically detected by installer. You need to provide the java location only if installer cannot detect it.

How to detect java location manually:

1. Login to trafodion' node
2. Use `ps -ef | grep java | grep hadoop | grep hbase` to determine what version HBase is running.

Example

```
ps -ef | grep java | grep hadoop | grep hbase
hbase      17302  17288   1 20:35 ?          00:00:10 /usr/jdk64/jdk1.7.0_67/bin/java
-Dproc_master -XX:OnOutOfMemoryError=kill -9 %p -Dhdp.version=2.3.6.0-3796
-XX:+UseConcMarkSweepGC -XX:ErrorFile=/var/log/hbase/hs_err_pid%p.log
-Djava.io.tmpdir=/tmp -verbose:gc -XX:+PrintGCDetails -XX:+PrintGCDateStamps
-Xloggc:/var/log/hbase/gc.log-201606302035 -Xmx1024m -XX:PermSize=128m
-XX:MaxPermSize=128m -Dhbase.log.dir=/var/log/hbase -Dhbase.log.file=hbase-hbase-master
-ip-172-31-56-238.log -Dhbase.home.dir=/usr/hdp/current/hbase-master/bin/..
-Dhbase.id.str=hbase -Dhbase.root.logger=INFO,RFAS -Djava.library.path=/usr/hdp/2.3.6.0
-3796/hadoop/lib/native/Linux-amd64-64:/usr/hdp/2.3.6.0-3796/hadoop/lib/native
-Dhbase.security.logger=INFO,RFAS org.apache.hadoop.hbase.master.HMaster start
```

The Java location is: `/usr/jdk64/jdk1.7.0_67`

2.3.2. Data Nodes

Trafodion is installed on all data nodes in your Hadoop cluster. Data nodes can be automatically detected by installer while installing on a HDP/CDH cluster.

You need to record hostname for each node when you install Trafodion on Apache Hadoop. For example, refer to `/etc/hosts`.

Example

```
$ cat /etc/hosts
127.0.0.1    localhost localhost.localdomain localhost4 localhost4.localdomain4
::1         localhost localhost.localdomain localhost6 localhost6.localdomain6

172.31.56.238      node-1.internal node-1
172.31.61.110      node-2.internal node-2
```

Record the node names in a comma-separated list `node-1, node-2` or in regular expression mode `node-[1-2]`

2.3.3. Distribution Manager URL

The Installer interacts with the Distribution Manager (for example, Apache Ambari or Cloudera Manager) to modify the Hadoop configuration.

Example

Apache Ambari URL

```
http://myhost.com:8080
```

2.4. Run Installer

You run the Installer once you've collected the base information as described in [Collect Information](#) above.

Please refer to [Guided Install](#) for the **example** of installing Trafodion on a two-node Cloudera Hadoop cluster.

Chapter 3. Introduction

Trafodion is a Hadoop add-on service that provides transactional SQL on top of HBase. Typically, you use Trafodion as the database for applications that require Online Transaction Processing (OLTP), Operational Data Store (ODS), and/or strong reporting capabilities. You access Trafodion using standard JDBC and ODBC APIs.

You may choose whether to add Trafodion to an existing Hadoop environment or to create a standalone Hadoop environment specifically for Hadoop.

This guide assumes that a Hadoop environment exists upon which your provisioning Trafodion. Refer to [Hadoop Software](#) for information about what Hadoop software is required Trafodion.

3.1. Security Considerations

The following users and principals need be considered for Trafodion:

- **Provisioning User:** A Linux-level user that performs the Trafodion provisioning tasks. This user ID requires `sudo` access and passwordless ssh among the nodes where Trafodion is installed. In addition, this user ID requires access to Hadoop distribution, HDFS, and HBase administrative users to change respective environment's configuration settings per Trafodion requirements. Refer to [Trafodion Provisioning User](#) for more information about the requirements and usage associated with this user ID.
- **Runtime User:** A Linux-level user under which the Trafodion software runs, default name is `trafodion`. This user ID must be registered as a user in the Hadoop Distributed File System (HDFS) to store and access objects in HDFS, HBase, and Hive. In addition, this user ID requires passwordless access among the nodes where Trafodion is installed. Refer to [Trafodion Runtime User](#) for more information about this user ID.
- **Trafodion Database Users:** Trafodion users are managed by Trafodion security features (grant, revoke, etc.), which can be integrated with LDAP if so desired. These users are referred to as **database users** and do not have direct access to the operating system. Refer to [LDAP](#) for details on enabling LDAP for authenticating database users. Refer to [Register User](#), [Grant](#), and other SQL statements in the [Trafodion SQL Reference Manual](#) for more information about managing Trafodion Database Users.

If your environment has been provisioned with Kerberos, then the following additional information is required.

- **KDC admin principal:** Trafodion requires administrator access to Kerberos to create principals and keytabs for the `trafodion` user, and to look-up principal names for HDFS and HBase keytabs. Refer to [Kerberos](#) for more information about the requirements and usage associated with this principal.

- **HBase keytab location:** Trafodion requires administrator access to HBase to grant required privileges to the `trafodion` user. Refer to [Kerberos](#) for more information about the requirements and usage associated with this keytab.
- **HDFS keytab location:** Trafodion requires administrator access to HDFS to create directories that store files needed to perform SQL requests such as data loads and backups. Refer to [Kerberos](#) for more information about the requirements and usage associated with this keytab.

If your environment is using LDAP for authentication, then the following additional information is required.

- **LDAP username for database root access:** When Trafodion is installed, it creates a predefined database user referred to as the `DB__ROOT` user. In order to connect to the database as database root, there must be a mapping between the database user `DB__ROOT` and an LDAP user. Refer to [LDAP](#) for more information about this option.
- **LDAP search user name:** Trafodion optionally requests an LDAP username and password in order to perform LDAP operations such as LDAP search. Refer to [LDAP](#) for more information about this option.

3.2. Provisioning Options

Trafodion includes two options for installation: a plug-in integration with Apache Ambari and command-line installation scripts.

The Ambari integration provides support for Hortonworks Hadoop distributions, while the command-line Trafodion Installer supports Cloudera and Hortonworks Hadoop distributions, and for select vanilla Hadoop installations.

The Trafodion Installer supports Linux distributions SUSE and RedHat/CentOS. There are, however, some differences. Prerequisite software packages are not installed automatically on SUSE.

The Trafodion Installer automates many of the tasks required to install/upgrade Trafodion, from downloading and installing required software packages and making required configuration changes to your Hadoop environment via creating the Trafodion runtime user ID to installing and starting Trafodion. It is, therefore, highly recommend that you use the Trafodion Installer for initial installation and upgrades of Trafodion. These steps are referred to as "Script-Based Provisioning" in this guide. Refer to [Trafodion Installer](#) that provides usage information.

3.3. Provisioning Activities

Trafodion provisioning is divided into the following main activities:

- **Requirements:** Activities and documentation required to install the Trafodion software. These activities include tasks such as understanding hardware and operating system requirements, Hadoop requirements, what software packages

that need to be downloaded, configuration settings that need to be changed, and user ID requirements.

- **Prepare:** Activities to prepare the operating system and the Hadoop ecosystem to run Trafodion. These activities include tasks such as installing required software packages, configure the Trafodion Installation User, gather information about the Hadoop environment, and the modify configuration for different Hadoop services.
- **Install:** Activities related to installing the Trafodion software. These activities include tasks such as unpacking the Trafodion tar files, creating the Trafodion Runtime User, creating Trafodion HDFS directories, installing the Trafodion software, and enabling security features.
- **Upgrade:** Activities related to the upgrading the Trafodion software. These activities include tasks such as shutting down Trafodion and installing a new version of the Trafodion software. The upgrade tasks vary depending on the differences between the current and new release of Trafodion. For example, an upgrade may or may not include an upgrade of the Trafodion metadata.
- **Activate:** Activities related to starting the Trafodion software. These actives include basic management tasks such as starting and checking the status of the Trafodion components and performing basic smoke tests.
- **Remove:** Activities related to removing Trafodion from your Hadoop cluster.

3.4. Provisioning Master Node

All provisioning tasks are performed from a single node in the cluster, which can be any node as long as it has access to the Hadoop environment you're adding Trafodion to. This node is referred to as the "**Provisioning Master Node**" in this guide.

The Trafodion Provisioning User must have access to all other nodes from the Provisioning Master Node in order to perform provisioning tasks on the cluster.

3.5. Trafodion Installer

The Trafodion Installer is a set of scripts automates most of the tasks requires to install/upgrade Trafodion. You download the Trafodion Installer tar file from the Trafodion [download](#) page. Next, you unpack the tar file.

Example

```
$ mkdir $HOME/trafodion-installer
$ cd $HOME/trafodion-downloads
$ tar -zxvf apache-trafodion-pyinstaller-2.1.0-incubating.tar.gz -C $HOME/trafodion-
installer
$
```

The Trafodion Installer supports two different modes:

1. **Guided Setup:** Prompts for information as it works through the installation/upgrade process. This mode is recommended for new users.
2. **Automated Setup:** Required information is provided in a pre-formatted ini configuration file, which is provided via a command argument when running the Trafodion Installer thereby suppressing all prompts. This ini configuration file only exists on the **Provisioning Master Node**, please secure this file or delete it after you installed Trafodion successfully.

A template of the configuration file is available here within the installer directory:

`configs/db_config_default.ini`. Make a copy of the file in your directory and populate the needed information.

Automated Setup is recommended since it allows you to record the required provisioning information ahead of time.

Refer to [Automated Setup](#) for information about how to populate this file.

3.5.1. Usage

The following shows help for the Trafodion Installer.

```
$ ./db_install.py -h
*****
Trafodion Installation ToolKit
*****
Usage: db_install.py [options]
Trafodion install main script.

Options:
-h, --help                show this help message and exit
-c FILE, --config-file=FILE
                           Json format file. If provided, all install prompts
                           will be taken from this file and not prompted for.
-u USER, --remote-user=USER
                           Specify ssh login user for remote server,
                           if not provided, use current login user as default.
-v, --verbose              Verbose mode, will print commands.
--silent                  Do not ask user to confirm configuration result
--enable-pwd              Prompt SSH login password for remote hosts.
                           If set, 'sshpass' tool is required.
--build                   Build the config file in guided mode only.
--reinstall               Reinstall Trafodion without restarting Hadoop.
--apache-hadoop           Install Trafodion on top of Apache Hadoop.
--offline                 Enable local repository for offline installing
                           Trafodion.
```

3.5.2. Install vs. Upgrade

The Trafodion Installer automatically detects whether you're performing an install or an upgrade by looking for the Trafodion Runtime User in the `/etc/passwd` file.

- If the user ID doesn't exist, then the Trafodion Installer runs in install mode.
- If the user ID exists, then the Trafodion Installer runs in upgrade mode.
- If `--reinstall` option is specified, then the Trafodion Installer will not restart Hadoop. It's only available when you reinstall the same release version, otherwise an error will be reported during installation.

3.5.3. Guided Setup

By default, the Trafodion Installer runs in Guided Setup mode, which means that it prompts you for information during the install/upgrade process.

Refer to the following sections for examples:

- [Guided Install](#)
- [Guided Upgrade](#)

3.5.4. Automated Setup

The `--config-file` option runs the Trafodion in Automated Setup mode.

Before running the Trafodion Installer with this option, you do the following:

1. Copy the `db_config_default.ini` file.

Example

```
cp configs/db_config_default.ini my_config
```

2. Edit the new file using information you collect in the [Gather Configuration Information](#) section in the [Prepare](#) chapter.
3. Run the Trafodion Installer in Automated Setup Mode

Example

```
./db_install.py --config-file my_config
```



Your Trafodion Configuration File contains the password for the Trafodion Runtime User and for the Distribution Manager. Therefore, we recommend that you secure the file in a manner that matches the security policies of your organization.

Example: Quick start using a Trafodion Configuration File

The Trafodion Installer supports a minimum configuration to quick start your installation in two steps. 1. Copy Trafodion server binary file to your installer directory.

```
cp /path/to/apache-trafodion_server-2.1.0-RH-x86_64-incubating.tar.gz python-installer/
```

1. Modify configuration file `my_config`, add the Hadoop Distribution Manager URL in `mgr_url`.

```
mgr_url = 192.168.0.1:8080
```

Once completed, run the Trafodion Installer with the `--config-file` option.

Example: Creating a Trafodion Configuration File

Using the instructions in [Gather Configuration Information](#) in the [Prepare](#) chapter, record the information and edit `my_config` to contain the following:

```
[dbconfigs]
# NOTICE: if you are using CDH/HDP hadoop distro,
# you can only specifiy management url address for a quick install

#####
# Common Settings
#####

# trafodion username and password
traf_user = trafodion
traf_pwd = traf123
# trafodion user's home directory
home_dir = /home
# the directory location of trafodion binary
# if not provided, the default value will be {package_name}-{version}
traf_dirname =

# trafodion used java(JDK) path on trafodion nodes
# if not provided, installer will auto detect installed JDK
java_home =

# cloudera/ambari management url(i.e. http://192.168.0.1:7180 or just 192.168.0.1)
```

```

# if 'http' or 'https' prefix is not provided, the default one is 'http'
# if port is not provided, the default port is cloudera port '7180'
mgr_url = 192.168.0.1:8080
# user name for cloudera/ambari management url
mgr_user = admin
# password for cloudera/ambari management url
mgr_pwd = admin
# set the cluster number if multiple clusters managed by one Cloudera manager
# ignore it if only one cluster being managed
cluster_no = 1

# trafodion tar package file location
# no need to provide it if the package can be found in current installer's directory
traf_package =
# the number of dcs servers on each node
dcs_cnt_per_node = 4

# scratch file location, seperated by comma if more than one
scratch_locs = $TRAF_VAR

# start trafodion instance after installation completed
traf_start = Y

#####
# DCS HA configuration
#####

# set it to 'Y' if enable DCS HA
dcs_ha = N
# if HA is enabled, provide floating ip, network interface and the hostname of backup
dcs master nodes
dcs_floating_ip =
# network interface that dcs used
dcs_interface =
# backup dcs master nodes, seperated by comma if more than one
dcs_backup_nodes =

#####
# Offline installation setting
#####

# set offline mode to Y if no internet connection
offline_mode = N
# if offline mode is set, you must provide a local repository directory with all needed
RPMs
local_repo_dir =

#####
# LDAP security configuration
#####

# set it to 'Y' if enable LDAP security
ldap_security = N
# LDAP user name and password to be assigned as DB admin privilege

```

```

db_admin_user = admin
db_admin_pwd = traf123
# LDAP user to be assigned DB root privileges (DB__ROOT)
db_root_user = trafodion
# if LDAP security is enabled, provide the following items
ldap_hosts =
# 389 for no encryption or TLS, 636 for SSL
ldap_port = 389
ldap_identifiers =
ldap_encrypt = 0
ldap_certpath =

# set to Y if user info is needed
ldap_userinfo = N
# provide if ldap_userinfo = Y
ladp_user =
ladp_pwd =

#####
# Kerberos security configuration
#####
# if kerberos is enabled in your hadoop system, provide below info

# KDC server address
kdc_server =
# include realm, i.e. admin/admin@EXAMPLE.COM
admin_principal =
# admin password for admin principal, it is used to create trafodion user's principal
and keytab
kdcadmin_pwd =

```

Once completed, run the Trafodion Installer with the `--config-file` option.

Refer to the following sections for examples:

- [Automated Install](#)
- [Automated Upgrade](#)

3.6. Trafodion Provisioning Directories

Trafodion stores its provisioning information in the following directories on each node in the cluster:

- `/etc/trafodion`: Configuration information.

Chapter 4. Requirements

Trafodion requires an x86 version of Linux.

The current release of Trafodion has been tested with:

- 64-bit Red Hat Enterprise Linux (RHEL) or CentOS 6.5 - 6.8
- Cloudera CDH 5.4 - 5.7
- Hortonworks HDP 2.3 - 2.4

Other OS releases may work, too. The Trafodion project is currently working on better support for more distribution and non-distribution versions of Hadoop.

4.1. General Cluster and OS Requirements and Recommendations

64-bit x86 instruction set running a Linux distribution is required. Further, Trafodion assumes an environment based on the requirements of the tested Hadoop distributions/services.

4.1.1. Hardware Requirements and Recommendations

Single-Node Cluster

It is possible to run Trafodion on a single-node sandbox environment. Typically, any sandbox running a Hadoop distribution can be used. A typical single-node configuration uses 4-8 cores with 16 GB of memory, and 20 GB free disk space.

Multi-Node Cluster

For multi-node end-user clusters, your typical HBase environment should suffice for Trafodion. Typically, memory configuration range between 64-128 GB per node with minimum requirement of 16 GB. The cluster size can span from 1 to n nodes; a minimum of two nodes is recommended. A minimum of two cores is required regardless of whether you're deploying Trafodion on a bare-metal or virtual environment.

Recommended configurations:

Attribute	Guidance
Processors per Node	<ul style="list-style-type: none"> • Small: 2 cores • Medium: 4 cores • Large: 8+ cores
Memory per Node	<ul style="list-style-type: none"> • Small: 16 GB • Medium: 64 GB • Large: 128 GB
Concurrency:Nodes	<ul style="list-style-type: none"> • Two Small Nodes: Four concurrent queries • Two Medium Nodes: 64 concurrent queries • Two Large Nodes: 256 concurrent queries

4.1.2. OS Requirements and Recommendations

Please verify these requirements on each node you will install Trafodion on:

Function	Requirement	Verification Guidance
Linux	64-bit version of Red Hat(RHEL) or CentOS 6.5-6.8	<code>cat /etc/redhat-release</code>
sshd	The <code>ssh</code> daemon is running on each node in the cluster.	<ul style="list-style-type: none"> • <code>ps aux grep sshd</code> • <code>sudo netstat -plant grep :22</code>
ntpd	The <code>ntp</code> daemon is running and synchronizing time on each node in the cluster.	<ul style="list-style-type: none"> • <code>ps aux grep ntp</code> • <code>ntpq -p</code>
FQDN	<ul style="list-style-type: none"> • <code>/etc/hosts</code> is set up for fully-qualified node names (FQDN). • <code>/etc/resolv.conf</code> is configured to use a name server. 	<ul style="list-style-type: none"> • <code>hostname --fqdn</code> shows the fully-qualified node name, if any. • The fully-qualified node name is part of the <code>/etc/hosts</code> file. • <code>host -T <FQDN></code> (responds if using a DNS server, times out otherwise) • Simply <code>ssh</code> among nodes using <code>ssh <FQDN></code>.
Port Availability	The Linux Kernel Firewall (iptables) has either been disabled or ports required by Trafodion have been opened.	<ul style="list-style-type: none"> • <code>lsmod grep ip_tables</code> checks whether <code>iptables</code> is loaded. If not, no further checking is needed. • <code>sudo iptables -nL grep <port></code> checks the configuration of a port. An empty response indicates no rule for the port, which often means the port is not open.
passwordless ssh	The user name used to provision Trafodion must have passwordless ssh access to all nodes.	<code>ssh</code> to the nodes, ensure that no password prompt appears.
sudo privileges	The user name used to provision Trafodion must sudo access to a number of root functions	<code>sudo echo "test"</code> on each node.
bash	Available for shell-script execution.	<code>bash --version</code>
java	Available to run the Trafodion software. Same version as HBase is using.	<code>java --version</code>
perl	Available for script execution.	<code>perl --version</code>
python	Available for script execution.	<code>python --version</code>

Function	Requirement	Verification Guidance
yum	Available for installs, updates, and removal of software packages.	<code>yum --version</code>
rpm	Available for installs, updates, and removal of software packages.	<code>rpm --version</code>
scp	Available to copy files among nodes in the cluster.	<code>scp --help</code>

4.1.3. IP Ports

The following table lists the default ports used by the different Trafodion components plus the configuration file and configuration attribute associated with each port setting.

Default Port	Configuration File	Configuration Entry	Required	Range	Protocol	Comment
4200	rest-site.xml	trafodion.rest.port	Yes	1	REST	Trafodion REST Server.
4201	rest-site.xml	trafodion.rest.https.port	Yes	1	HTTPS	Trafodion REST Server (HTTPS).
23400	dcs-site.xml	dcs.master.port	Yes	<i>n</i>	binary	Start of Trafodion DCS port range. (37800 for Trafodion 1.1)
24400	dcs-site.xml	dcs.master.info.port	Yes	1	HTTP	DCS master web GUI. (40010 for Trafodion 1.1)
24410	dcs-site.xml	dcs.server.info.port	Yes	<i>n</i>	HTTP	Start of range for DCS server web GUIs. (40030 for Trafodion 1.1)
50030	mapred-site.xml	mapred.job.tracker.http.address	No	1	HTTP	MapReduce Job Tracker web GUI.
50070	hdfs-site.xml	dfs.http.address	No	1	HTTP	HDFS Name Node web GUI.
50075	hdfs-site.xml	dfs.datanode.http.address	No	1	HTTP	HDFS Data Node web GUI.
50090	hdfs-site.xml	dfs.secondary.http.address	No	1	HTTP	HDFS Secondary Name Node web GUI.
60010	hbase-site.xml	hbase.master.info.port	No	1	HTTP	HBase Master web GUI.
60030	hbase-site.xml	hbase.regionserver.info.port	No	1	HTTP	HBase Region Server web GUI.

There are two port ranges used by Trafodion.

- 23400 is a range, to allow multiple mxosrvr processes on each node. Allow a range of a few ports, enough to cover all the servers per node that are listed in the "servers" file in the DCS configuration directory.
- 24410 is a range as well, enough to cover the DCS servers per node, usually 1 or 2.

On top of the ports identified above, you also need the ports required by your Hadoop distribution. For example:

- [Cloudera Ports](#)
- [Hortonworks Ports](#)

If you have Kerberos or LDAP enabled, then ports required by these products need to be opened as well.

Although not all the ports will be used on every node of the cluster, you need to open most of them for all the nodes in the cluster that have Trafodion, HBase, or HDFS servers on them.

4.2. Prerequisite Software

4.2.1. Hadoop Software

Trafodion runs as an add-on service on Hadoop distributions. The following Hadoop services and their dependencies must be installed and running on the cluster where you intend to install Trafodion:

- Hadoop Distributed File System (HDFS)
- ZooKeeper
- HBase
- Hive
- Apache Ambari (Hortonworks) or Cloudera Manager (Cloudera) with associated embedded databases.

4.2.2. Software Packages

In addition to the software packages required to run different Hadoop services listed above (for example, `Java`), Trafodion requires supplementary software to be installed on the cluster before it is installed. These are Linux tools that are not typically packaged as part of the core Linux distribution.



For RedHat/CentOS, the Trafodion Installer automatically attempts get a subset of these packages over the Internet. If the cluster's access to the Internet is disabled, then you need to manually download the packages and make them available for installation.

Package	Usage	Installation
EPEL	Add-on packages to completed the Linux distribution.	Download Fedora RPM
pdsh	Parallelize shell commands during install and Trafodion runtime utilities.	<code>yum install pdsh</code>
sqlite	Internal configuration information managed by the Trafodion Foundation component.	<code>yum install sqlite</code>
expect	Not used?	<code>yum install expect</code>

Package	Usage	Installation
perl-DBD-SQLite	Allows Perl scripts to connect to SQLite.	yum install perl-DBD-SQLite
perl-Params-Validate	Validates method/function parameters in Perl scripts.	yum install perl-Params-Validate
perl-Time-HiRes	High resolution alarm, sleep, gettimeofday, interval timers in Perl scripts.	yum install perl-Time-HiRes
protobuf	Data serialization.	yum install protobuf
xerces-c	C++ XML parsing.	yum install xerces-c
gzip	Data compress/decompress.	yum install gzip
apr-devel	Support files used to build applications using the APR library.	yum install apr-devel
apr-util-devel	Support files used to build applications using the APR utility library.	yum install apr-util-devel

The Trafodion Installer requires both Internet access/Offline mode to install the required software packages. Specify `db_install.py --offline` to use the offline install feature. Before that, you need to prepare a local repository folder for all the above dependencies. To create a local repository, be sure the `createrepo` package is installed, then run `createrepo` command in your rpm folder.

```
$ createrepo -d .
```

4.3. Trafodion User IDs and Their Privileges

4.3.1. Trafodion Runtime User

The `trafodion:trafodion` user ID is created as part of the installation process. The default password is: `traf123`.

Trafodion requires that either HDFS ACL support or Kerberos is enabled. The Trafodion Installer will enable HDFS ACL and Kerberos support. Refer to [Kerberos](#) for more information about the requirements and usage of Kerberos in Trafodion. Refer to [Apache HBase™ Reference Guide](#) for security in HBase.



Do **not** create the `trafodion:trafodion` user ID in advance. The Trafodion Installer uses the presence of this user ID to determine whether you're doing an installation or upgrade.

4.3.2. Trafodion Provisioning User

Typically, the Trafodion Installer is used for Trafodion installations. It requires access to the user IDs documented below.

Linux Installation User

The user ID that performs the Trafodion installation steps. Typically, this User ID runs the Trafodion Installer.

Requirements:

- User name or group cannot be `trafodion`.
- Passwordless ssh access to all nodes in the cluster.
- Internet access to download software packages.
- `requiretty` must be disabled in `/etc/sudoers`.
- `sudo`¹ access to:
 - Download and install software packages.
 - Modify `/etc/sudoers.d` (allow the `trafodion` user to modify floating IP: `ip` and `arping`).
 - Create the `trafodion` user ID and group.
 - Install Trafodion software into the HBase environment.
 - Run Java version command on each node in the cluster.
 - Run Hadoop version command on each node in the cluster.
 - Run HBase version command on each node in the cluster.
 - Create directories and files in `/etc/trafodion`:
 - Invoke `su` to execute commands as other users; for example, `trafodion`.
 - Edit `sysctl.conf` and activate changes using `sysctl -p`:
 - Modify kernel limits.
 - Reserve IP ports.

¹ `sudo` is **required** in the current release of Trafodion. This restriction may be relaxed in later releases. Alternative mechanisms for privileged access (such as `sudo` alternative commands) are not supported.

Distribution Manager User

A user ID that can change the configuration using Apache Ambari or Cloudera Manager. The Trafodion Installer makes REST request to perform configuration and control functions to the distribution manager using this user ID.

Requirements:

- Administrator user name and password.
- URL to Distribution Manager's REST API.

HDFS Administrator User

The HDFS super user. Required to create directories and change security settings, as needed. The Trafodion Installer uses `su` to run commands under this user ID.

Requirements:

- HDFS Administrator user name.
- Write access to home directory on the node where the Distribution Manager is running.
- For Kerberos enabled installations, location of the keytab for the HDFS service principal.

HBase Administrator User

The HBase super user. Required to change directory ownership in HDFS. For Kerberos enabled installations, the HBase super user is needed to grant the `trafodion` user create, read, write, and execute privileges.

Requirements:

- HBase Administrator user name and group.
- Read access to `hbase-site.xml`.
- For Kerberos enabled installations, location of the keytab for the HBase service principal.

Kerberos Administrator User

The Kerberos administrator. Required to create Trafodion principals and keytabs on a cluster where Kerberos is enabled.

Requirements:

- Kerberos Administrator admin name including the realm.
- Kerberos Administrator password

4.4. Recommended Configuration Changes

The following configuration changes are recommended but not required.



The Trafodion Installer does **not** make these changes.

4.4.1. Recommended Security Changes

The `trafodion` user ID should not be given other `sudo` privileges than what's specified in this manual. Also, we recommend that this user ID is locked (`sudo passwd -l trafodion`) once the installation/upgrade activity has been completed. Users that need issue commands as the `trafodion` ID should do so using `sudo`; for example, `sudo -u trafodion -i`.

4.4.2. Recommended HDFS Configuration Changes

These settings are configured in the `hadoop-env.sh` file.

Property	Recommended Setting	Guidance
DataNode Java Heap Size	2 GB	Use this setting for a large configuration.
NameNode Java Heap Size	2 GB	Use this setting for a large configuration.
Secondary NameNode Java Heap Size	2 GB	Use this setting for a large configuration.

4.4.3. Recommended HBase Configuration Changes

Configuration Property	Recommended Setting	Guidance
<code>hbase.rpc.timeout</code>	10 minutes	This setting depends on the tables' size. Sixty (60) seconds is the default. Increase this value for big tables. Make it the same value as <code>hbase.client.scanner.timeout.period</code> . We have found that increasing the setting to six-hundred (600) seconds will prevent many of the timeout-related errors we encountered, such as <code>OutOfOrderNextException</code> errors.
<code>hbase.client.scanner.timeout.period</code>	10 minutes	Similar to the <code>hbase.rpc.timeout</code> setting. Sixty (60) seconds is the default. Depending on the size of a user table, we have experienced timeout failures on <code>count(*)</code> and <code>update statistics</code> commands from this setting. The underlying issue is the length of the execution of the coprocessor within HBase. NOTE: HBase uses the smaller of <code>hbase.rpc.timeout</code> and <code>hbase.client.scanner.timeout.period</code> to calculate the scanner timeout.
<code>hbase.snapshot.master.timeoutMillis</code> and <code>hbase.snapshot.region.timeout</code>	10 minutes	HBase's default setting is 60000 milliseconds. If you experience timeout issues with HBase snapshots when you use the Trafodion Bulk Loader or other statements, you can set the value for these two HBase properties to 10 minutes (600,000 milliseconds).

Configuration Property	Recommended Setting	Guidance
hbase.hregion.max.filesize	107374182400 bytes	HBase's default setting is 10737418240 (10 GB). We have increased the setting to 107374182400 (100 GB), which reduces the number of HStoreFiles per table and appears to reduce disruptions to active transactions from region splitting.
hbase.hstore.blockingStoreFiles	10	http://gbif.blogspot.com/2012/07/optimizing-writes-in-hbase.html
hbase.regionserver.handler.count	<num>	This setting should match the number of concurrent sessions (mxosrvr). The default is 10.

Chapter 5. Prepare

You need to prepare your Hadoop environment before installing Trafodion.

1. [Install Optional Workstation Software](#)
2. [Configure Installation User ID](#)
3. [Disable requiretty](#)
4. [Verify OS Requirements and Recommendations](#)
5. [Configure Kerberos](#)
6. [Configure LDAP Identity Store](#)
7. [Gather Configuration Information](#)
8. [Install Required Software Packages](#)

5.1. Install Optional Workstation Software

If you are using a Windows workstation, then the following optional software helps installation process. We recommended that you pre-install the software before continuing with the Trafodion installation:

- putty and puttygen (download from [PuTTY web site](#))
- VNC client (download from [RealVNC web site](#))
- Firefox or Chrome browser
- SFTP client to transfer files from your workstation to the Linux server: WinSCP or FileZilla

5.2. Configure Installation User ID

If using the command-line Installer, Trafodion installation requires a user ID with these attributes:

- `sudo` access per the requirements documented in [Linux Installation User](#).
- passwordless ssh to all nodes on the cluster where Trafodion will be installed.



You may need to request permission from your cluster-management team to obtain this type of access.

The following example shows how to set up your user ID to have "passwordless ssh" abilities.

Do the following on the Provision Master Node:

```
echo -e 'y\n' | ssh-keygen -t rsa -N "" -f $HOME/.ssh/id_rsa
cat $HOME/.ssh/id_rsa.pub >> $HOME/.ssh/authorized_keys
echo localhost $(cat /etc/ssh/ssh_host_rsa_key.pub) >> $HOME/.ssh/known_hosts
echo "NoHostAuthenticationForLocalhost=yes" >> $HOME/.ssh/config
chmod 600 $HOME/.ssh/config
chmod 600 $HOME/.ssh/authorized_keys
chmod 700 $HOME/.ssh/
```

After running these commands, do the following:

- If necessary, create the `$HOME/.ssh` directory on the other nodes in your cluster and secure it private to yourself (`chmod 700`).
- If necessary, create the `$HOME/.ssh/authorized_keys` file on the other nodes in your cluster. Secure it with `chmod 600 $HOME/.ssh/authorized_keys`.
- Copy the content of the `$HOME/.ssh/id_rsa.pub` file on the Provisioning Master Node and append the to the `$HOME/.ssh/authorized_keys` file on the other nodes in your cluster.
- `ssh` to the other nodes in the cluster. Answer `y` to the prompt asking you whether to continue the connection. This adds the node to the `$HOME/.ssh/known_hosts` file completing the passwordless `ssh` setup.

5.3. Disable requiretty

If using the command-line Installer, you need to disable `requiretty` in `/etc/sudoers` on all nodes in the cluster to ensure that `sudo` commands can be run from inside the installation scripts.

Comment out the `Defaults requiretty` setting in the `/etc/sudoers` file to ensure that the `requiretty` option is NOT being used.

5.4. Verify OS Requirements and Recommendations

Please ensure that the [OS Requirements and Recommendations](#) are met for each node in the cluster where you intend to install Trafodion.

5.5. Configure Kerberos

If your Hadoop installation has enabled Kerberos, then Trafodion needs to have Kerberos enabled. If not, then Trafodion will not run. If you plan to enable Kerberos in Trafodion, then you need to have access to a KDC (Kerberos Key Distribution Center) and administration credentials so you can create the necessary Trafodion principals and keytabs.

If you wish to manually set up and activate Kerberos principals and keytabs, then refer to the section on [Kerberos](#).

5.6. Configure LDAP Identity Store

If you plan to enable security features in Trafodion, then you need to have an LDAP identity store available to perform authentication. The Trafodion Installer prompts you to set up an authentication configuration file that points to an LDAP server (or servers), which enables security (that is, authentication and authorization) in the Trafodion database.

If you wish to manually set up the authentication configuration file and enable security, then refer to the section on [LDAP](#).

5.7. Gather Configuration Information

You need to gather/decide information about your environment to aid installation Trafodion for the Trafodion Installer. (Listed in alphabetical order to make it easier to find information when referenced in the install and upgrade instructions.)

ID	Information	Default	Notes
admin_principal	Kerberos admin principal to manage principals and keytabs	None	Required if Kerberos is enabled.
cluster_no	Cluster number if multiple clusters managed in ClouderaManager	1	Not required in Hortonworks distro
dcscnt_per_node	Number of concurrent client sessions per node.	4	This number specifies the concurrent sessions per node to be supported. Each session could require up to 1GB of physical memory. The number can be changed post-installation. For more information, refer to the Trafodion Client Installation Guide .
dcsha	Whether to run DCS in high-availability (HA) mode.	N	If Y, you need to provide below dcs configurations.
db_admin_user	LDAP name used to connect as database admin user	admin	Required when LDAP is enabled.
db_root_user	LDAP name used to connect as database root user	trafodion	Required when LDAP is enabled.
dcbackup_nodes	List of nodes where to start the backup DCS Master components.	None	Required when LDAP is enabled. Comma separated FQDN list.

ID	Information	Default	Notes
dcf_floating_ip	IP address if running DCS in HA mode.	None	Required when LDAP is enabled. An FQDN name or IP address.
dcf_interface	Interface type used for dcf_floating_ip.	None	Required when LDAP is enabled. For example, eth0.
home_dir	Root directory under which the trafodion home directory should be created.	/home	Example If the home directory of the trafodion user is /opt/home/trafodion, then specify the root directory as /opt/home.
java_home	Location of Java 1.7.0_65 or higher (JDK).	auto detected	Fully qualified path of the JDK. For example: /usr/java/jdk1.7.0_67-cloudera
kdcadmin_pwd ¹	Password for kerberos admin principal	None	Should be removed from configuration file or secured after install.
kdc_server ¹	Location of host where Kerberos server exists	None	Required if Kerberos enabled.
ldap_security ¹	Whether to enable simple LDAP authentication.	N	If Y, then you need to provide below ldap configurations.
ldap_encrypt ¹	LDAP Encryption Level.	0	0: Encryption not used, 1: SSL, 2: TLS
ldap_certpath ¹	Full path to TLS certificate.	None	Required if ldap_encrypt = 1 or 2.
ldap_hosts ¹	List of nodes where LDAP Identity Store servers are running.	None	Comma separated. FQDN format.
ldap_identifiers ¹	List of LDAP unique identifiers.	None	Comma separated.
ldap_port ¹	Port used to communicate with LDAP Identity Store.	None	Examples: 389 for no encryption or TLS, 636 for SSL.
ldap_userinfo	Whether to use LDAP Search user name.	N	If Y, then you need to provide ldap_user and ldap_pwd.
ldap_user ¹	LDAP Search user name.	None	Required if you need additional LDAP functionality such as LDAPSearch. If so, must provide ldap_pwd too.
ldap_pwd ¹	Password for ldap_user.	None	If ldap_userinfo is required.
local_repo_dir	folder location of Trafodion local repository	None	Required if offline_mode = Y. A local folder with all trafodion rpm dependencies and repodata. For example: /opt/trafodion_repo
mgr_url	FQDN and port for the Distribution Manager's REST API.	None	Include http:// or https:// as applicable. If no prefix, default is http://. Specify in the form: <IP-address>:<port> or <node name>:<port> Example: https://vm-1.yourcompany.local:8080
mgr_user	Administrator user name for Apache Ambari or Cloudera Manager.	admin	A user that can change configuration and restart services via the distribution manager's REST API.
mgr_pwd	Administrator password for Apache Ambari or Cloudera Manager.	admin	A user that can change configuration and restart services via the distribution manager's REST API.
offline_mode	Whether to install Trafodion without internet connection.	N	If Y, then you need to provide local directory in local_repo_dir

ID	Information	Default	Notes
scratch_locs	Overflow scratch files location for large queries that cannot fit in memory.	\$TRAF_VAR	Comma seperated if more than one folder, it should be set in a large disk
traf_dirname	Target folder name for the Trafodion software.	apache-trafodion-{version}	Trafodion is installed in this directory under \$HOME on all nodes in \$NODE_LIST.
traf_package	The location of the Trafodion server package tar file.	auto detected in installer folder	The package file contains the Trafodion server, DCS, and REST software
traf_pwd	The password used for the Trafodion runtime user ID.	traf123	Must be 6-8 characters long.
traf_start	Whether to start Trafodion after install/upgrade.	Y	traf_user

1. Refer to [Enable Security](#) for more information about these security settings.

5.8. Install Required Software Packages

5.8.1. Download and Install Packages

This step is required if you're:

- Installing Trafodion on SUSE.
- Can't download the required software packages using the Internet.

If none of these situations exist, then we highly recommend that you use the Trafodion Installer.

You perform this step as a user with `root` or `sudo` access.

Install the packages listed in [Software Packages](#) above on all nodes in the cluster.

5.9. Download Trafodion Binaries

You download the Trafodion binaries from the Trafodion [Download](#) page. Download the following packages:

Command-line Installation

- Trafodion Installer
- Trafodion Server tar file

Ambari Installation

- Trafodion Ambari RPM
- Trafodion Server RPM




You can download and install the Trafodion Clients once you've installed and activated Trafodion. Refer to the [Trafodion Client Install Guide](#) for instructions.

Chapter 6. Install with Ambari

This method of installation uses RPM packages rather than tar files. There are two packages:

- `traf_ambari` - Ambari management pack (plug-in) that is installed on the Ambari Server node
- `apache-trafodion_server` - Trafodion package that is installed on every data node

You can either set up a local yum repository (requires a web server) or install the RPMs manually on each node.



6.1. Local Repository

On your web server host, be sure the **createrepo** package is installed. Copy the two RPM files into a directory served to the web and run the createrepo command.

```
$ createrepo -d .
```

The command must be used to update repo meta-data any time new RPMs are added or replaced.

Note the Trafodion repository URL for later use.

6.2. Install Ambari Management Pack for Trafodion

On your Ambari server host:

1. If Ambari Server is not already installed, be sure to download a yum repo file for Ambari. For example: [Ambari-2.4.2 repo](#).
2. Add a yum repo file with the URL of your local repository, or copy the traf_ambari RPM locally.
3. Install the Trafodion Ambari management pack RPM. Ambari server will be installed as a dependency, if not already installed.

```
$ sudo yum install traf_ambari
```

4. Set-up Ambari

- a. If Ambari server was previously running, restart it.

```
$ sudo ambari-server restart
```

- b. If Ambari server was not previously running, initialize and start it.

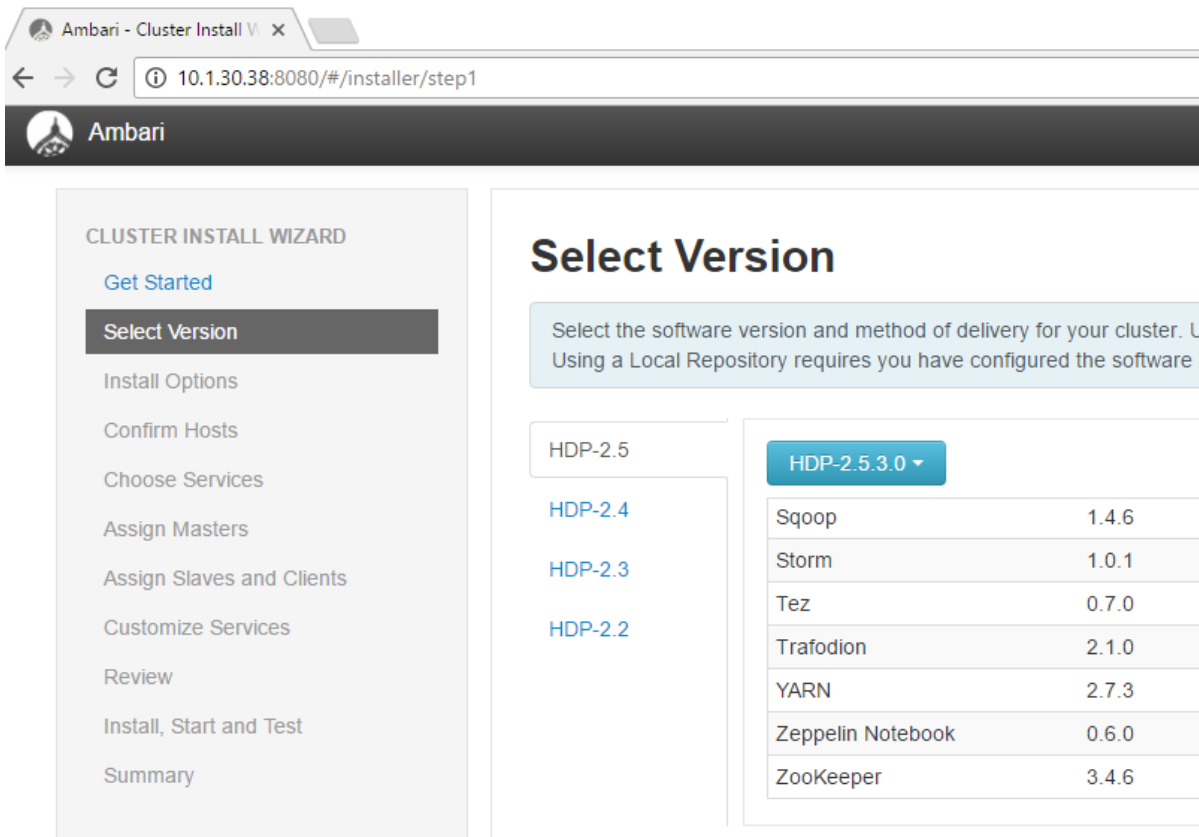
```
$ sudo ambari-server setup  
...  
$ sudo ambari-server start
```

6.3. Install Trafodion

Unlike the command-line installer, Trafodion can be provisioned at time of creating a new cluster.

6.3.1. Initial Cluster Creation

If you are creating a new cluster and you have the Trafodion server RPM hosted on a local yum repository, then create the cluster as normal, and select a stack version that includes the Trafodion service.



When selecting HDP stack version, Ambari prompts for the repository URLs, be **sure to update the Trafodion URL to the URL for your local repository.**

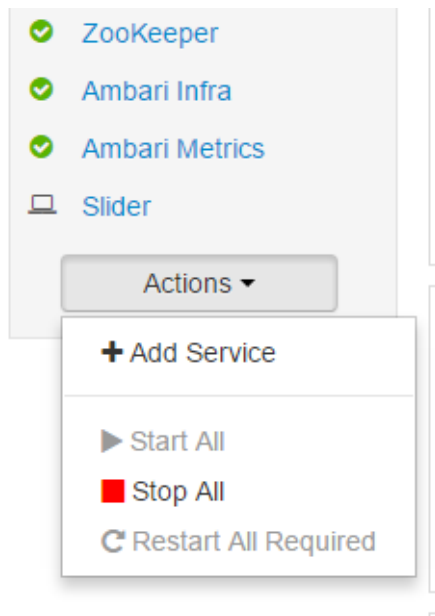
redhat6	HDP-2.5	http://public-repo-1.hortonworks.com/HDP/centos6/2.x/u	Remove
	HDP-UTILS-1.1.0.21	http://public-repo-1.hortonworks.com/HDP-UTILS-1.1.0.21	
	Trafodion-2.1.0	http://no.such.server.org/repo/RH6/2.1.0	

If you plan to install the server RPM manually, set the Trafodion URL to some other valid repository, otherwise the the invalid URL will cause errors later. Any valid repo URL will do. You may copy the HDP URL and paste it into the Trafodion URL. In this case, do not select the Trafodion service on the "Choose Services" step. First, create a cluster without the Trafodion service and follow instructions below for an existing cluster.

6.3.2. Existing Cluster

If you are not using a local yum repository, manually copy the apache-trafodion_server RPM to each data node and install it using yum install.

Using Ambari, select the cluster and then from the Actions menu choose "Add Service" and select Trafodion.

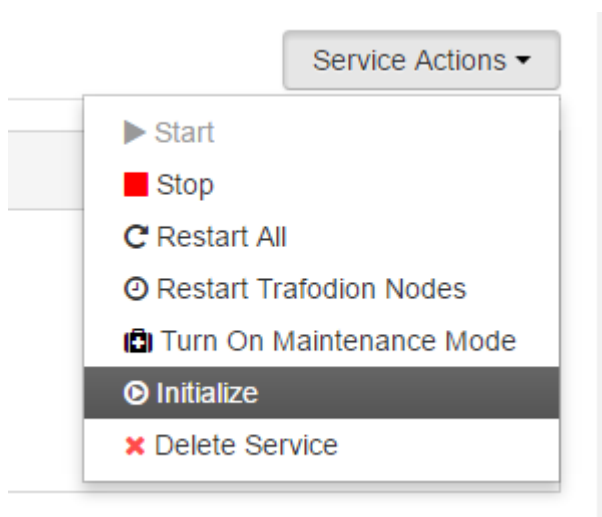


You may get a configuration warning regarding a hbase-site setting. In this case, follow the warning message instructions to go back to the HBase advanced configuration, hbase-site section and use the button to revert to the recommended value for "hbase.coprocessor.region.classes".

Several services will need to be restarted before the Trafodion service can be started. Choose cluster "Actions" menu option "Restart All Required".

6.4. Initialize Meta-Data

After Trafodion is installed and started, it must be initialized before use. Select the Trafodion service and then from the Service Actions menu choose "Initialize" and confirm the action.



6.5. Upgrade from Prior Version

First, be sure to stop the trafodion service prior to upgrading.

Upgrade the traf_ambari package and re-start ambari-server service.

```
$ sudo ambari-server restart
```

Upgrade Entire HDP Stack

Ambari requires version upgrades to be an entire stack, not single components. If you want to include a trafodion upgrade in entire stack version upgrade, first upgrade the traf_ambari package and restart Ambari server.

Once that is done, proceed with the upgrade following Ambari instructions. Be sure to define the Trafodion-2.2 URL that points to a repo server with the new apache-trafodion_server package.

Upgrade Only Trafodion

In this case, Ambari will not upgrade a single component, so the apache-trafodion_server package must be installed on each node directly (via yum command). Once that is done, the service can be started as usual.

Chapter 7. Install

This chapter describes how to use the Trafodion Installer to install Trafodion. You use the Trafodion Provisioning ID to run the Trafodion Installer.

7.1. Unpack Installer

You should already have downloaded the Trafodion Binaries per the instructions in the [Download Trafodion Binaries](#) in the [Prepare](#) chapter. If not, please do so now.

The first step in the installation process is to unpack the Trafodion Installer tar file.

Example

```
$ mkdir $HOME/trafodion-installer
$ cd $HOME/trafodion-downloads
$ tar -zxvf apache-trafodion-pyinstaller-2.1.0-incubating.tar.gz -C $HOME/trafodion-
installer
$
```

7.2. Guided Install

The Trafodion Installer prompts you for the information you collected in the [Gather Configuration Information](#) step in the [Prepare](#) chapter.

The following example shows a guided install of Trafodion on a two-node Cloudera Hadoop cluster that does not have Kerberos nor LDAP installed.

Example

1. Run the Trafodion Installer in guided mode.

```
$ cd $HOME/trafodion-installer/python-installer
$ ./db_install.py
*****
Trafodion Installation ToolKit
*****
Enter HDP/CDH web manager URL:port, (full URL, if no http/https prefix, default
prefix is http://): 192.168.0.31:7180
Enter HDP/CDH web manager user name [admin]:
Enter HDP/CDH web manager user password:
Confirm Enter HDP/CDH web manager user password:

TASK: Environment Discover
*****

Time Cost: 0 hour(s) 0 minute(s) 4 second(s)
Enter full path to Trafodion tar file [/data/python-installer/apache-
trafodion_server-2.1.0-RH-x86_64-incubating.tar.gz]:
Enter directory name to install trafodion to [apache-trafodion-2.1.0]:
Enter trafodion user password:
Confirm Enter trafodion user password:
Enter number of DCS client connections per node [4]:
Enter trafodion scratch file folder location(should be a large disk),
if more than one folder, use comma seperated [$TRAF_VAR]:
Start instance after installation (Y/N) [Y]:
Enable LDAP security (Y/N) [N]:
Enable DCS High Availability (Y/N) [N]:
*****

Final Configs
*****
+-----+
+-----+
+
| config type      | value
|
+-----+
+-----+
+
| dcs_cnt_per_node | 4
```

```

| dcs_ha          | N
| first_rsnode    | node-1
| hbase_user      | hbase
| hdfs_user       | hdfs
| home_dir        | /home
| java_home       | /usr/lib/jvm/java-1.7.0-openjdk.x86_64
| ldap_security   | N
| mgr_url         | http://192.168.0.31:7180
| mgr_user        | admin
| node_list       | node-1,node-2
| scratch_locs    | $TRAF_VAR
| traf_dirname    | apache-trafodion-2.1.0
| traf_package    | /data/python-installer/apache-trafodion_server-2.1.0-RH-x86_64-
incubating.tar.gz |
| traf_start      | Y
| traf_user       | trafodion
|
+-----+
+-----+
+
Confirm result (Y/N) [N]: y

** Generating config file to save configs ...

*****
Installation Start
*****

TASK: Environment Check
*****

Host [node-2]: Script [traf_check.py] ..... [
OK   ]

Host [node-1]: Script [traf_check.py] ..... [
OK   ]

TASK: Copy Trafodion package file

```

```

*****

Script [copy_files.py] ..... [
OK  ]

TASK: Trafodion user Setup
*****

Host [node-2]: Script [traf_user.py] ..... [
OK  ]

Host [node-1]: Script [traf_user.py] ..... [
OK  ]

TASK: Install Trafodion dependencies
*****

Host [node-2]: Script [traf_dep.py] ..... [
OK  ]

Host [node-1]: Script [traf_dep.py] ..... [
OK  ]

TASK: Install Trafodion package
*****

Host [node-2]: Script [traf_package.py] ..... [
OK  ]

Host [node-1]: Script [traf_package.py] ..... [
OK  ]

TASK: Environment Setup
*****

Host [node-1]: Script [traf_setup.py] ..... [
OK  ]

Host [node-2]: Script [traf_setup.py] ..... [
OK  ]

TASK: DCS/REST Setup
*****

Host [node-2]: Script [dcs_setup.py] ..... [

```



```

OK ]

Host [node-1]: Script [dcs_setup.py] ..... [
OK ]

TASK: Hadoop modification and restart
*****

***[INFO]: Restarting CDH services ...
Check CDH services restart status (timeout: 600 secs) .....
***[OK]: CDH services restart successfully!

***[INFO]: Deploying CDH client configs ...
Check CDH services deploy status (timeout: 300 secs) ..
***[OK]: CDH services deploy successfully!

Script [hadoop_mods.py] .....
[ OK ]

TASK: Set permission of HDFS folder for Trafodion user
*****

Host [node-1]: Script [hdfs_cmds.py] .....
[ OK ]

TASK: Sqconfig Setup
*****

Host [node-1]: Script [traf_sqconfig.py] .....
[ OK ]

TASK: Start Trafodion
*****

Host [node-1]: Script [traf_start.py] .....
[ OK ]

Time Cost: 0 hour(s) 7 minute(s) 45 second(s)
*****
Installation Complete
*****

```

2. Switch to the Trafodion Runtime User and check the status of Trafodion.

```

$ sudo su - trafodion
$ sqcheck
*** Checking Trafodion Environment ***

Checking if processes are up.
Checking attempt: 1; user specified max: 2. Execution time in seconds: 0.

The SQ environment is up!

Process          Configured      Actual          Down
-----
DTM               2               2
RMS               4               4
DcsMaster         1               1
DcsServer         2               2
mxosrvr          8               8
RestServer        1               1
$

```

Trafodion is now running on your Hadoop cluster. Please refer to the [Activate](#) chapter for basic instructions on how to verify the Trafodion management and how to perform basic management operations.

7.3. Automated Install

The `--config-file` option runs the Trafodion in Automated Setup mode. Refer to [Trafodion Installer](#) in the [Introduction](#) chapter for instructions of how you edit your configuration file.

Edit your config file using the information you collected in the [Gather Configuration Information](#) step in the [Prepare](#) chapter.

The following example shows an automated install of Trafodion on a two-node Hortonworks Hadoop cluster that does not have Kerberos nor LDAP enabled.

Example

1. Run the Trafodion Installer in Automated Setup mode.

```
$ cd $HOME/trafodion-installer/python-installer
$ ./db_install.py --config-file my_config --silent
*****
Trafodion Installation ToolKit
*****

** Loading configs from config file ...

TASK: Environment Discover
*****

Time Cost: 0 hour(s) 0 minute(s) 4 second(s)

*****
Installation Start
*****

TASK: Environment Check
*****

Host [node-2]: Script [traf_check.py] ..... [
OK   ]

Host [node-1]: Script [traf_check.py] ..... [
OK   ]

TASK: Copy Trafodion package file
*****

Script [copy_files.py] ..... [
OK   ]
```

```

TASK: Trafodion user Setup
*****

Host [node-2]: Script [traf_user.py] ..... [
OK  ]

Host [node-1]: Script [traf_user.py] ..... [
OK  ]

TASK: Install Trafodion dependencies
*****

Host [node-2]: Script [traf_dep.py] ..... [
OK  ]

Host [node-1]: Script [traf_dep.py] ..... [
OK  ]

TASK: Install Trafodion package
*****

Host [node-2]: Script [traf_package.py] ..... [
OK  ]

Host [node-1]: Script [traf_package.py] ..... [
OK  ]

TASK: Environment Setup
*****

Host [node-1]: Script [traf_setup.py] ..... [
OK  ]

Host [node-2]: Script [traf_setup.py] ..... [
OK  ]

TASK: DCS/REST Setup
*****

Host [node-2]: Script [dcs_setup.py] ..... [
OK  ]

Host [node-1]: Script [dcs_setup.py] ..... [

```

```

OK ]

TASK: Hadoop modification and restart
*****

***[INFO]: Restarting CDH services ...
Check CDH services restart status (timeout: 600 secs) .....
***[OK]: CDH services restart successfully!

***[INFO]: Deploying CDH client configs ...
Check CDH services deploy status (timeout: 300 secs) ..
***[OK]: CDH services deploy successfully!

Script [hadoop_mods.py] .....
[ OK ]

TASK: Set permission of HDFS folder for Trafodion user
*****

Host [node-1]: Script [hdfs_cmds.py] .....
[ OK ]

TASK: Sqconfig Setup
*****

Host [node-1]: Script [traf_sqconfig.py] .....
[ OK ]

TASK: Start Trafodion
*****

Host [node-1]: Script [traf_start.py] .....
[ OK ]

Time Cost: 0 hour(s) 7 minute(s) 45 second(s)
*****
Installation Complete
*****

```

2. Switch to the Trafodion Runtime User and check the status of Trafodion.

```

$ sudo su - trafodion
$ sqcheck
*** Checking Trafodion Environment ***

Checking if processes are up.
Checking attempt: 1; user specified max: 2. Execution time in seconds: 0.

The SQ environment is up!

Process          Configured      Actual          Down
-----
DTM               2               2
RMS               4               4
DcsMaster         1               1
DcsServer         2               2
mxosrvr          8               8
RestServer        1               1
$

```

Trafodion is now running on your Hadoop cluster. Please refer to the [Activate](#) chapter for basic instructions on how to verify the Trafodion management and how to perform basic management operations.

Chapter 8. Upgrade

This chapter describes how to use the Trafodion Installer to upgrade Trafodion. You use the Trafodion Provisioning ID to run the Trafodion Installer.

8.1. Download Updated Trafodion Binaries

You perform this step as the Trafodion Provisioning User.

You download the updated Trafodion binaries from the Trafodion [Download](#) page. Download the following packages:

- Trafodion Installer (if planning to use the Trafodion Installer)
- Trafodion Server

Refer to [Download Trafodion Binaries](#) in the [Prepare](#) chapter for examples.

8.2. Unpack Installer

You perform this step as the Trafodion Provisioning User.

You unpack the updated Trafodion Installer into a new directory.

Example

```
$ mkdir $HOME/trafodion-installer
$ cd $HOME/trafodion-downloads
$ tar -zxf apache-trafodion-pyinstaller-2.1.0-incubating.tar.gz -C $HOME/trafodion-
installer
$
```

8.3. Stop Trafodion

You perform this step as the Trafodion Runtime User.

Example

```

$ sudo su trafodion
$ sqstop
Shutting down the REST environment now
stopping rest.
Shutting down the DCS environment now
stopping master.
trafodion-1: stopping server.
trafodion-2: stopping server.
stopped $zlobsrv0
stopped $zlobsrv1
Shutting down (normal) the SQ environment!
Wed Feb 17 05:12:40 UTC 2016
Processing cluster.conf on local host trafodion-1
[$Z000KAE] Shell/shell Version 1.0.1 Apache_Trafodion Release 1.3.0 (Build release
[1.3.0-0-g5af956f_Bld2], date 20160112_1927)
[$Z000KAE] %ps
[$Z000KAE] NID,PID(os)  PRI  TYPE  STATES  NAME          PARENT        PROGRAM
[$Z000KAE] -----
[$Z000KAE] 000,00064198 000 WDG   ES--A-- $WDG000      NONE          sqwatchdog
[$Z000KAE] 000,00064199 000 PSD   ES--A-- $PSD000      NONE          pstartd
[$Z000KAE] 000,00064212 001 GEN   ES--A-- $TSID0       NONE          idtmsrv
[$Z000KAE] 000,00064242 001 DTM   ES--A-- $TM0         NONE          tm
[$Z000KAE] 000,00065278 001 GEN   ES--A-- $ZSC000      NONE          mxsscp
[$Z000KAE] 000,00065305 001 SSMP  ES--A-- $ZSM000      NONE          mxssmp
[$Z000KAE] 000,00001219 001 GEN   ES--A-- $Z0000ZU     NONE          mxosrvr
[$Z000KAE] 000,00001235 001 GEN   ES--A-- $Z00010A     NONE          mxosrvr
[$Z000KAE] 000,00001279 001 GEN   ES--A-- $Z00011J     NONE          mxosrvr
[$Z000KAE] 000,00001446 001 GEN   ES--A-- $Z00016B     NONE          mxosrvr
[$Z000KAE] 000,00024864 001 GEN   ES--A-- $Z000KAE     NONE          shell
[$Z000KAE] 001,00025180 000 PSD   ES--A-- $PSD001      NONE          pstartd
[$Z000KAE] 001,00025179 000 WDG   ES--A-- $WDG001      NONE          sqwatchdog
[$Z000KAE] 001,00025234 001 DTM   ES--A-- $TM1         NONE          tm
[$Z000KAE] 001,00025793 001 GEN   ES--A-- $ZSC001      NONE          mxsscp
[$Z000KAE] 001,00025797 001 SSMP  ES--A-- $ZSM001      NONE          mxssmp
[$Z000KAE] 001,00026587 001 GEN   ES--A-- $Z010LPM     NONE          mxosrvr
[$Z000KAE] 001,00026617 001 GEN   ES--A-- $Z010LQH     NONE          mxosrvr
[$Z000KAE] 001,00026643 001 GEN   ES--A-- $Z010LR8     NONE          mxosrvr
[$Z000KAE] 001,00026644 001 GEN   ES--A-- $Z010LR9     NONE          mxosrvr
shutdown
[$Z000KAE] %shutdown
exit
Issued a 'shutdown normal' request

Shutdown in progress

# of SQ processes: 0
SQ Shutdown (normal) from /home/trafodion Successful
Wed Feb 17 05:12:47 UTC 2016
$

```


8.4. Guided Upgrade

You perform this step as the Trafodion Provisioning User.

As in the case with an installation, the Trafodion Installer prompts you for the information you collected in the [Gather Configuration Information](#) step in the [Prepare](#) chapter.

From user's perspective, guided upgrade doesn't require the Trafodion runtime user's password. The rest of things have no difference with guided install. Please refer to [Guided Install](#) for the **example** of installing Trafodion on a two-node Cloudera Hadoop cluster.

8.5. Automated Upgrade

You perform this step as the Trafodion Provisioning User.

The `--config-file` option runs the Trafodion in Automated Setup mode. Refer to [Trafodion Installer](#) in the [Introduction](#) chapter for instructions of how you edit your configuration file.

At a minimum, you need to change the following settings:

- `traf_package`

Example

```
$ cd $HOME/trafodion-configuration
$ cp my_config my_config_2.0
$ # Pre edit content

traf_package = "/home/centos/trafodion-download/apache-trafodion-2.0.0-
incubating.tar.gz"

$ # Use your favorite editor to modify my_config_2.0
$ emacs my_config_2.0
$ # Post edit changes

traf_package = "/home/centos/trafodion-download/apache-trafodion-2.1.0-
incubating.tar.gz"
```



The Trafodion Installer performs the same configuration changes as it does for an installation, including restarting Hadoop services.

Please refer to [Automated Install](#) for the **example** of installing Trafodion on a two-node Cloudera Hadoop cluster.

Chapter 9. Activate

9.1. Manage Trafodion

You use the Trafodion runtime user ID to perform Trafodion management operations.

The following table provides an overview of the different Trafodion management scripts.

Component	Start	Stop	Status
All of Trafodion	sqstart	sqstop	sqcheck
RMS Server	rmsstart	rmsstop	rmscheck
REST Server	reststart	reststop	-
LOB Server	lobstart	lobstop	-
DCS (Database Connectivity Services)	dcsstart	dcsstop	dcscheck

Example: Start Trafodion

```
cd $TRAF_HOME/sql/scripts
sqstart
sqcheck
```

9.2. Validate Trafodion Installation

You can use sqlci (part of the base product) or trafci (requires separate install; see the [Trafodion Client Installation Guide](#)) to validate your installation.

9.2.1. Smoke Test

A simple smoke test to validate that Trafodion is functioning.

```
get schemas;
create table table1 (a int);
invoke table1;
insert into table1 values (1), (2), (3), (4);
select * from table1;
drop table table1;
exit;
```

Example

```

$ sqlci
Apache Trafodion Conversational Interface 1.3.0
Copyright (c) 2015 Apache Software Foundation
>>get schemas;

Schemas in Catalog TRAFODION
=====

SEABASE
_MD_
_LIBMGR_
_REPOS_

--- SQL operation complete.
>>create table table1 (a int);

--- SQL operation complete.
>>invoke table1;

-- Definition of Trafodion table TRAFODION.SEABASE.TABLE1
-- Definition current  Mon Feb 15 07:42:02 2016

(
    SYSKEY                                LARGEINT NO DEFAULT NOT NULL NOT DROPPABLE
    NOT SERIALIZED
, A                                INT DEFAULT NULL SERIALIZED
)

--- SQL operation complete.
>>insert into table1 values (1), (2), (3), (4);

--- 4 row(s) inserted.
>>select * from table1;

A
-----

      1
      2
      3
      4

--- 4 row(s) selected.
>>drop table table1;

--- SQL operation complete.
>>exit;
$

```

Assuming no errors, your installation has been successful. Next, do the following:

- Install the Trafodion client software. Refer to the [Trafodion Client Installation Guide](#).
- Use Trafodion. Refer to the [Trafodion SQL Reference Manual](#).
- Review other Trafodion documentation. Refer to the [Trafodion documentation](#).

9.3. Troubleshooting Tips

If you are not able to start up the environment or if there are problems running `sqlci` or `trafci`, then verify that the all the processes are up and running.

- `sqcheck` should indicate all processes are running.

If processes are not running as expected, then:

- `sqstop` to shut down Trafodion. If some Trafodion processes do not terminate cleanly, then run `killall`.
- `sqstart` to restart Trafodion.

If problems persist please review logs:

- `$TRAF_HOME/logs`: Trafodion logs.

Chapter 10. Remove

You use the Trafodion Provisioning User for these instructions.



You do not need to use the `db_uninstall.py` script before upgrading Trafodion. Instead, use the `db_install.py` script, which automatically upgrades the version of Trafodion. Please refer to the [Install](#) chapter for further instructions.

Run the commands from the first node of the cluster. You can also run it from any node, but you need to specify the hostnames of Trafodion cluster.

10.1. Stop Trafodion

Do the following:

```
su trafodion
cd $TRAF_HOME/sql/scripts or cds
sqstop
exit
```

Example

```

[admin@trafodion-1 ~]$ su trafodion
[trafodion@trafodion-1 scripts]$ cds
[trafodion@trafodion-1 scripts]$ sqstop
Shutting down the REST environment now
stopping rest.
Shutting down the DCS environment now
stopping master.
trafodion-1: stopping server.
trafodion-2: stopping server.
stopped $zlobsrv0
stopped $zlobsrv1
Shutting down (normal) the SQ environment!
Mon Feb 15 07:49:18 UTC 2016
Processing cluster.conf on local host trafodion-1
.
.
.

[$Z000HDS] 001,00024772 001 GEN  ES--A-- $Z010K7S      NONE      mxosrvr
[$Z000HDS] 001,00024782 001 GEN  ES--U-- $ZLOBSRV1     NONE      mxlobsrvr
shutdown
[$Z000HDS] %shutdown
exit
Issued a 'shutdown normal' request

Shutdown in progress

# of SQ processes: 0
SQ Shutdown (normal) from /home/trafodion/apache-trafodion-1.3.0-incubating-
bin/sql/scripts Successful
Mon Feb 15 07:49:26 UTC 2016
[trafodion@trafodion-1 scripts]$ exit
[admin@trafodion-1 ~]$

```

10.2. Run db_uninstall.py

The `db_uninstall.py` completely removes Trafodion, includes Trafodion user's home directory.

Example

```
[admin@trafodion-1 ~]$ cd $HOME/trafodion-installer/python-installer
[admin@trafodion-1 installer]$ ./db_uninstall.py
*****
    Trafodion Uninstall Start
*****
Uninstall Trafodion on [node-1 node-2] [N]: y

***[INFO]: Remove Trafodion on node [node-1] ...

***[INFO]: Remove Trafodion on node [node-2] ...
*****
    Trafodion Uninstall Completed
*****
[admin@trafodion-1 installer]$
```


Chapter 11. Enable Security

Trafodion supports user authentication with LDAP, integrates with Hadoop's Kerberos environment and supports authorization through database grant and revoke requests (privileges).

If this is an initial installation, both LDAP and Kerberos can be configured by running the Trafodion installer. If Trafodion is already installed, then both LDAP and Kerberos can be configured by running the Trafodion security installer.

- If Hadoop has enabled Kerberos, then Trafodion must also enable Kerberos.
- If Kerberos is enabled, then LDAP must be enabled.
- If LDAP is enabled, then database authorization (privilege support) is automatically enabled.
- If Kerberos is not enabled, then enabling LDAP is optional.

11.1. Configuring Trafodion for Kerberos

Kerberos is a protocol for authenticating a request for a service or operation. It uses the notion of a ticket to verify accessibility. The ticket is proof of identity encrypted with a secret key for the particular requested service. Tickets exist for a short time and then expire. Therefore, you can use the service as long as your ticket is valid (i.e. not expired). Hadoop uses Kerberos to provide security for its services, as such Trafodion needs to function properly with Hadoop that have Kerberos enabled.

11.1.1. Kerberos configuration file

It is assumed that Kerberos has already been set up on all the nodes by the time Trafodion is installed. This section briefly discusses the Kerberos configuration file for reference.

The Kerberos configuration file defaults to `/etc/krb5.conf` and contains, among other attributes:

```
* log location: location where Kerberos errors and other information are logged
* KDC location: host location where the KDC (Key Distribution Center) is located
* admin server location: host location where the Kerberos admin server is located
* realm: the set of nodes that share a Kerberos database
* ticket defaults: contains defaults for ticket lifetimes, encoding, and other
attributes
```

You need to have access to a Kerberos administrator account to enable Kerberos for Trafodion. The following is an example request that lists principals defined in the Kerberos database that can be used to test connectivity:

```
kadmin -p 'kdcadmin/admin' -w 'kdcadmin123' -s 'kdc.server' -q 'listprincs'
* -p (principal): please replace 'kdcadmin/admin' with your admin principal
* -w (password): please replace 'kdcadmin123' with the password for the admin principal
* -s (server location): please replace 'kdc.server' with your KDC admin server location
* -q (command): defines the command to run, in this case principals are returned
```

11.1.2. Ticket Management

When Kerberos is enabled in Trafodion, the security installation process:

- Adds a Trafodion principal in Kerberos, one per node with the name `trafodion/hostname@realm`.
- Creates a keytab for each principal and distributes the keytab to each node. The keytab name is the same for all nodes and defaults to a value based on the distribution, for example: `etc/trafodion/keytabs/trafodion.service.keytab`.
- Performs a "kinit" on all nodes in the cluster for the `trafodion` user.
- Adds commands to perform "kinit" and to start the ticket renewal procedure to the `trafodion .bashrc` scripts on each node.

The ticket renewal service renews tickets up until the maximum number of renewals allowed. So if your ticket lifetime is one day and the number of renewals is seven days, the ticket renewal service automatically renews tickets six times. Once the ticket expires, it must be initialized again to continue running Trafodion. Connecting to each node as the `trafodion` user initializes the ticket if one does not exist.

TBD - A future update will include details on how tickets can be managed at the cluster level.

11.1.3. Kerberos installation

The Trafodion installation scripts automatically determine if Kerberos is enabled on the node. If it is enabled, then the environment variable `SECURE_HADOOP` is set to "Y".

The following are questions that will be asked related to Kerberos:

- Enter KDC server address, default is []: – no default
- Enter admin principal (include realm), default is []: - no default
- Enter password for admin principal:



KDC admin password will be saved only in configuration file `db_config.bakYYMMDD_HHMM` in installer folder when installation completed. You can delete this file for secure perspective. NOTE: Keytab files are auto detected by installer in CDH/HDP cluster. NOTE: Installer doesn't support kerberos enabled Apache Hadoop for this release.

11.2. Configuring LDAP

Trafodion does not manage user names and passwords internally but supports authentication via directory servers using the OpenLDAP protocol, also known as LDAP servers. You can configure the LDAP servers during installation by answering the Trafodion Installer's prompts. To configure LDAP after installation run the Trafodion security installer directly. Installing LDAP also enables database authorization (privilege support).

Once authentication and authorization are enabled, Trafodion allows users to be registered in the database and allows privileges on objects to be granted to users and roles (which are granted to users). Trafodion also supports component-level (or system-level) privileges, such as `MANAGE_USERS`, which can be granted to users and roles. Refer to [Manage Users](#) below.



If you do not enable LDAP in Trafodion, then a client interface to Trafodion may request a user name and password, but Trafodion ignores the user name and password entered in the client interface, and the session runs as the database **root** user, `DB__ROOT`, without restrictions. If you want to restrict users, restrict access to certain users only, or restrict access to an object or operation, then you must enable security, which enforces authentication and authorization.

11.2.1. Configuring LDAP Servers

To specify the LDAP server(s) to be used for authentication, you need to configure the text file `.traf_authentication_config`, located (by default) in `$TRAF_HOME/sql/scripts`. This file is a flat file, organized as a series of attribute/value pairs. Details on all the attributes and values accepted in the authentication configuration file and how to configure alternate locations can be found in [.traf_authentication_config](#) below.

A sample template file is located in `$TRAF_HOME/sql/scripts/traf_authentication_config`.

Attributes and values in the authentication configuration file are separated with a colon immediately following the attribute name. In general, white space is ignored but spaces may be relevant in some values. Attribute names are always case insensitive. Multiple instances of an attribute are specified by repeating the attribute name and providing the new value. For attributes with only one instance, if the attribute is repeated, the last value provided is used.

```

Attribute1: valueA
Attribute2: valueB
Attribute1: valueC

```

If `Attribute1` has only one instance, `valueC` is used, otherwise, `valueA` and `valueC` are both added to the list of values for `Attribute1`.

Attributes are grouped into sections; this is for future enhancements. Attributes are declared in the `LOCAL` section, unless otherwise specified.



Section names, attribute names, and the general layout of the authentication configuration file are subject to change in future versions of Trafodion and backward compatibility is not guaranteed.

Specification of your directory server(s) requires at a minimum:

Setting	Description	Example
LDAP Host Name(s)	One or more names of hosts that support the OpenLDAP protocol must be specified. Trafodion attempts to connect to all provided host names during the authentication process. The set of user names and passwords should be identical on all hosts to avoid unpredictable results. The attribute name is <code>LDAPHostName</code> .	<code>LDAPHostName:</code> <code>ldap.company.com</code>
LDAP Port Number	Port number of the LDAP server. Typically this is 389 for servers using no encryption or TLS, and 636 for servers using SSL. The attribute name is <code>LDAPPort</code> .	<code>LDAPPort: 389</code>
LDAP Unique Identifier	Attribute(s) used by the directory server that uniquely identifies the user name. You may provide one or more unique identifier specifiers.	<code>UniqueIdentifier:</code> <code>uid=,ou=users,dc=com</code>
Encryption Level	A numeric value indicating the encryption scheme used by your LDAP server. Values are: <ul style="list-style-type: none"> • 0: Encryption not used • 1: SSL • 2: TLS 	<p><code>LDAPSSL: 2</code></p> <p>If your LDAP server uses TLS you must specify a file containing the certificate used to encrypt the password. By default the Trafodion software looks for this file in <code>\$TRAF_HOME/cacerts</code>, but you may specify a fully qualified filename, or set the environment variable <code>CACERTS_DIR</code> to another directory. To specify the file containing the certificate, you set the value of the attribute <code>TLS_CACERTFilename</code>, located in the Defaults section.</p> <p>Example</p> <pre> ` TLS_CACERTFilename: mycert.pem TLS_CACertFilename: /usr/etc/cert.pem ` </pre>

Setting	Description	Example
Search username and password	Some LDAP servers require a known user name and password to search the directory of user names. If your environment has that requirement, provide these "search" values.	LDAPSearchDN: lookup@company.com LDAPSearchPwd: Lookup123

There are additional optional attributes that can be used to customize Trafodion authentication. As mentioned earlier, they are described in [.traf_authentication_config](#) below.

You can test the authentication configuration file for syntactic errors using the `ldapconfigcheck` tool. If you have loaded the Trafodion environment (`sqenv.sh`), then the tool automatically checks the file at

`$TRAF_HOME/sql/scripts/.traf_authentication_config`. If not, you can specify the file to be checked.

Example

```
ldapconfigcheck -file myconfigfile
File myconfigfile is valid.
```

If an error is found, then the line number with the error is displayed along with the error. Please refer to [ldapconfigcheck](#) below for more information.



The authentication configuration file needs to be propagated to all nodes, but there is a script that does that for you described later. For now, you can test your changes on the local node.

You can test the LDAP connection using the utility `ldapcheck`. To use this utility the Trafodion environment must be loaded (`sqenv.sh`), but the Trafodion instance does not need to be running. To test the connection only, you can specify any user name, and a name lookup is performed using the attributes in `.traf_authentication_config`.

```
ldapcheck --username=fakename@company.com
User fakename@company.com not found
```

If `ldapcheck` reports either that the user was found or the user was not found, the connection was successful. However, if an error is reported, either the configuration file is not setup correctly, or there is a problem either with your LDAP server or the connection to the server. You can get additional error detail by including the `--verbose` option. Please refer to [ldapcheck](#) for more information.

If you supply a password, `ldapcheck` attempts to authenticate the specified username and password. The example below shows the password for illustrative purposes, but to avoid typing the password on the command line, leave the password blank (`--password=`) and the utility prompts for the password with no echo.

```
ldapcheck --username=realuser@company.com --password=StrongPassword  
Authentication successful
```

11.2.2. Generate Trafodion Certificate

Trafodion clients such as `trafci` encrypt the password before sending it to Trafodion. A self-signed certificate is used to encrypt the password. The certificate and key are generated when the `sqgen` script is invoked. By default, the files `server.key` and `server.crt` are located in `$HOME/sqcert`. If those files are not present and since Trafodion clients does not send unencrypted passwords, then you need to manually generate those files. To do so, run the script `sqcertgen` located in `$TRAF_HOME/sql/scripts`. The script runs `openssl` to generate the certificate and key.

To run `openssl` manually, follow the example:

```
openssl req -x509 -nodes -days 365 -subj
'/C=US/ST=California/L=PaloAlto/CN=host.domain.com/O=Some Company/OU=Service
Connection'
- newkey rsa:2048 -keyout server.key -out server.crt
```

Option	Description
-x509	Generate a self-signed certificate.
-days <validity of certificate>	Make the certificate valid for the days specified.
-newkey rsa:<bytes>	Generate a new private key of type RSA of length 1024 or 2048 bytes.
-subj <certificateinfo>	Specify the information that is incorporated in the certificate. Each instance in a cluster should have a unique common name(CN).
-keyout <filename>	Write the newly generated RSA private key to the file specified.
-nodes	It is an optional parameter that specifies NOT to encrypt the private key. If you encrypt the private key, then you must enter the password every time the private key is used by an application.
-out <filename>	Write the self-signed certificate to the specified file.

Both the public (`server.crt`) and private (`server.key`) files should be placed in the directory `$HOME/sqcert`. If you do not want to use the `HOME` directory or if you want to use different names for the private and/or public key files, then Trafodion supports environment variables to specific the alternate locations or names.

- Trafodion first checks the environment variables `SQCERT_PRIVKEY` and `SQCERT_PUBKEY`. If they are set, Trafodion uses the fully qualified filename value of the environment variable.

You can specify either one filename environment variable or both.

- If at least one filename environment variable is not set, Trafodion checks the value of the environment variable `SQCERT_DIR`. If set, then the default filename `server.key` or `server.crt` is appended to the value of the environment variable `SQCERT_DIR`.
- If the filename environment variable is not set and the directory environment variable is not set, then Trafodion uses the default location (`$HOME/sqcert`) and the default filename.

11.2.3. Creating the LDAP configuration file

The `.traf_authentication_config` file is user to enable the Trafodion security features.

File Location

By default, the Trafodion authentication configuration file is located in

`$TRAF_HOME/sql/scripts/.traf_authentication_config`. If you want to store the configuration file in a different location and/or use a different filename, then Trafodion supports environment variables to specify the alternate location/name.

Trafodion firsts checks the environment variable `TRAFAUTH_CONFIGFILE`. If set, the value is used as the fully-qualified Trafodion authentication configuration file.

If the environment variable is not set, then Trafodion next checks the variable `TRAFAUTH_CONFIGDIR`. If set, the value is prepended to `.traf_authentication_config` and used as the Trafodion authentication file.

If neither is set, Trafodion defaults to `$TRAF_HOME/sql/scripts/.traf_authentication_config`.

Template

```
# To use authentication in Trafodion, this file must be configured
# as described below and placed in $TRAF_HOME/sql/scripts and be named
# .traf_authentication_config.
#
# NOTE: the format of this configuration file is expected to change in the
# next release of Trafodion. Backward compatibility is not guaranteed.
#
SECTION: Defaults
    DefaultSectionName: local
    RefreshTime: 1800
    TLS_CACERTFilename:
SECTION: local

# If one or more of the LDAPHostName values is a load balancing host, list
# the name(s) here, one name: value pair for each host.
    LoadBalanceHostName:

# One or more identically configured hosts must be specified here,
# one name: value pair for each host.
    LDAPHostName:

# Default is port 389, change if using 636 or any other port
    LDAPPort:389

# Must specify one or more unique identifiers, one name: value pair for each
    UniqueIdentifier:

# If the configured LDAP server requires a username and password to
# to perform name lookup, provide those here.
    LDAPSearchDN:
    LDAPSearchPwd:

# If configured LDAP server requires TLS(1) or SSL (2), update this value
    LDAPSSL:0

# Default timeout values in seconds
    LDAPNetworkTimeout: 30
    LDAPTimeout: 30
    LDAPTimeLimit: 30

# Default values for retry logic algorithm
    RetryCount: 5
    RetryDelay: 2
    PreserveConnection: No
    ExcludeBadHosts: Yes
    MaxExcludeListSize: 3
```

Configuration Attributes

Attribute Name	Purpose	Example Value	Notes
LDAPHostName	Host name of the local LDAP server.	ldap.master.com	If more than one LDAPHostName entry is provided, then Trafodion attempts to connect with each LDAP server before returning an authentication error. Also see the description related to RetryCount and RetryDelay entries.
LDAPPort	Port number of the local LDAP server.	345	Must be numeric value. Related to LDAPSSL entry. Standard port numbers for OpenLDAP are as follows: <ul style="list-style-type: none"> • Non-secure: 389 • SSL: 636 • TLS: 389
LDAPSearchDN	If a search user is needed, the search user distinguished name is specified here.	cn=aaabbb, dc=demo, dc=net	If anonymous search is allowed on the local server, then this attribute does not need to be specified or can be specified with no value (blank). To date, anonymous search is the normal approach used.
LDAPSearchPWD	Password for the LDAPSearchDN value. See that entry for details.	welcome	None.
LDAPSSL	A numeric value specifying whether the local LDAP server interface is unencrypted or TLS or SSL. Legal values are 0 for unencrypted, 1 for SSL, and 2 for TLS. For SSL/TLS, see the section below on Encryption Support.	0	None.
UniqueIdentifier	The directory attribute that contains the user's unique identifier.	uid=,ou=Users,dc=demo,dc=net	To account for the multiple forms of DN supported by a given LDAP server, specify the UniqueIdentifier parameter multiple times with different values. During a search, each UniqueIdentifier is tried in the order it is listed in the configuration file.

Attribute Name	Purpose	Example Value	Notes
LDAPNetworkTimeout	Specifies the timeout (in seconds) after which the next LDAPHostName entry is tried, in case of no response for a connection request. This parameter is similar to NETWORK_TIMEOUT in ldap_conf(5). Default value is 30 seconds.	20	The value must be a positive number or -1. Setting this to -1 results in an infinite timeout.
LDAPTimeLimit	Specifies the time to wait when performing a search on the LDAP server for the user name. The number must be a positive integer. This parameter is similar to TIMELIMIT in ldap_conf(5). Default value is 30 seconds.	15	The server may still apply a lower server-side limit on the duration of a search operation.
LDAPTimeout	Specifies a timeout (in seconds) after which calls to synchronous LDAP APIs aborts if no response is received. This parameter is similar to TIMEOUT in ldap_conf(5). Default value is 30 seconds.	15	The value must be a positive number or -1. Setting this to -1 results in an infinite timeout.
RetryCount	Number of attempts to establish a successful LDAP connection. Default is 5 retries before returning an error.	10	When a failed operation is retried, it is attempted with each configured LDAP server, until the operation is successful or the number of configured retries is exceeded.
RetryDelay	Specifies the number of seconds to delay between retries. Default value is 2 seconds. See description of RetryCount.	1	None.
PreserveConnection	Specifies whether the connection to LDAP server is maintained (YES) or closed (NO) once the operation finishes. Default value is NO.	YES	None.

Attribute Name	Purpose	Example Value	Notes
RefreshTime	Specifies the number of seconds that must have elapsed before the configuration file is reread. Default is 1800 (30 minutes).	3600	If set to zero, the configuration file is never read. The connectivity servers must be restarted for changes to take effect if this value is zero. This attribute is not specific to either configuration and must be defined in the DEFAULTS section.
TLS_CACERTFilename	Specifies the location of the certificate file for the LDAP server(s). Filename can either be fully qualified or relative to \$CACERTS_DIR.	cert.pem	This attribute applies to both configurations. If a configuration does not require a certificate, then this attribute is ignored. This attribute must be defined in the DEFAULTS section.
DefaultSectionName	Specifies the configuration type that is assigned to a user by the REGISTER USER command if no authentication type is specified. In the initial Trafodion release, only one configuration is supported.	LOCAL	This attribute must be defined in the DEFAULTS section. If the DefaultSectionName attribute is specified, then a section by that name (or equivalent) must be defined in .traf_ldapconfig. Legal values are LOCAL and ENTERPRISE. This syntax is likely to change.

11.2.4. Verifying configuration and users through ldapcheck

Usage

```

ldapcheck  [<option>]...
<option> ::= --help|-h           display usage information
            --username=<LDAP-username>
            --password[=<password>]
            --primary             Use first configuration
            --local               Use first configuration
            --enterprise          Use first configuration
            --secondary          Use second configuration
            --remote              Use second configuration
            --cluster             Use second configuration
            --verbose             Display non-zero retry counts
                                and LDAP errors

```

Considerations

- Aliases for primary include enterprise and local. Aliases for secondary include cluster and remote. If no configuration is specified, primary is assumed.
- The equals sign is required when supplying a value to username or password.

- To be prompted for a password value with no echo, specify the password argument but omit the equals sign and value.
- Passwords that contain special characters may need to be escaped if the password is specified on the command line or within a script file.
- If the password keyword is not specified, only the username is checked. The tool can therefore be used to test the LDAP configuration and connection to the configured LDAP server(s) without knowing a valid username or password.

11.2.5. Verifying contents of configuration file through ldapconfigcheck

This page describes the `ldapconfigcheck` tool, which validates the syntactic correctness of a Trafodion authentication configuration file. Trafodion does not need to be running to run the tool.

Considerations

If the configuration filename is not specified, then the tool looks for a file using environment variables. Those environment variables and the search order are:

1. `TRAFAUTH_CONFIGFILE`

A fully qualified name is expected.

2. `TRAFAUTH_CONFIGDIR`

Filename `.traf_authentication_config/` is appended to the specified directory

3. `TRAF_HOME`

`/sql/scripts/.traf_authentication_config` is appended to the value of `TRAF_HOME`.

Errors

One of the following is output when the tool is run. Only the first error encountered is reported.

Code	Text
0	File <code>filename</code> is valid.
1	File <code>filename</code> not found.
2	File: <code>filename</code>
	Invalid attribute name on line line-number .
3	File: <code>filename</code>
	Missing required value on line line-number .
4	File: <code>filename</code>
	Value out of range on line line-number .
5	File: <code>filename</code>
	Open of <code>traf_authentication_config</code> file failed.
6	File: <code>filename</code>
	Read of <code>traf_authentication_config</code> file failed.
7	No file provided. Either specify a file parameter or verify environment variables.
8	TLS was requested in at least one section, but <code>TLS_CACERTFilename</code> was not provided.
9	Missing host name in at least one section.
	Each LDAP connection configuration section must provide at least one host name.
10	Missing unique identifier in at least one section.
	Each LDAP connection configuration section must provide at least one unique identifier.
11	At least one LDAP connection configuration section must be specified.
12	Internal error parsing <code>.traf_authentication_config</code> .

11.3. Manage Users

Kerberos is enabled for installations that require a secure Hadoop environment. LDAP is enabled to enforce authentication for any user connecting to Trafodion. The Trafodion database enforces privileges on the database, database schemas, database objects (table, views, etc) and database operations. Privileges are enforced when authorization is enabled. When LDAP or Kerberos is enabled, authorization is automatically enabled.

To determine the status of authentication and authorization, bring up `sqlci` and perform "env;".

```
>>env;
-----
Current Environment
-----
AUTHENTICATION      enabled
AUTHORIZATION        enabled
CURRENT DIRECTORY    /.../incubator-trafodion/install/installer
LIST_COUNT           4294967295
LOG FILE
MESSAGEFILE          /.../incubator-trafodion/core/sqf/export/ ...
MESSAGEFILE LANG     US English
MESSAGEFILE VRSN     {2016-06-14 22:27 LINUX:host/user}
SQL CATALOG          TRAFODION
SQL SCHEMA           SCH
SQL USER CONNECTED   user not connected
SQL USER DB NAME     SQLUSER1
SQL USER ID          33367
TERMINAL CHARSET     ISO88591
TRANSACTION ID
TRANSACTION STATE    not in progress
WARNINGS             on
```

Once authorization is enabled, there is one predefined database user called DB__ROOT associated with your specified LDAP username. Please connect to the database and this user and register users that will perform database administration. The database admin can then connect and setup required users, roles, and privileges.

TBD - A future update should include a pointer to the security best practices guide.

To learn more about how to register users, grant object and component privileges, and manage users and roles, please see the [Trafodion SQL Reference Manual](#).